

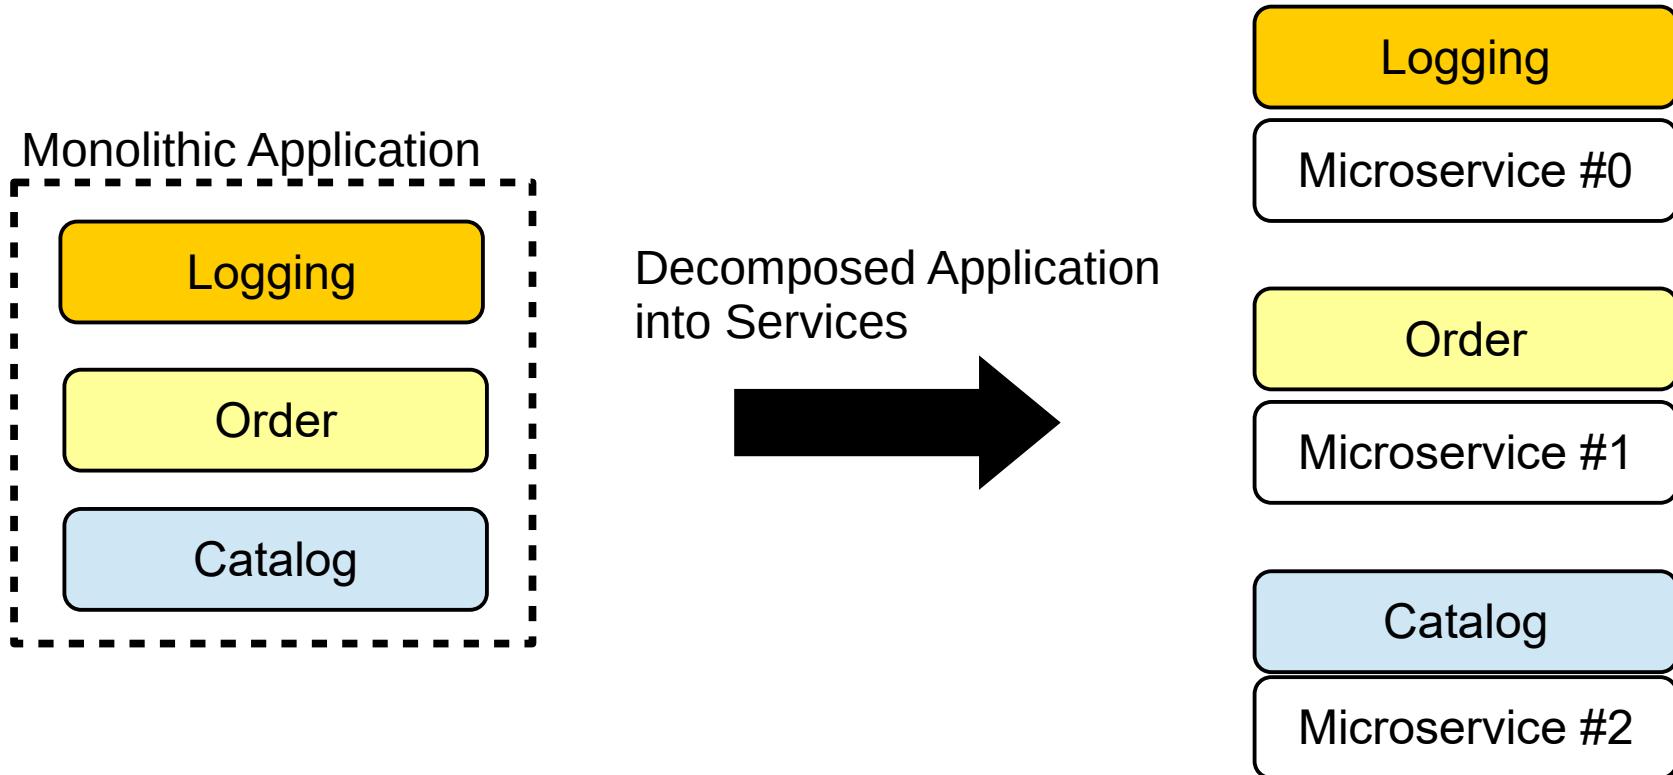


Leveraging Virtio-fs and Virtio-vsocket in Toro Unikernel

www.torokernel.io

Matias Vara Larsen
matiasevara@gmail.com

What are microservices?



How are microservices deployed?

Each VM requires its own OS

The use of VMs to host microservices allows to isolate different services

Microservice #0

Service instance per VM

Microservice #1

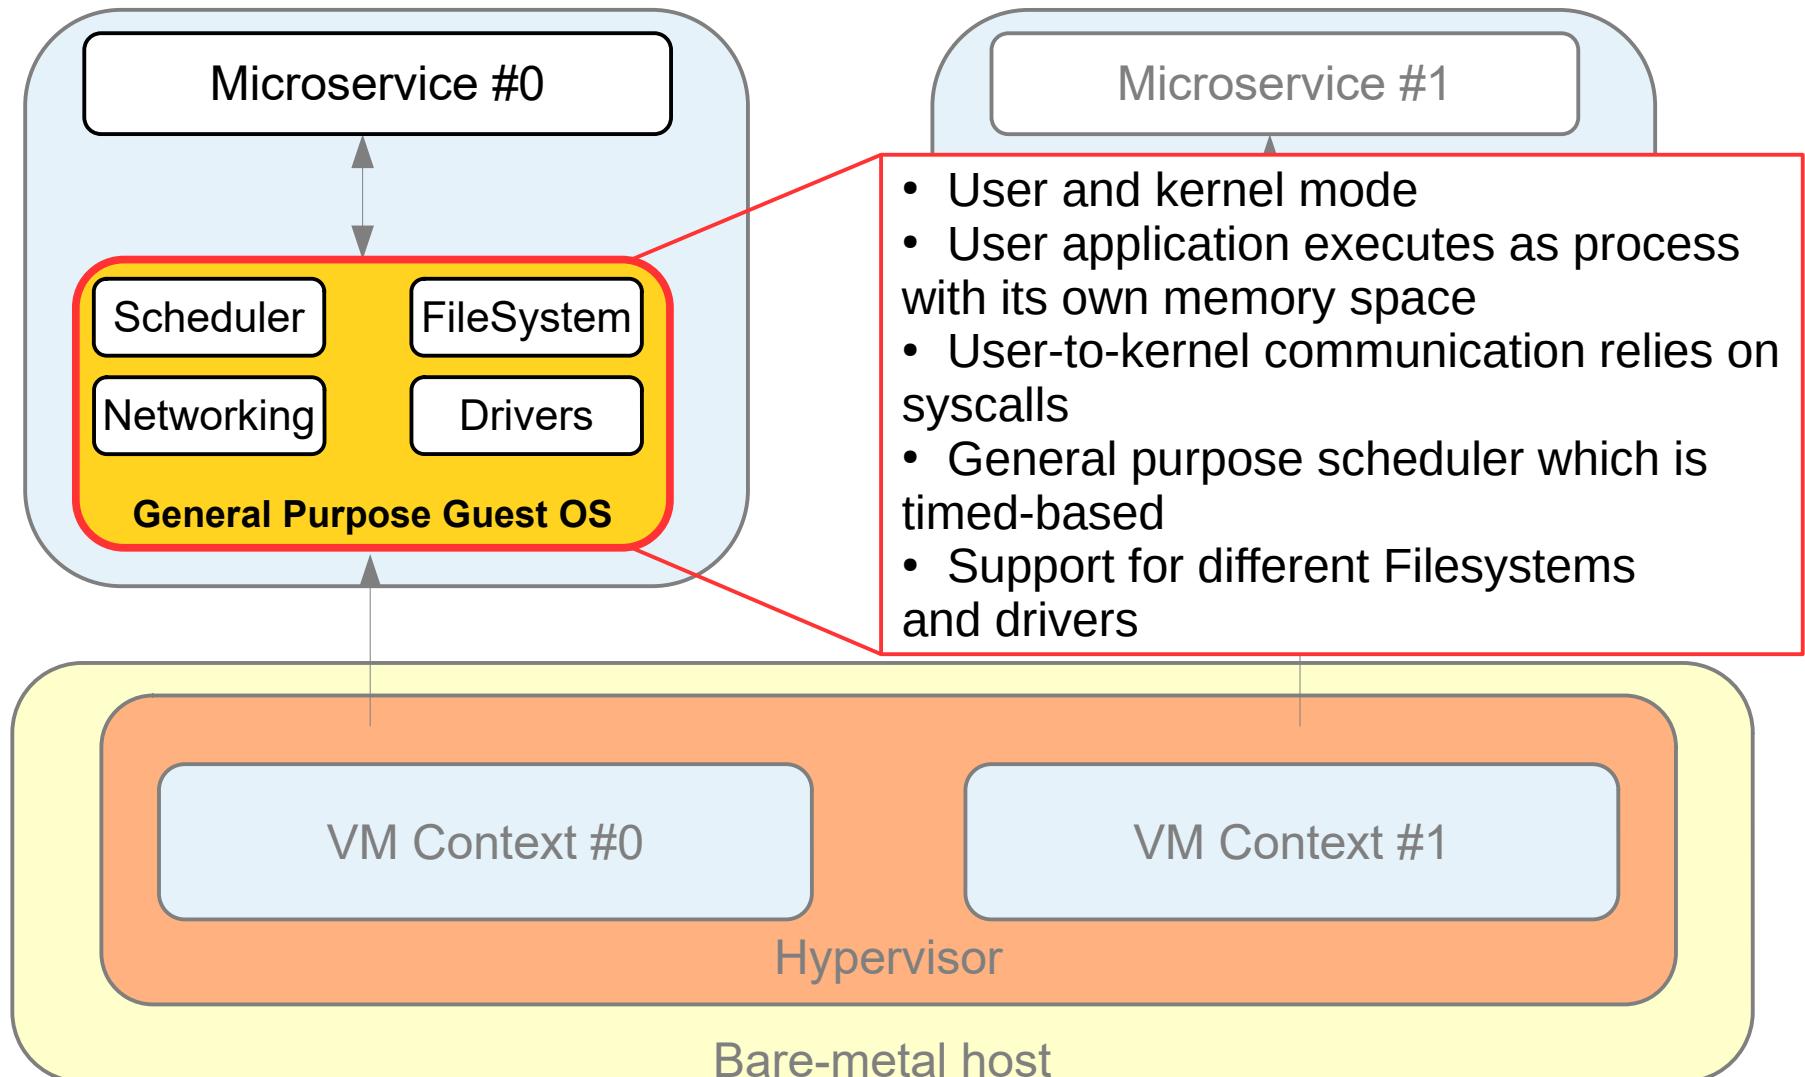
VMs are hosted on a Baremetal Host or as instances on a Cloud Provider, e.g, AWS or GCE

VM Context #0

VM Context #1

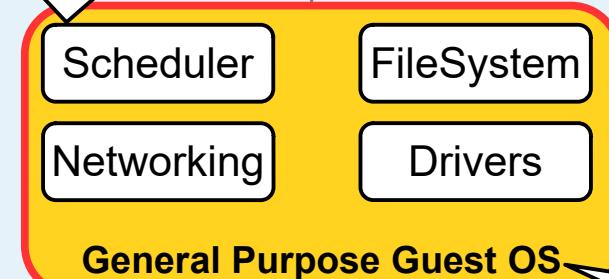
Hypervisor

Bare-metal host

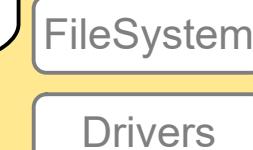


Guests consume a lot of resources, e.g., memory, cpu, on-disk image

#0



A different set of drivers may be needed depending on the Cloud provider device model



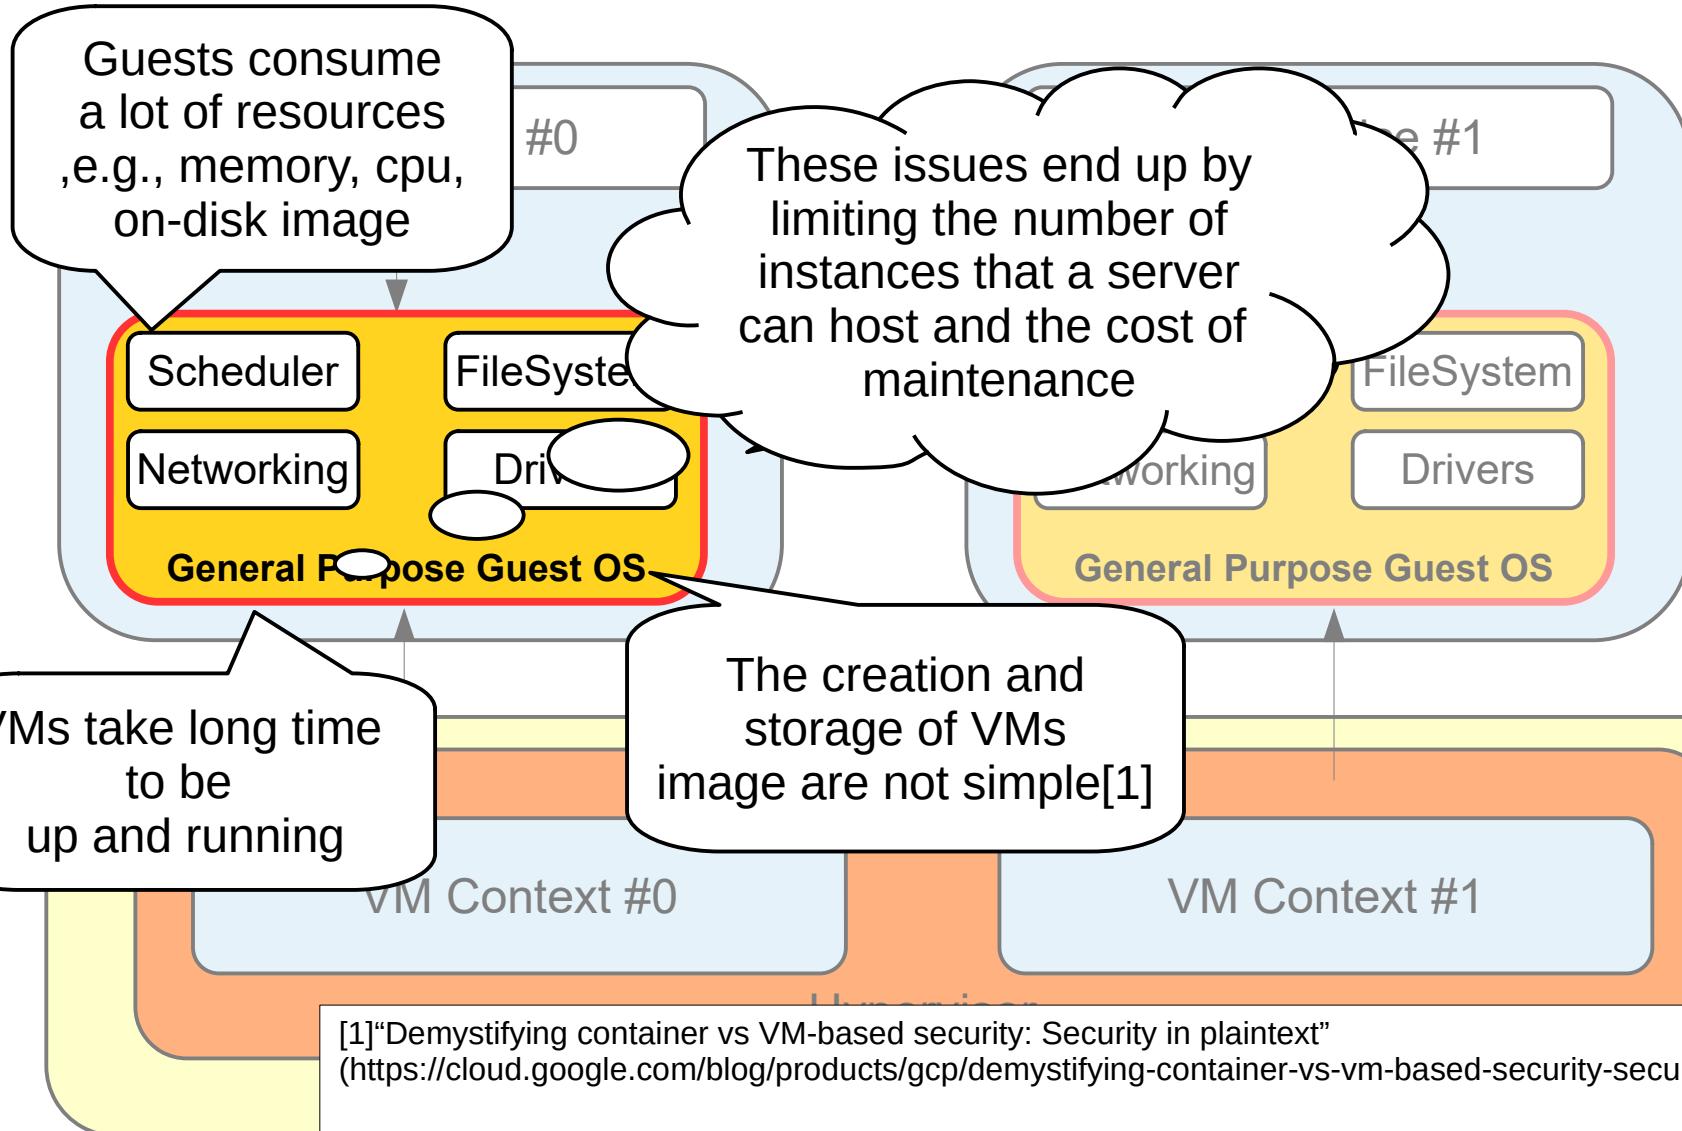
VMs take long time to be up and running

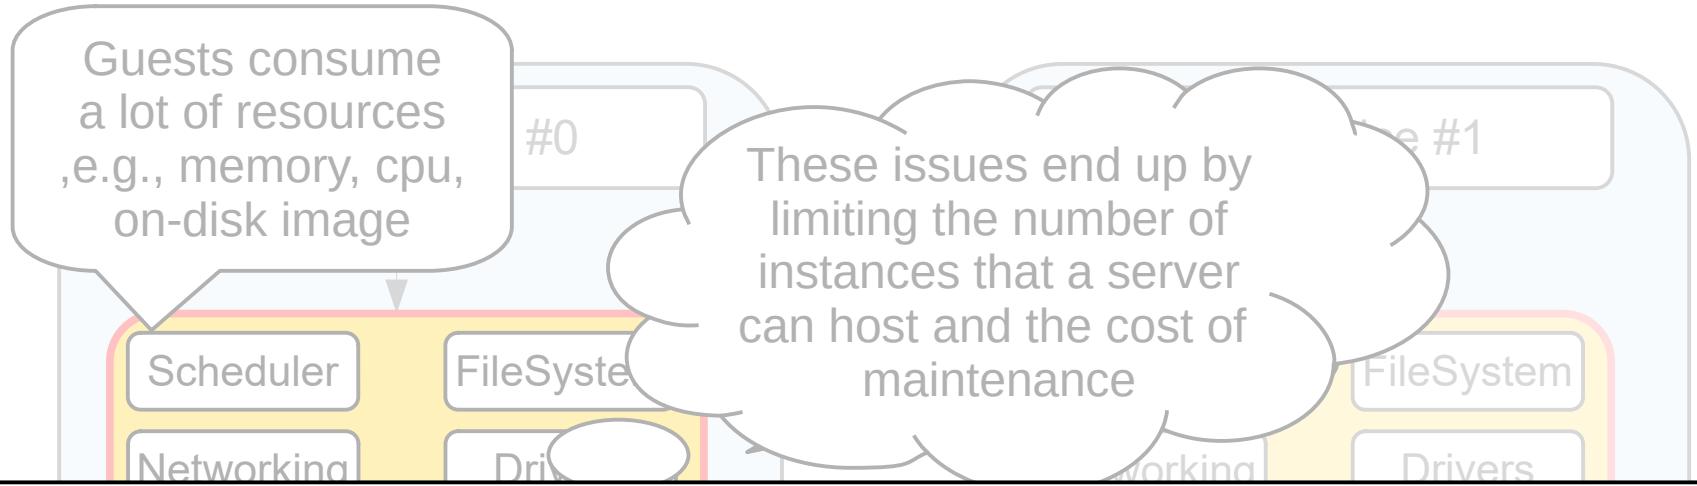
VM Context #0

The creation and storage of VMs image are not simple[1]

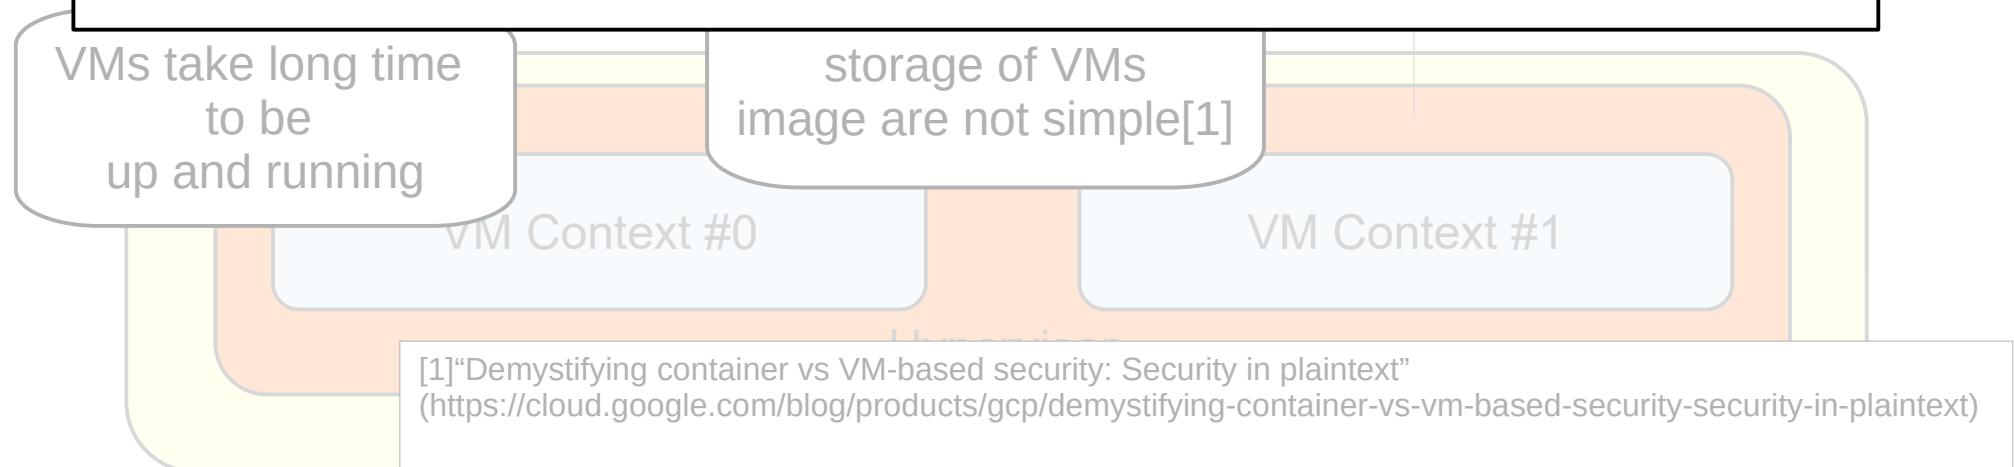
VM Context #1

[1]“Demystifying container vs VM-based security: Security in plaintext”
(<https://cloud.google.com/blog/products/gcp/demystifying-container-vs-vm-based-security-security-in-plaintext>)

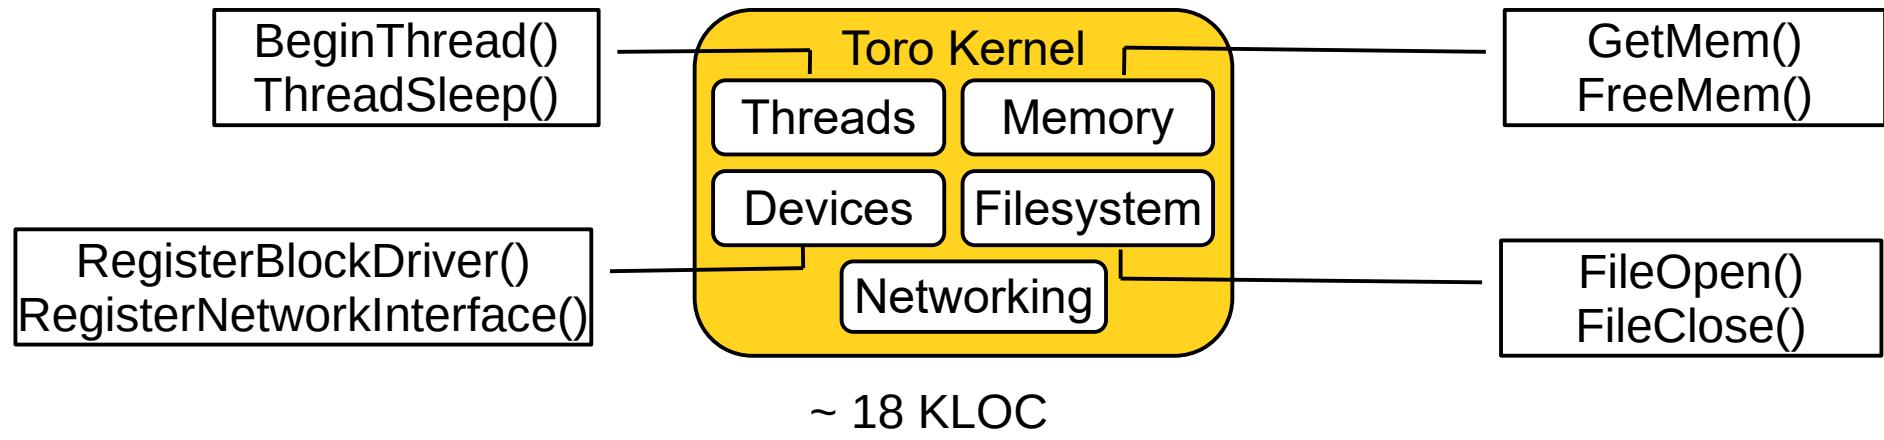




Toro is an application-oriented unikernel that allows **microservices** to run efficiently in **VMs** thus leveraging the strong isolation VMs provide.



Application-oriented Kernel

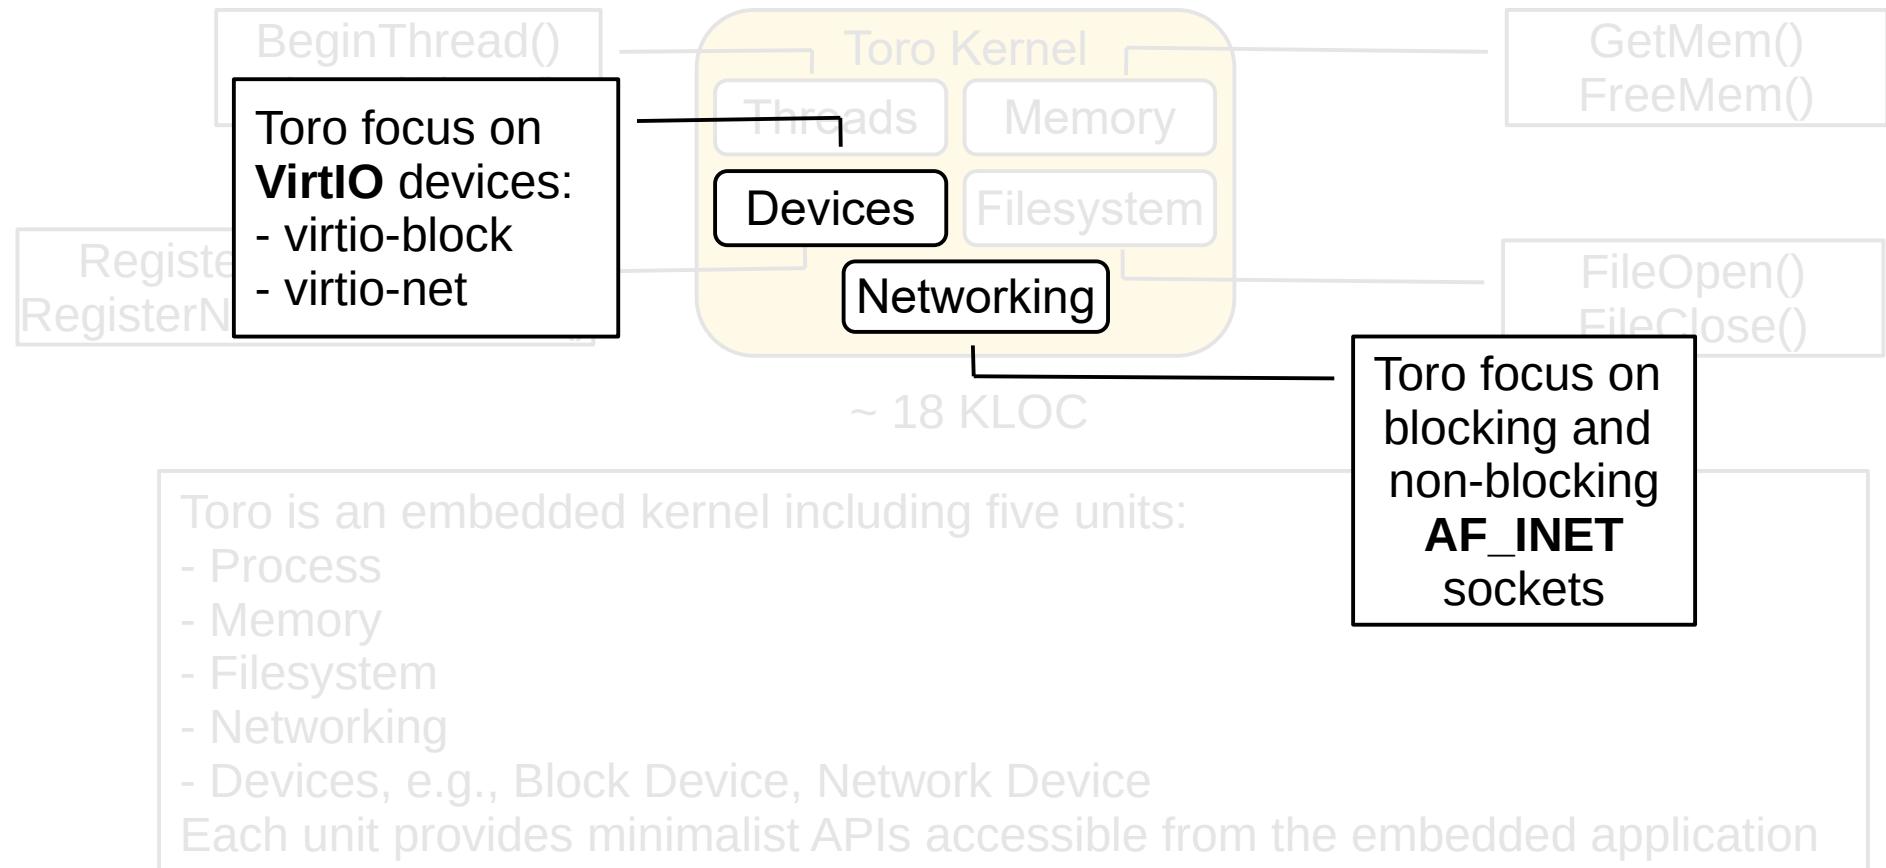


Toro is an embedded kernel including five units:

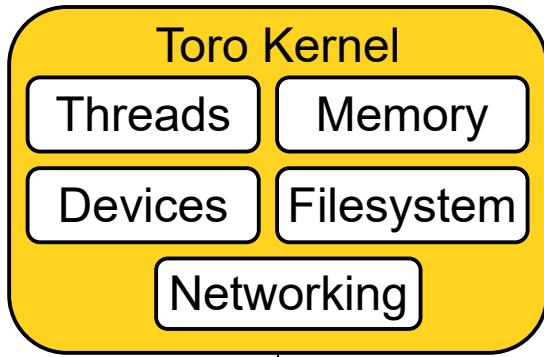
- Process
- Memory
- Filesystem
- Networking
- Devices, e.g., Block Device, Network Device

Each unit provides minimalist APIs accessible from the embedded application

Application-oriented Kernel



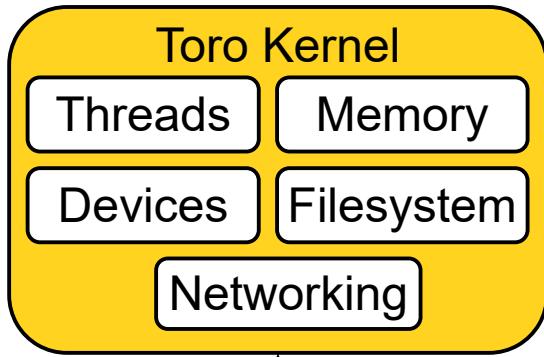
Application-oriented Kernel



- User application and kernel units are compiled in a single binary
- The application includes only the component required

Microservice

Application-oriented Kernel



- User application and kernel units are compiled in a single binary
- The command

program WebServerAppliance;
uses

Memory,
Filesystem,
Threads,
Networking,
Fat,
Virtio-blk,
Virtio-net;

Begin

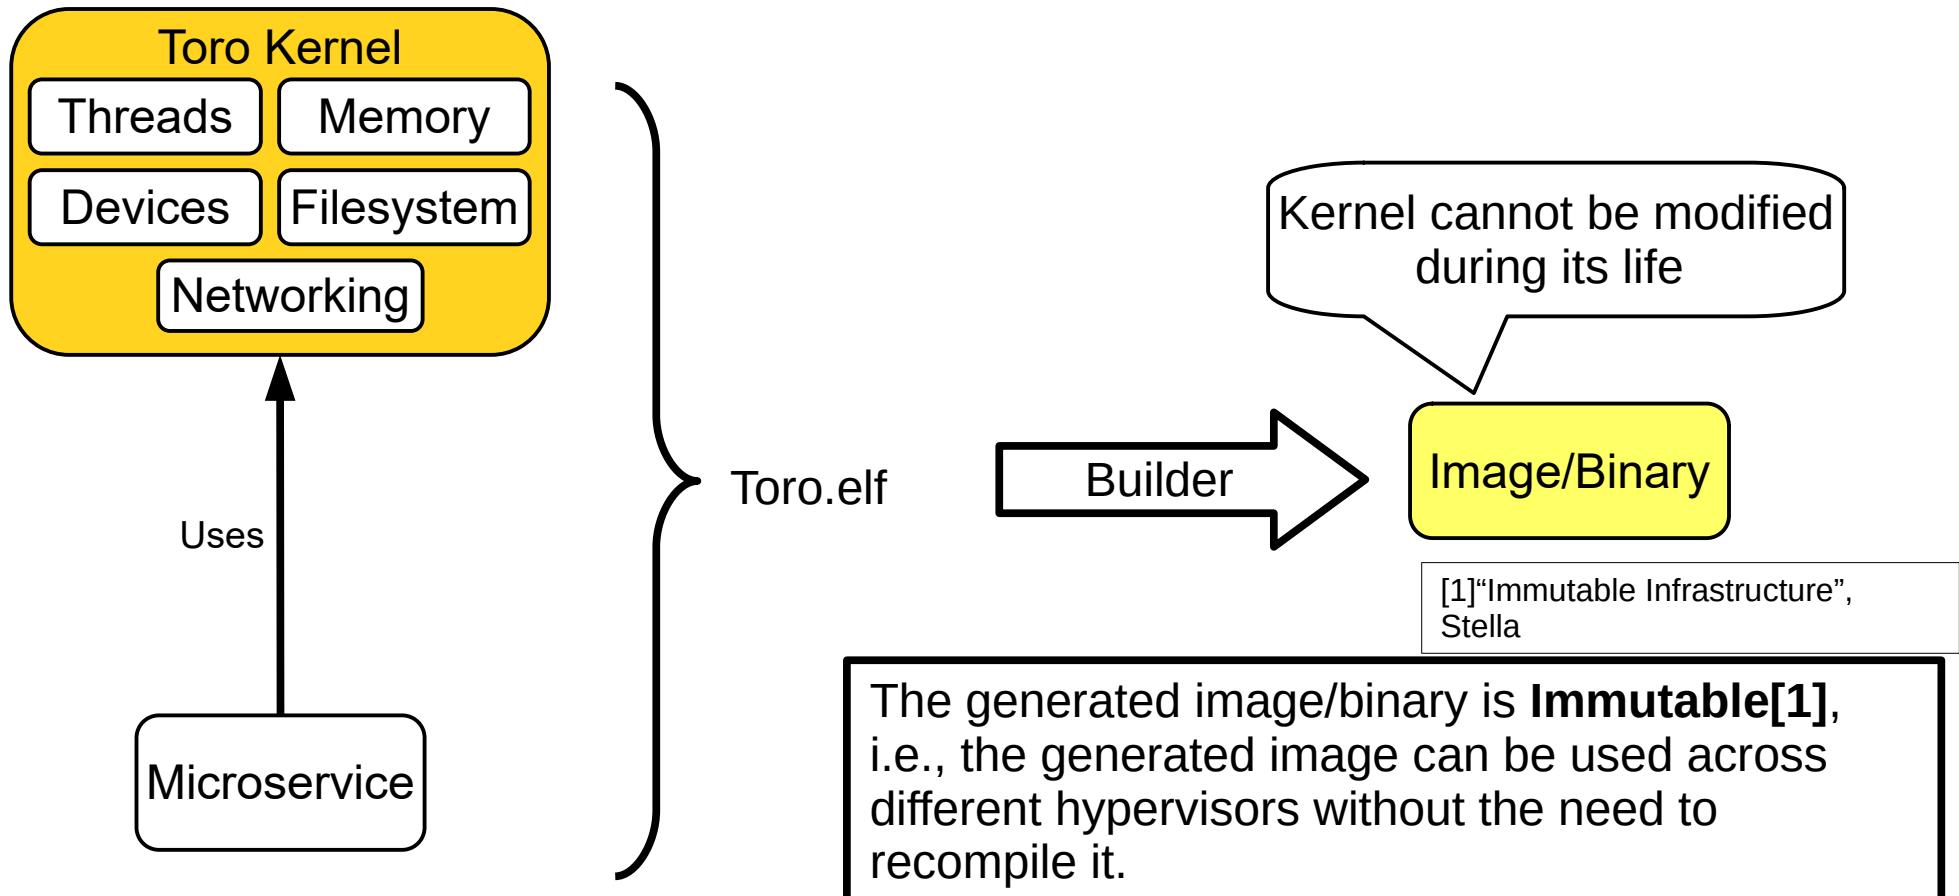
//

// Your Code Goes Here

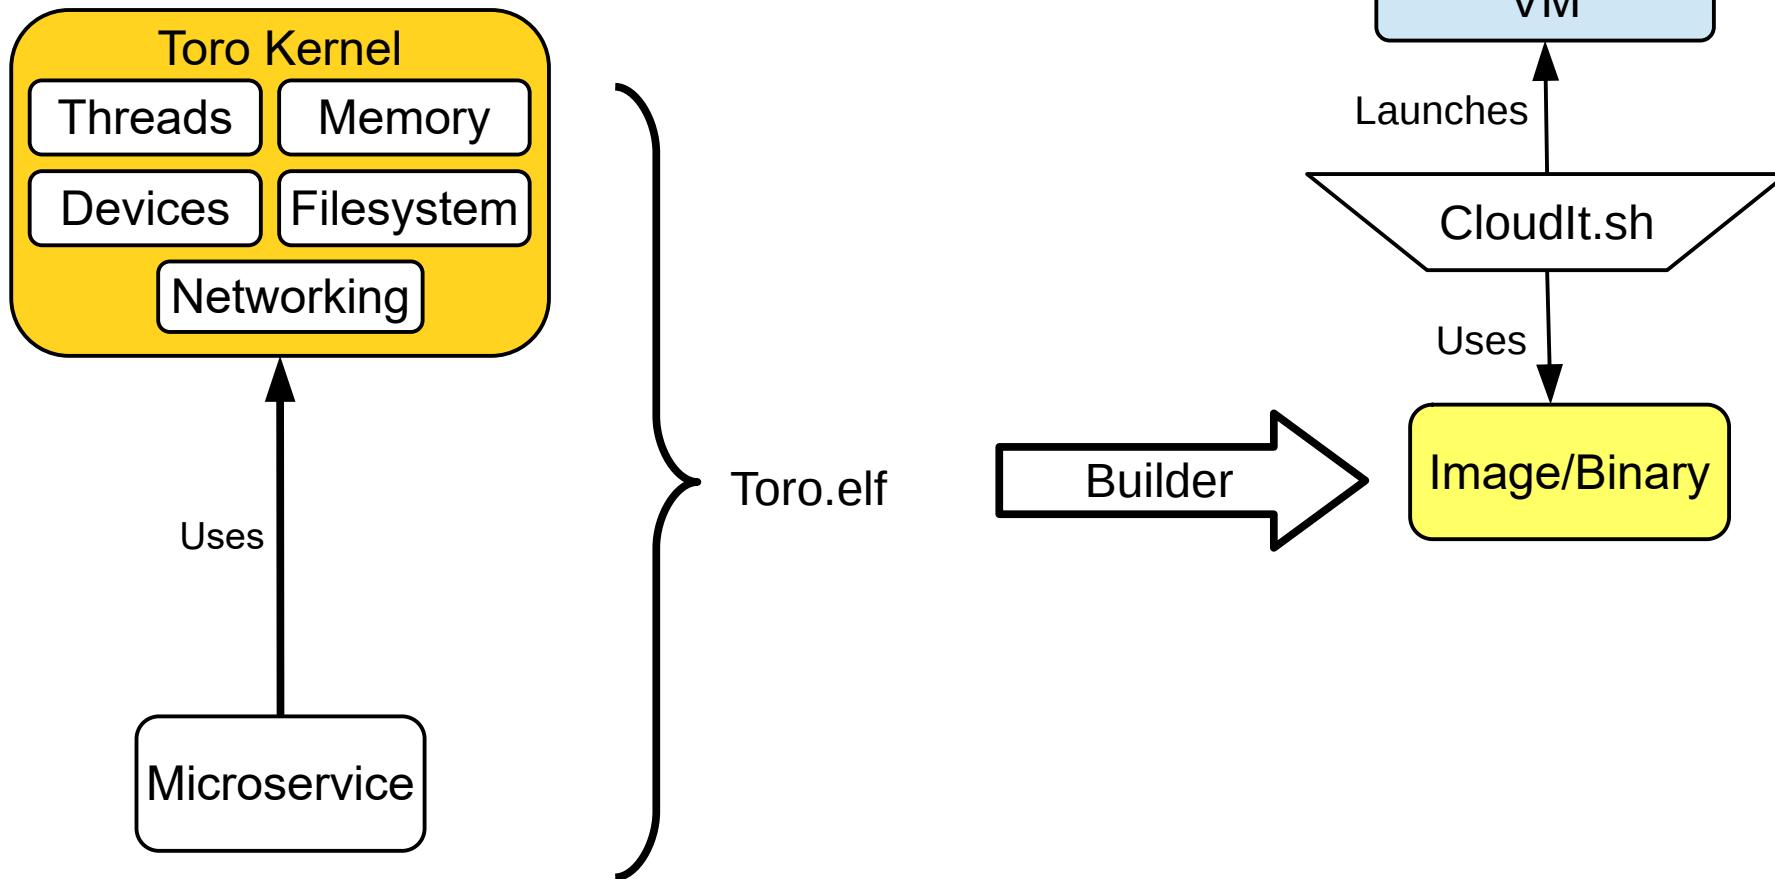
//

End.

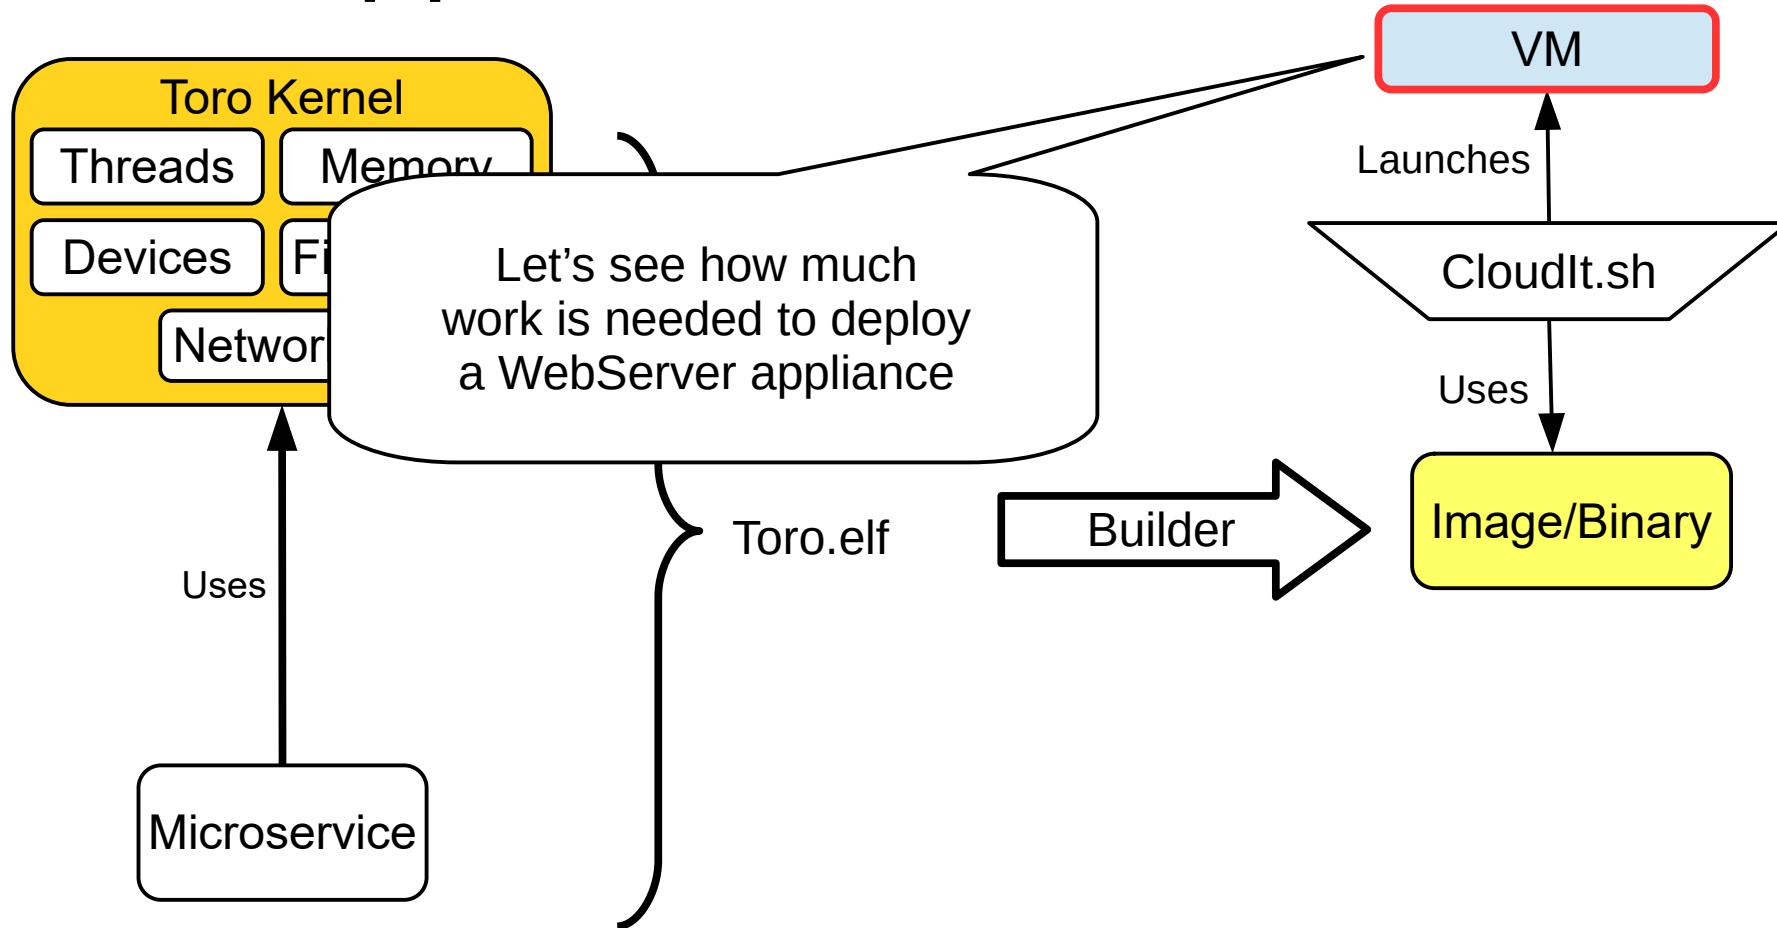
Application-oriented Kernel



Application-oriented Kernel



Application-oriented Kernel

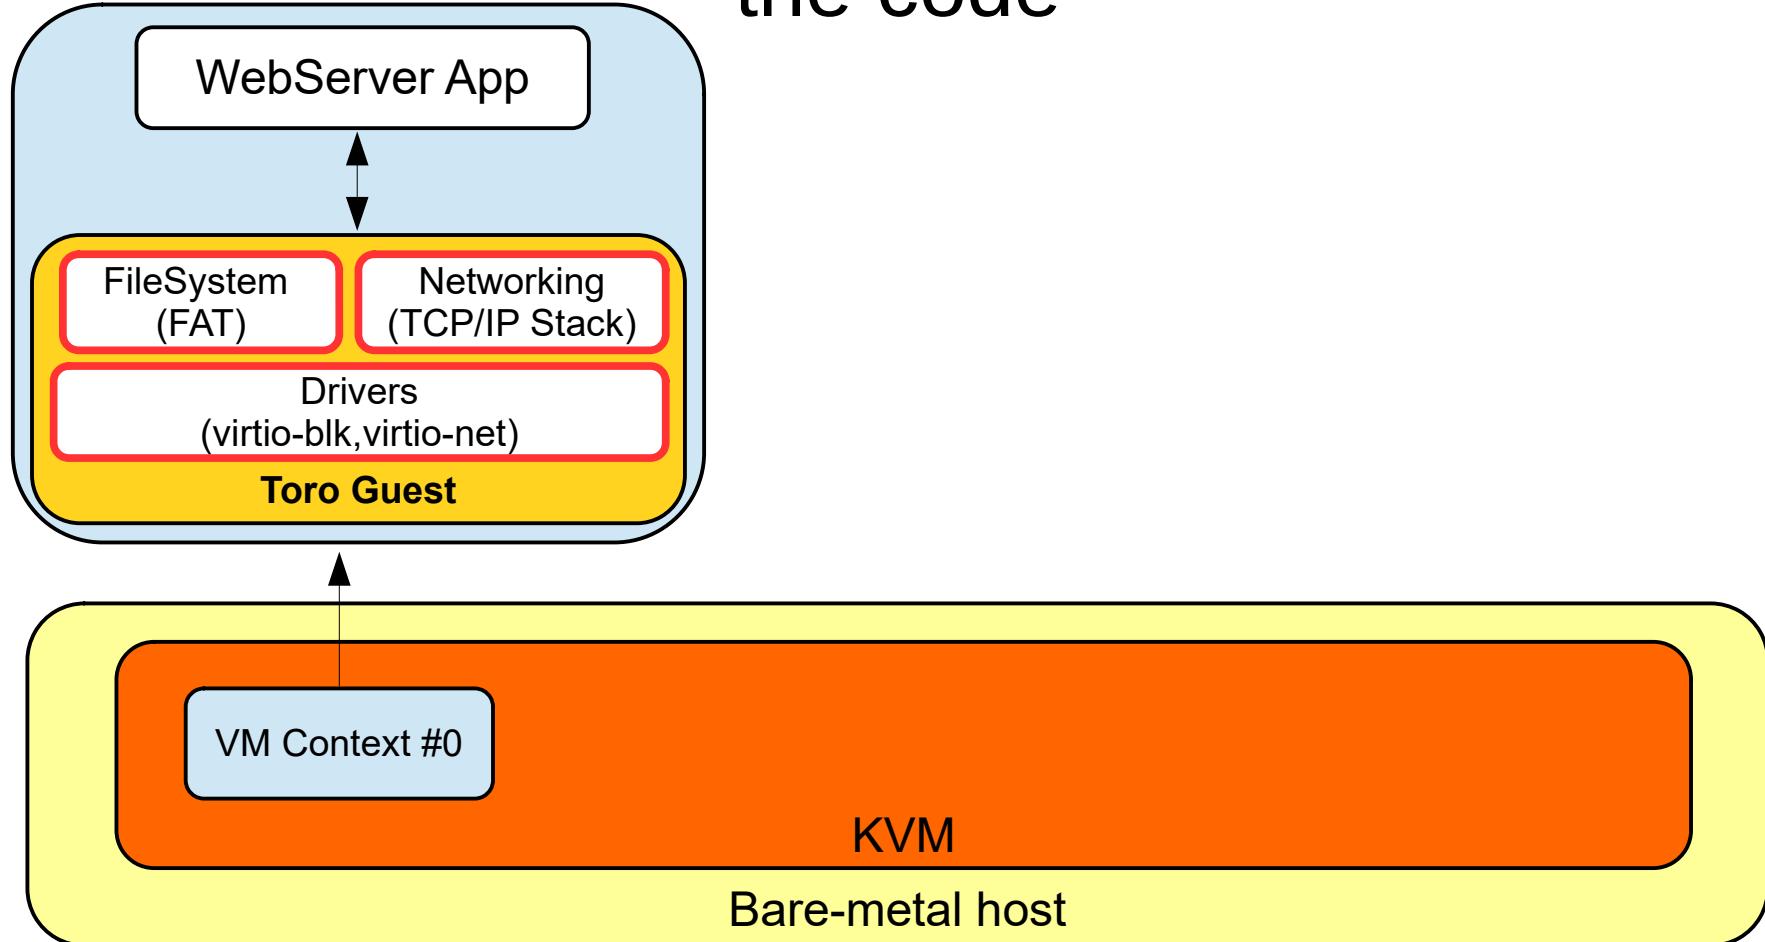


The WebServer Appliance

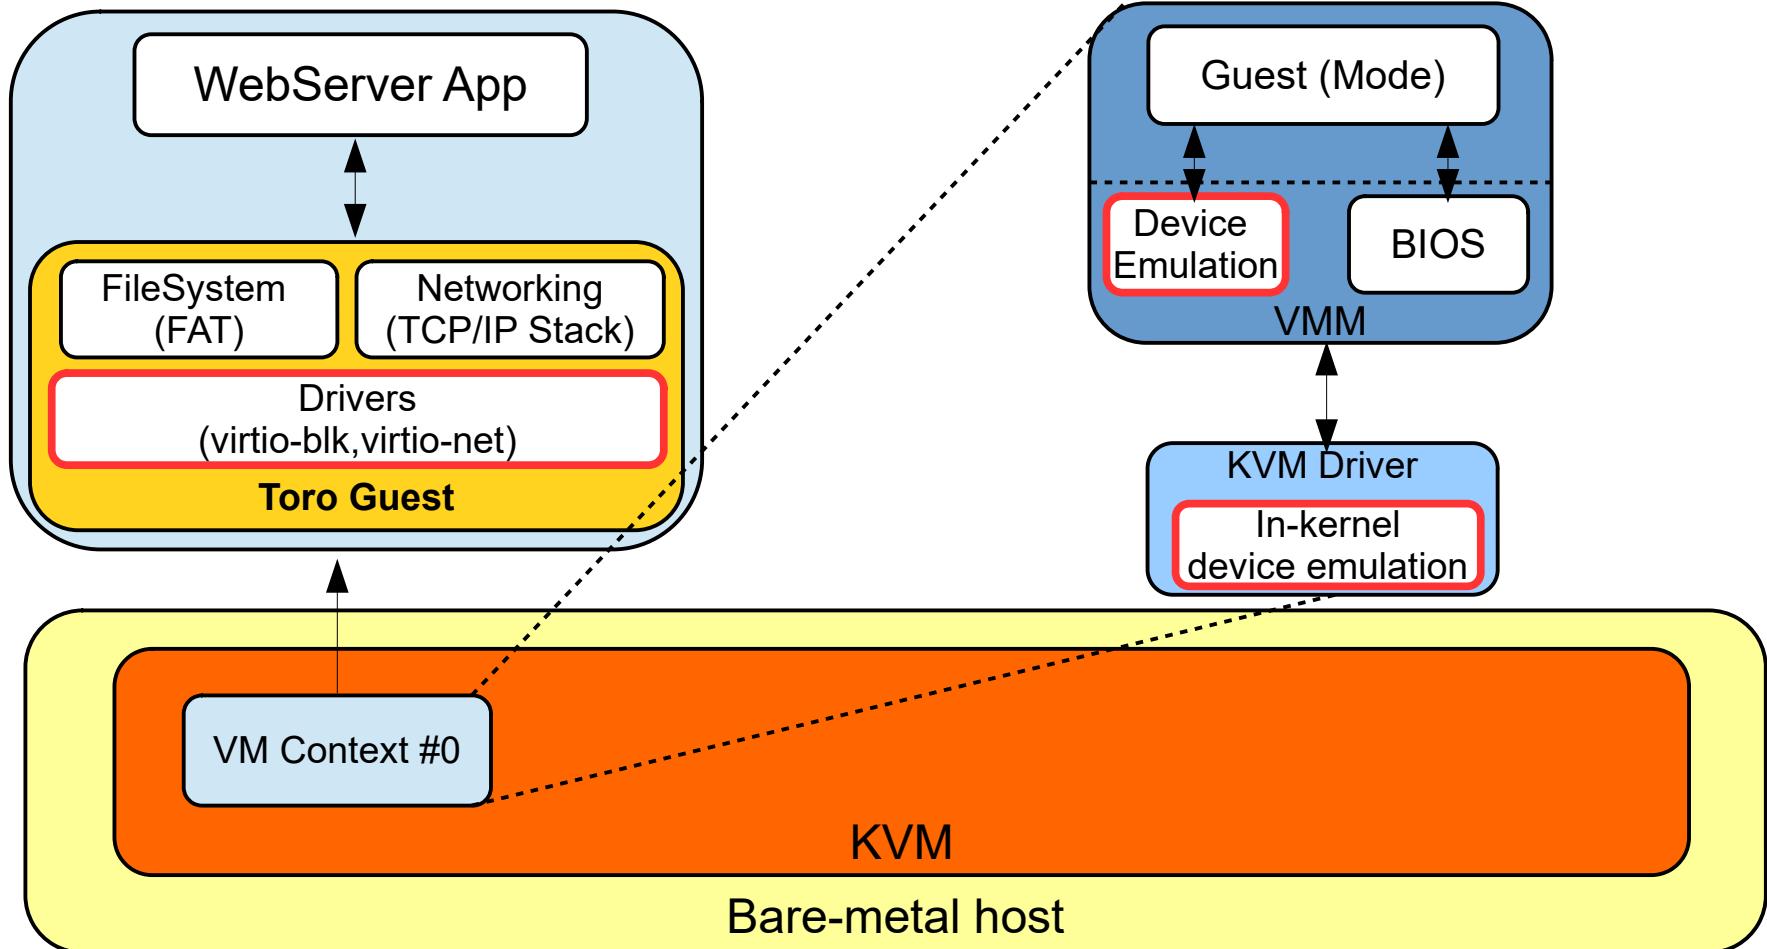
- Simple microservice that serves files by using the HTTP protocol
 - Find it at <https://github.com/torokernel/torokernel> among other examples
 - This appliance is used to host Toro's website (<http://www.torokernel.io> and click on “View on Toro”)

How the appliance is setup?

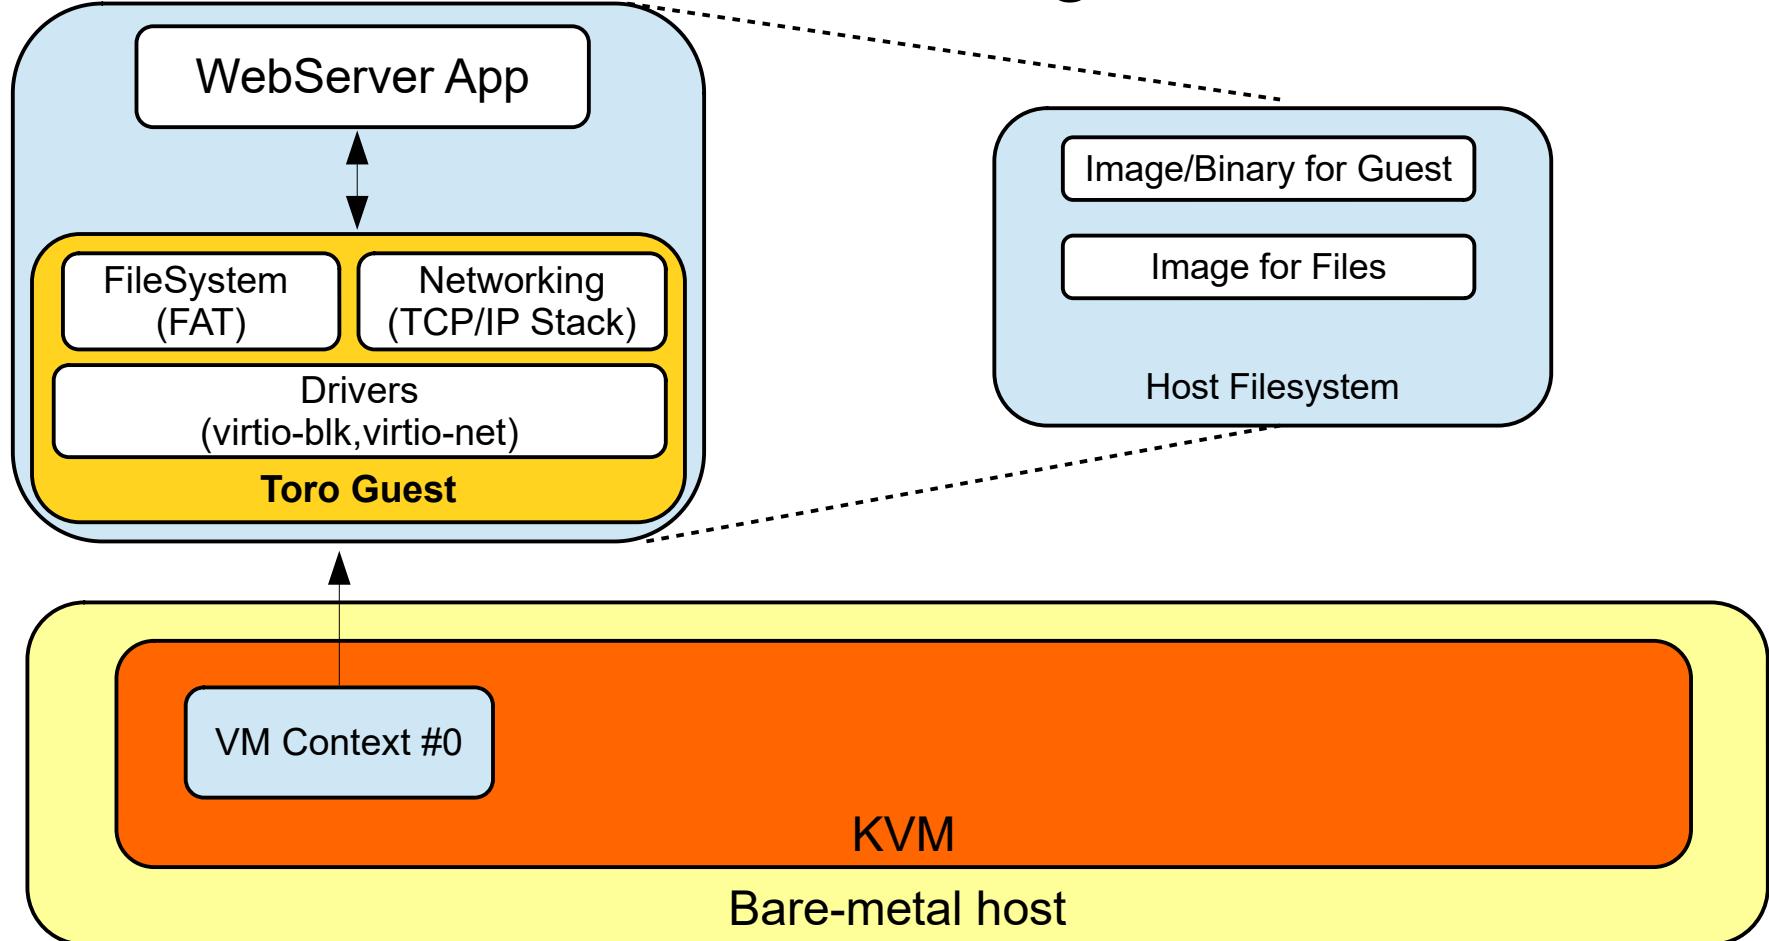
the code



How the appliance is setup? device model

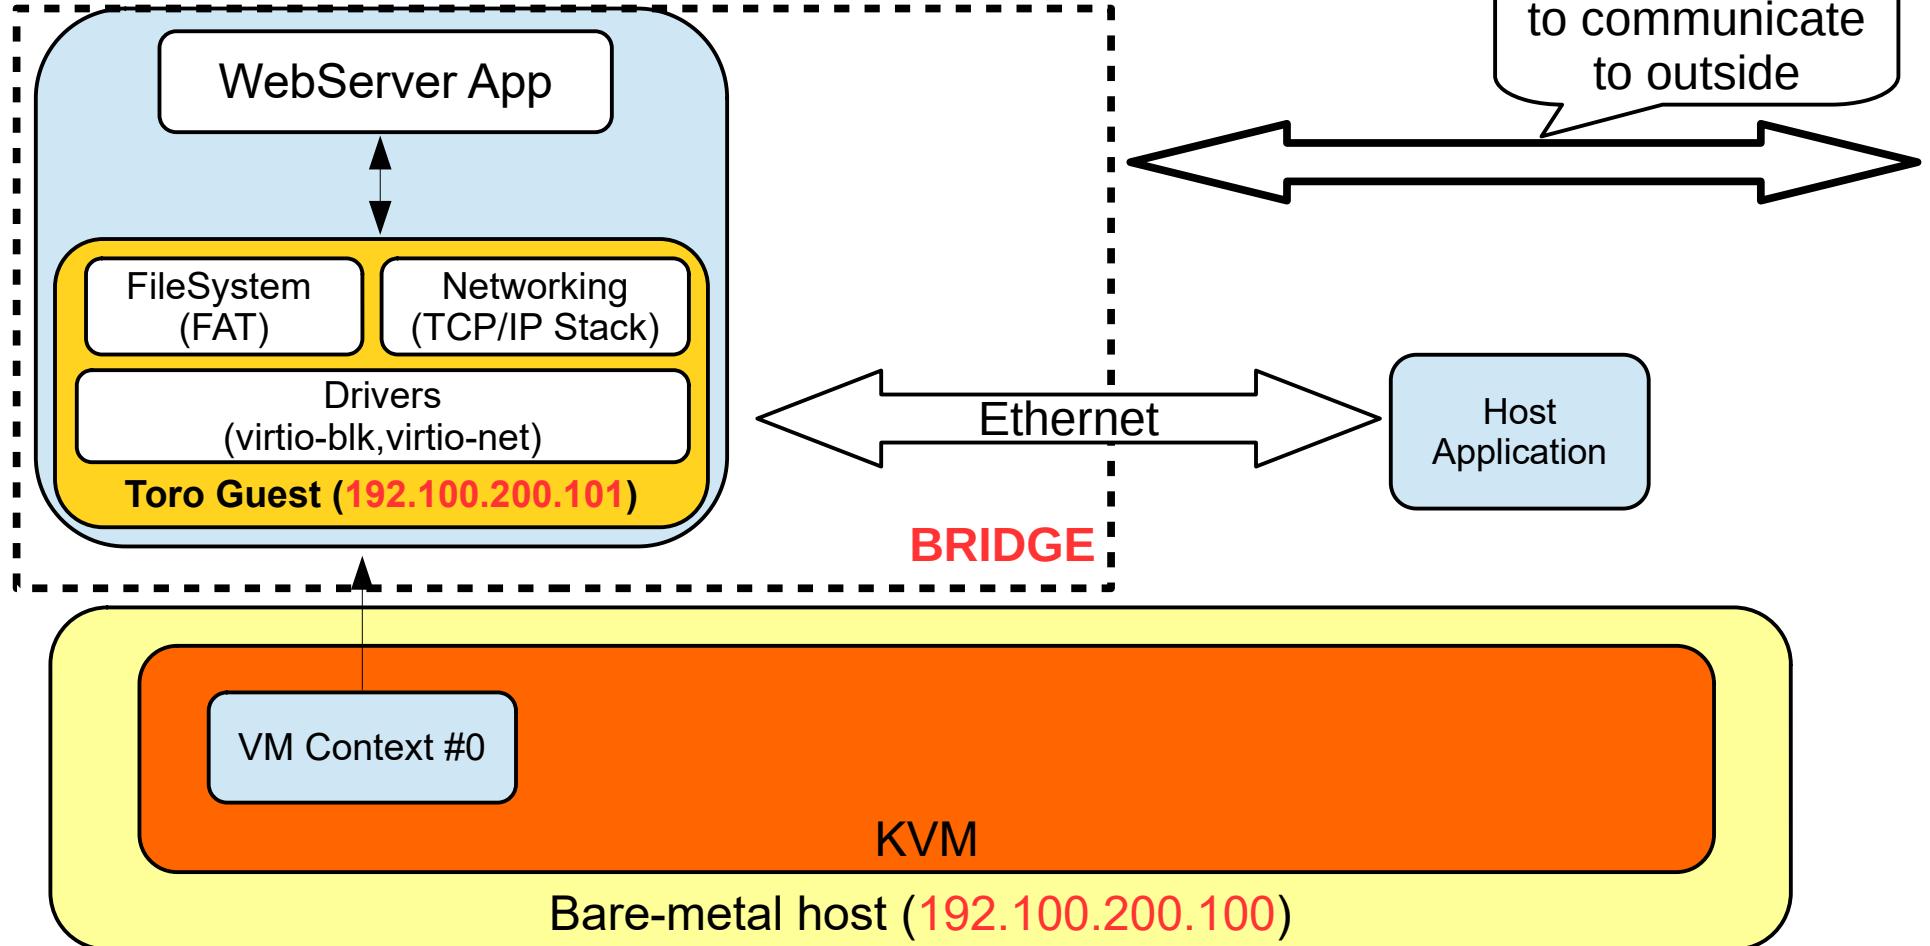


How the appliance is setup? on-disk images



How the appliance is setup?

ethernet configuration



The Static WebServer

drawbacks

- Disk images consume memory and on-disk space, e.g., each guest has its own image
- The use of a TCP/IP stack requires configurations, e.g., bridge, an IP per guest, guest drivers, devices
- The use of more devices increases the attack surface
- Sharing of files between guests and host is hard
- Relying on a specific FS is not good for immutable images

The Static WebServer

drawbacks

- Disk images consume memory and on-disk space, e.g., each guest has its own image
- The use of a TCP/IP stack requires configurations, e.g., bridge, an IP per guest, guest drivers, devices
- The use of more devices increases the attack surface
- Sharing of files between guests and host is hard
- Relying on a specific FS is not good for ~~im~~



Can we do it better?!

The Static WebServer drawbacks

- Disk images consume memory and on-disk space, e.g., each guest has its own image
- The use of a TCP/IP stack requires configuration, e.g.

We propose to use of **virtio-fs** for filesystem and **virtio-vsocket** for networking to simplify toro unikernel code, reduce attack surface and ease appliance configuration

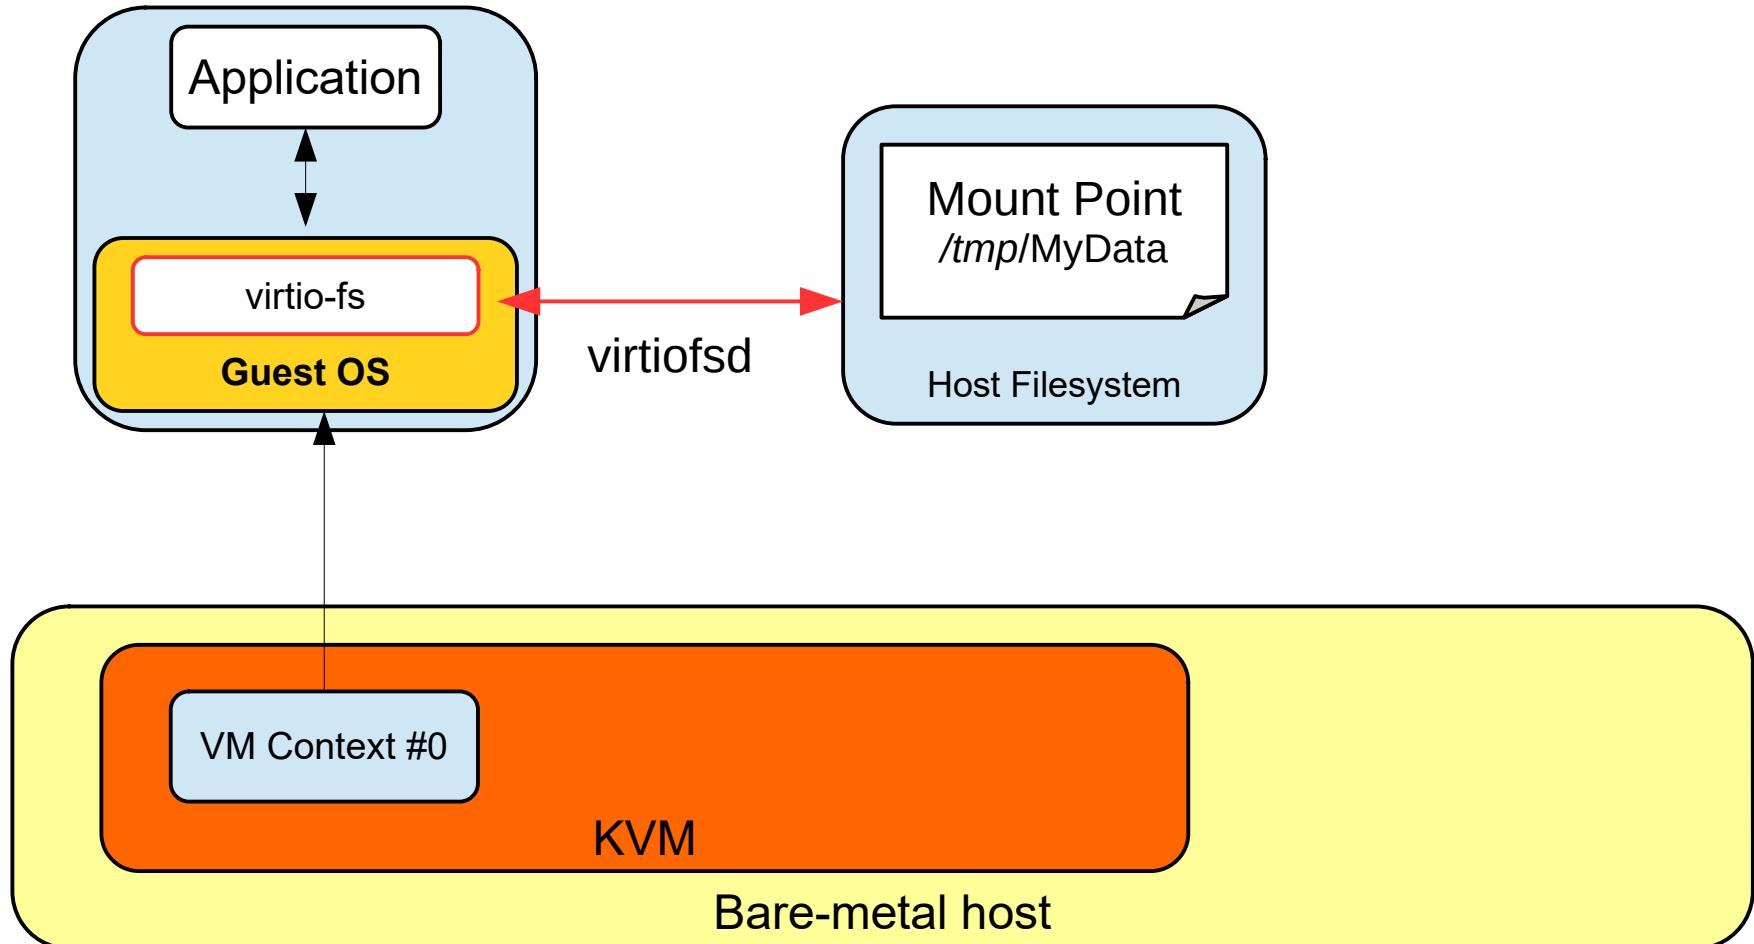
- Sharing of files between guests and host is hard
- Relying on a specific FS is not good for imm

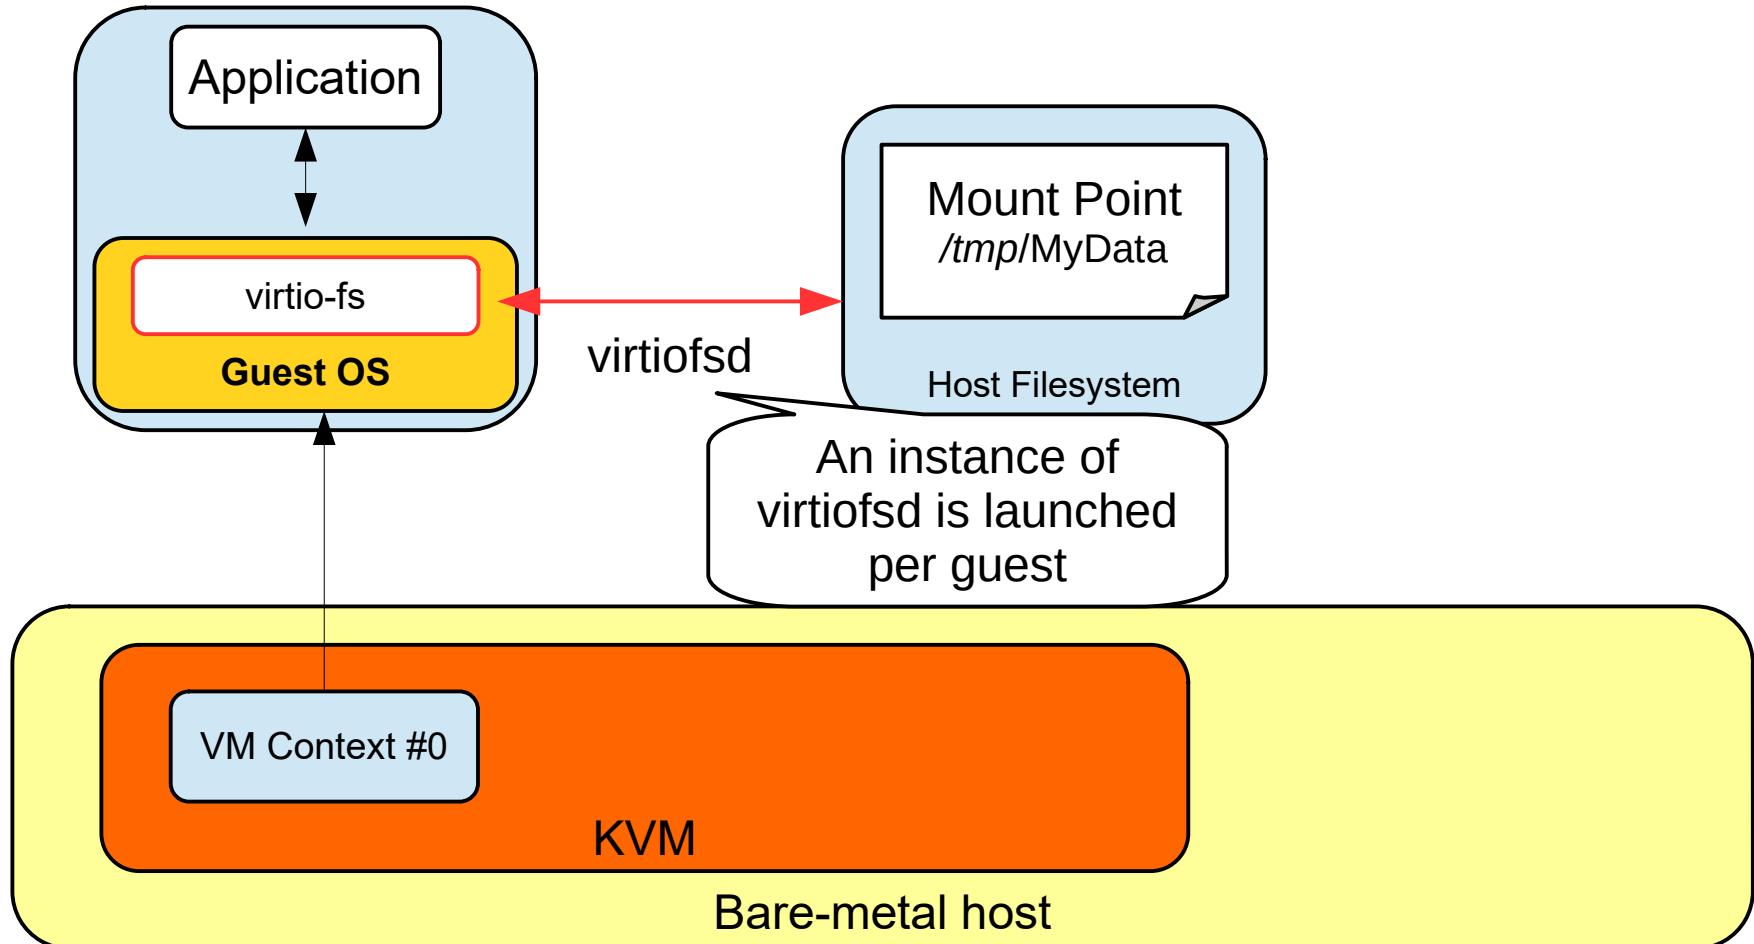


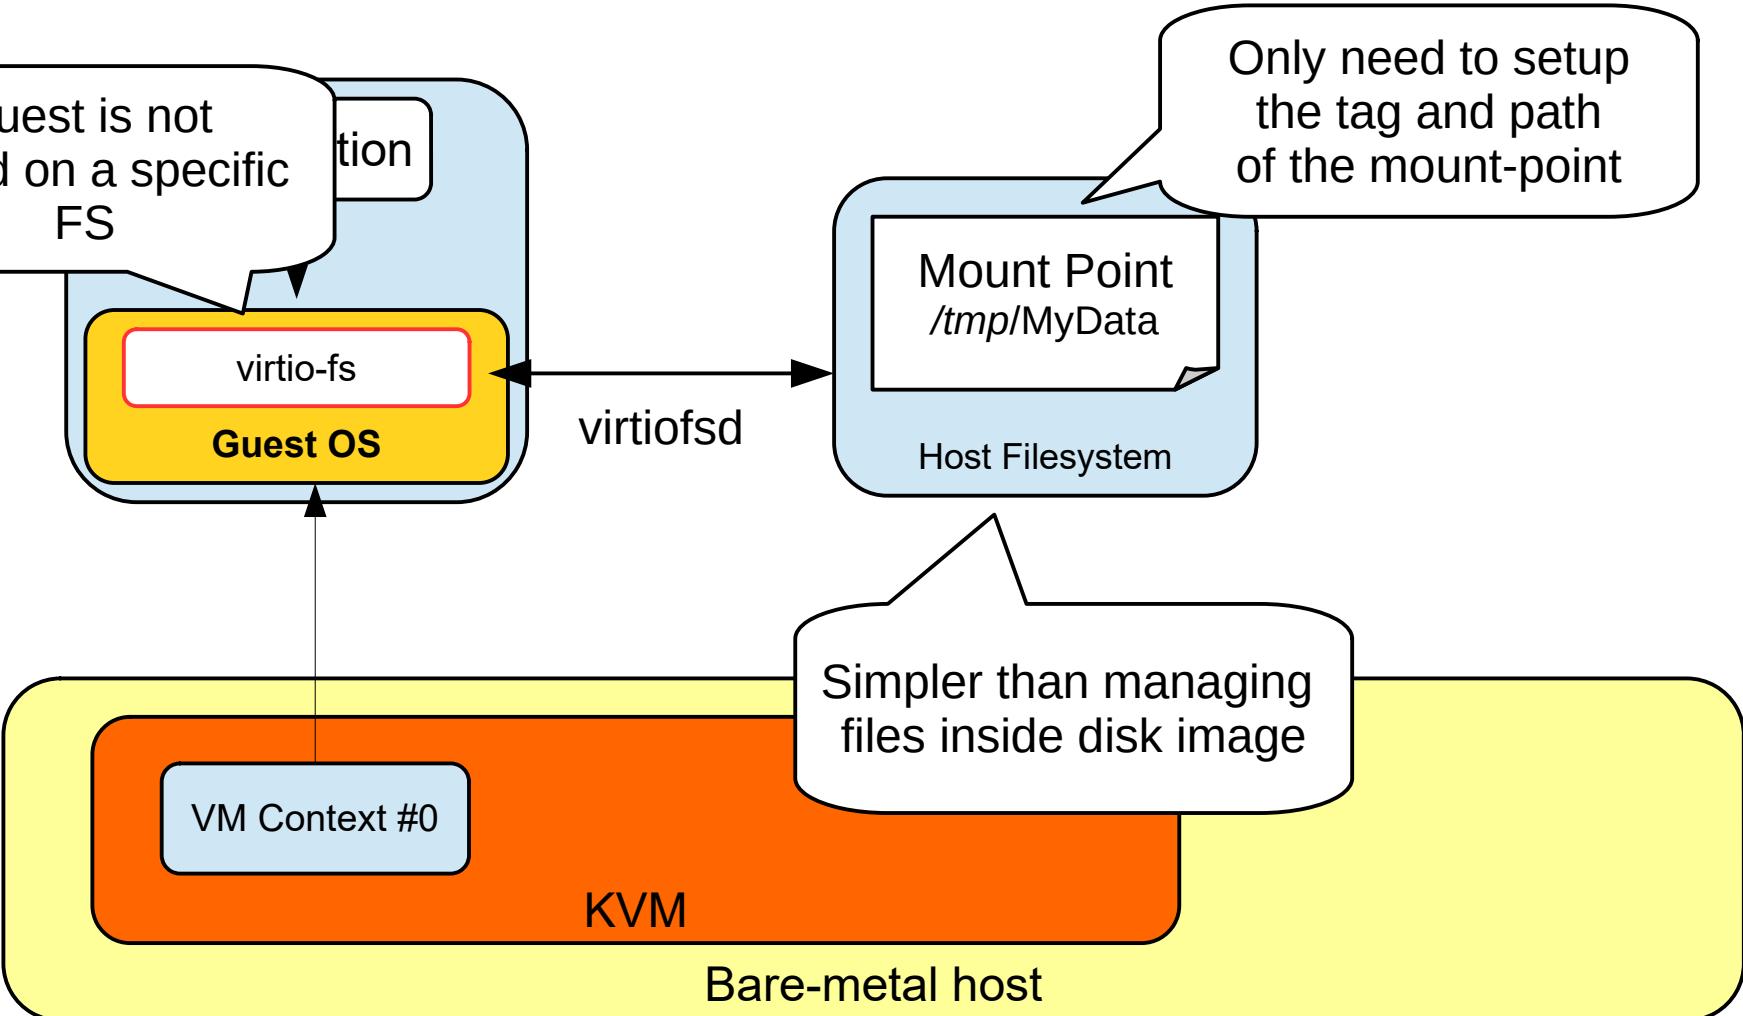
Can we do it better?!

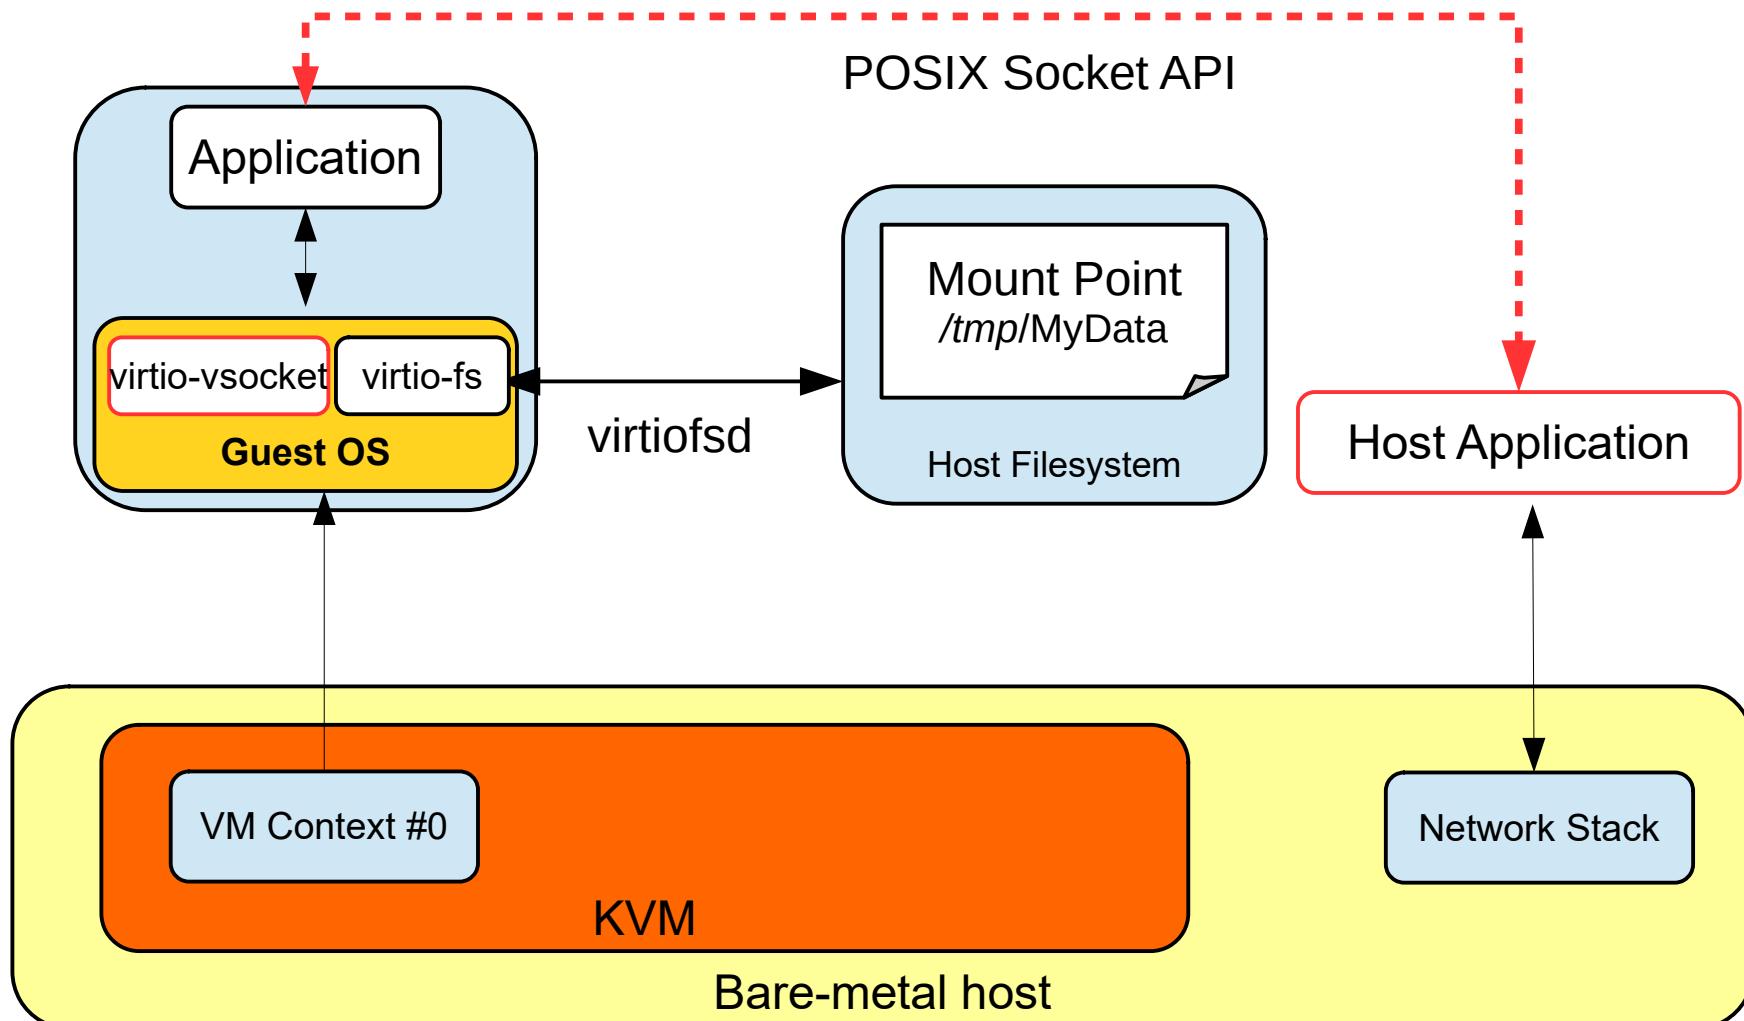
What are Virtio-fs and Virtio-vsocket?

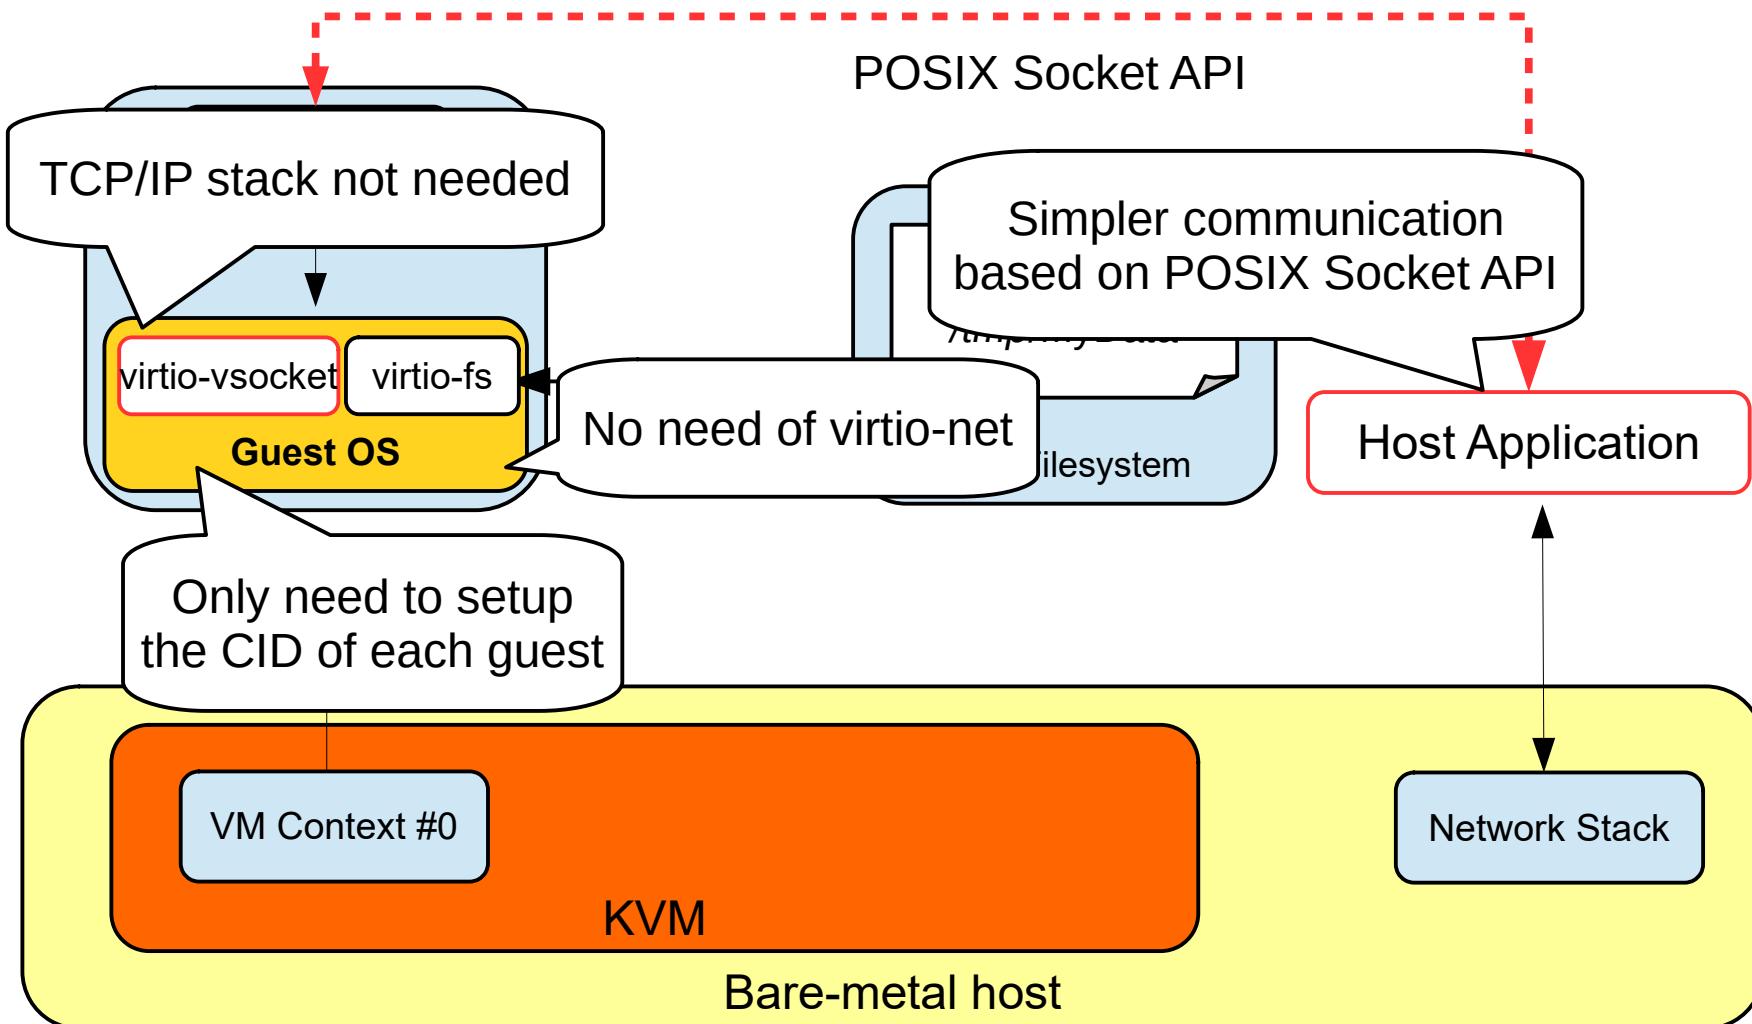
- Virtio devices
- Virtio-fs is a shared file system that allows VMs to access a directory tree on the host (more information at <https://virtio-fs.gitlab.io/>)
- Virtio-vsocket allows communication between guest and host by using POSIX Sockets API (more information at <https://wiki.qemu.org/Features/VirtioVsock>)











Results

- Reduce Networking source code
 - -1126 LoC → ~ 48% LoC less in Networking
- User appliance has no change
- Reduce FileSystem source code
 - -1046 LoC → ~ 54% LoC less in FileSytem
- Simplified appliance deployment
 - Launch virtiofsd and set the CID per instance
- Communication outside host is provided by an user application named VSocketProxy
 - This is ongoing work
- Currently deployed at <http://www.torokernel.io> and then go to “View on Toro”

Demo

Future Work

- Microvm
- Improve drivers, e.g., zero-copy in virtio-vsocket
- Formal specification of VirtIO
- Commit VSocketProxy

QA



Use Case: WebServer

