




# Toro, a Dedicated (Uni)kernel for Microservices


Dr. Matias E. Vara Larsen

# What are microservices?

1. Monolithic Application




2. Decomposed Application into Services



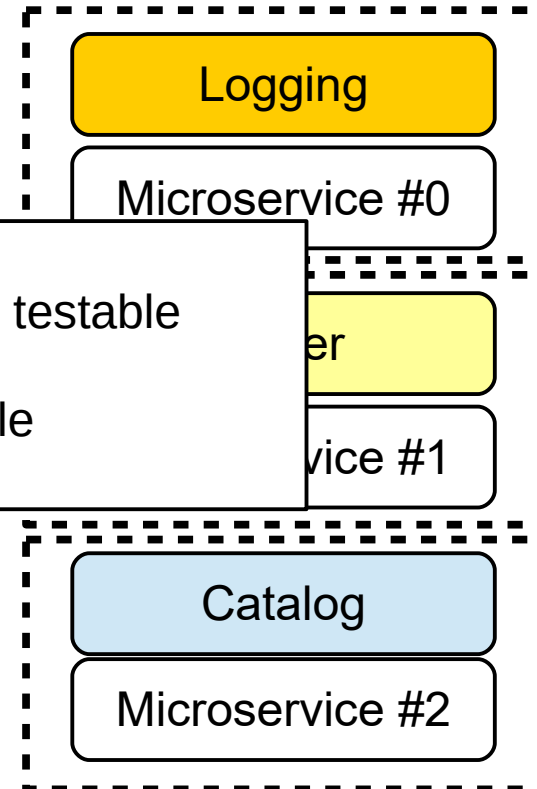
# What are microservices?

## 1. Monolithic Application

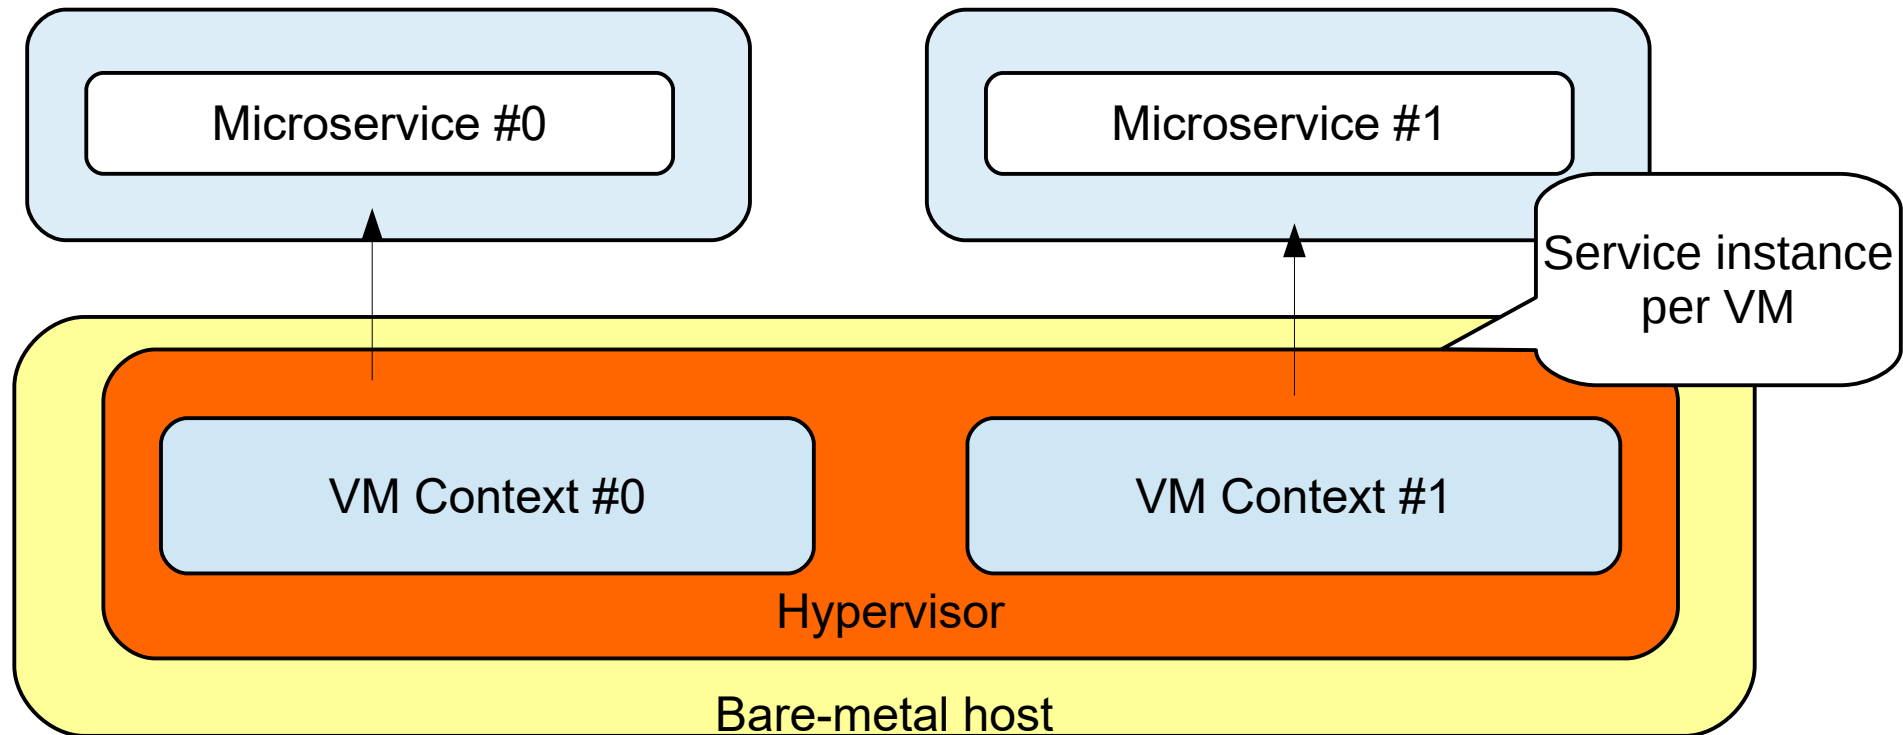


- Highly maintainable and testable
- Loosely coupled
- Independently deployable

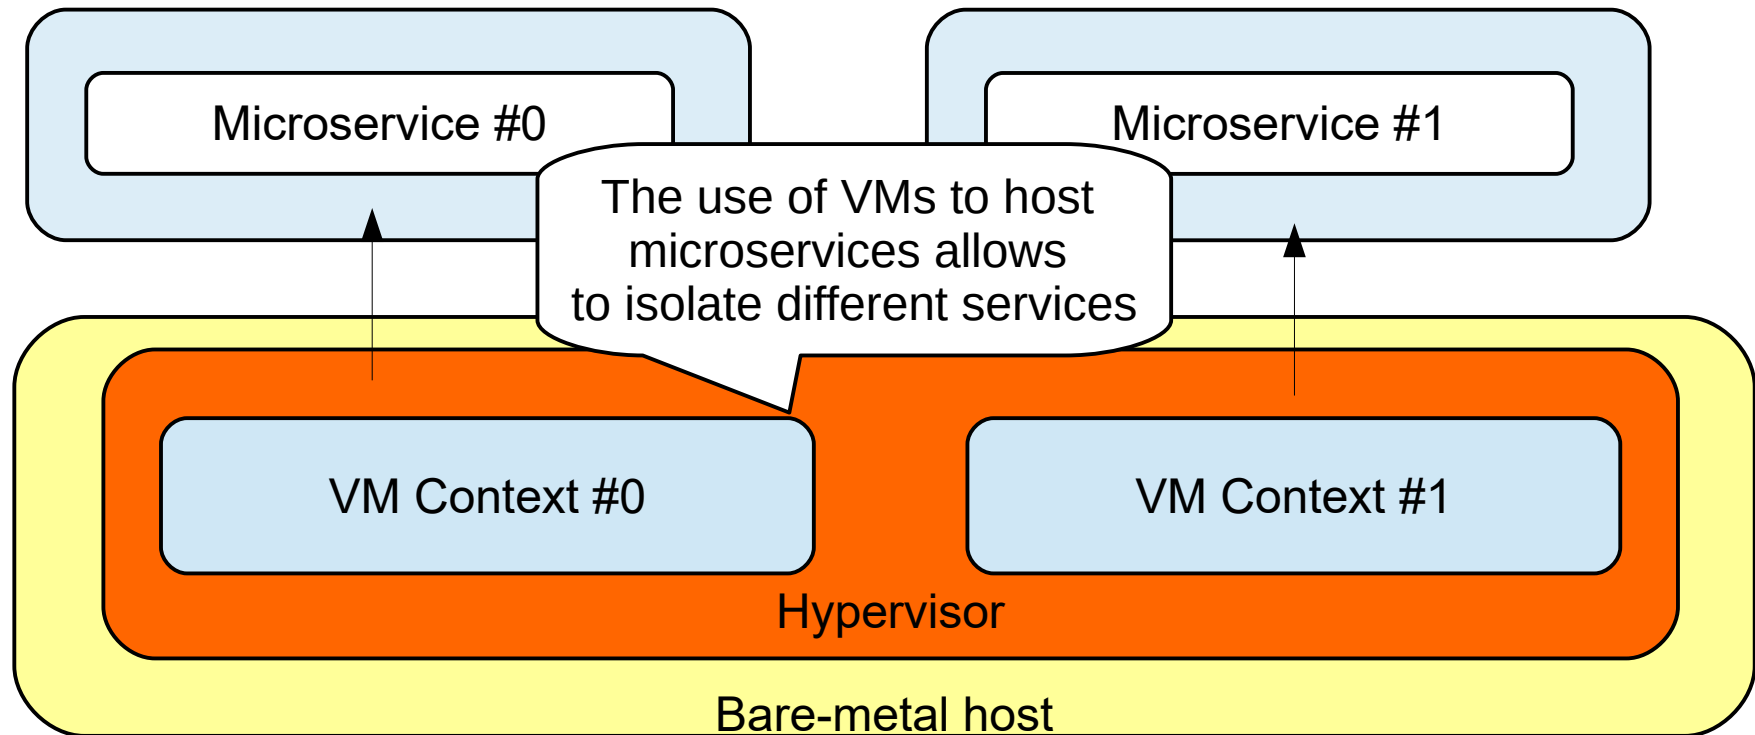
## 2. Decomposed Application into Services

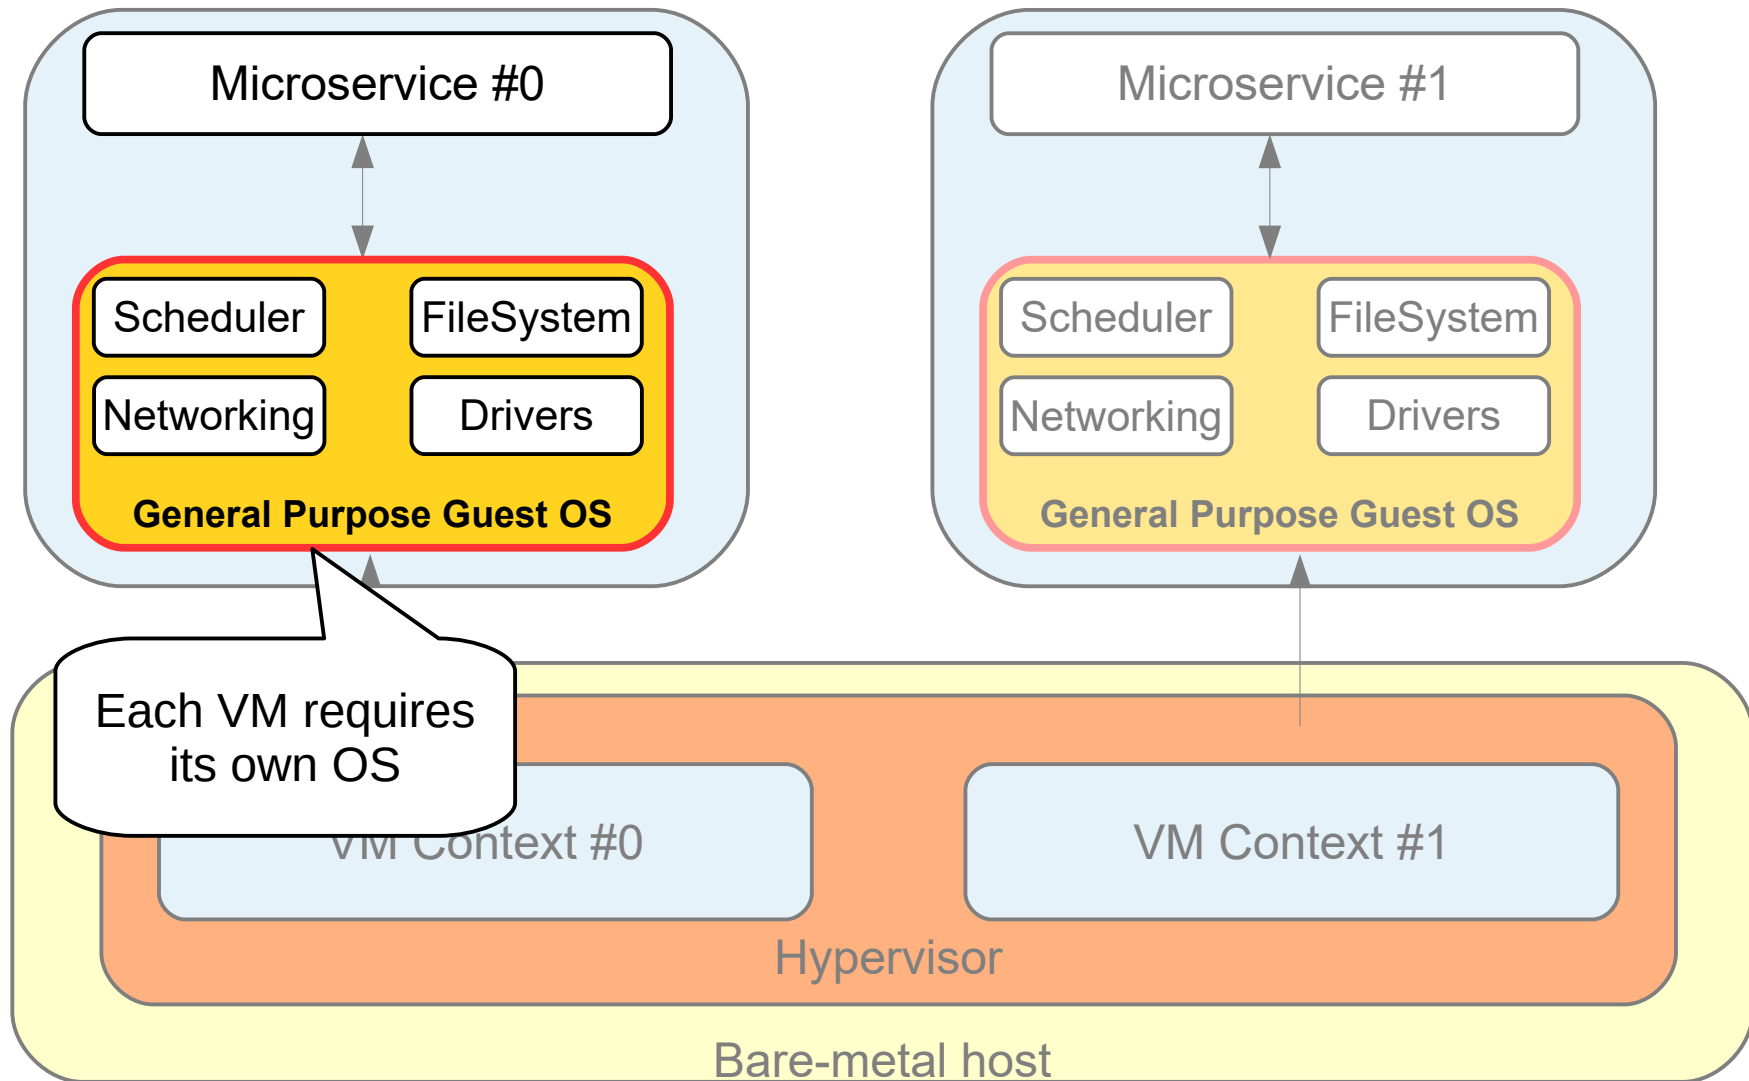


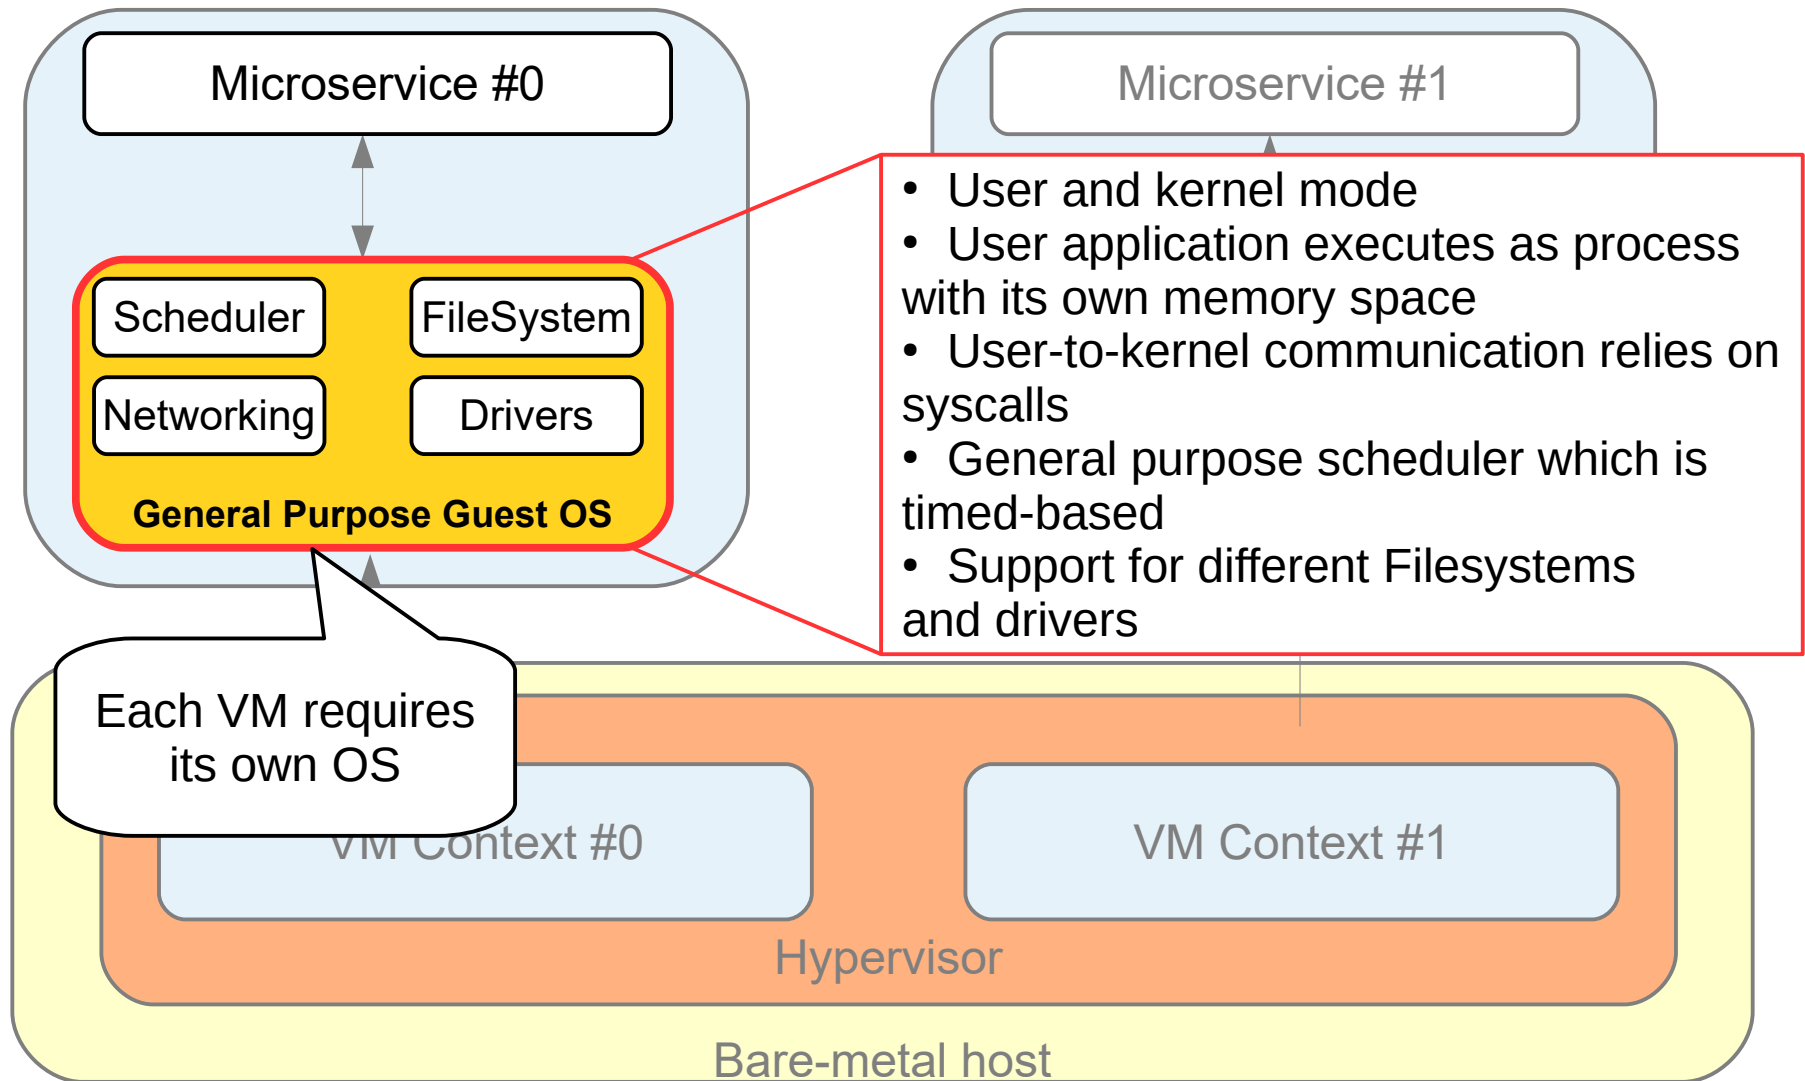
# How are microservices deployed?

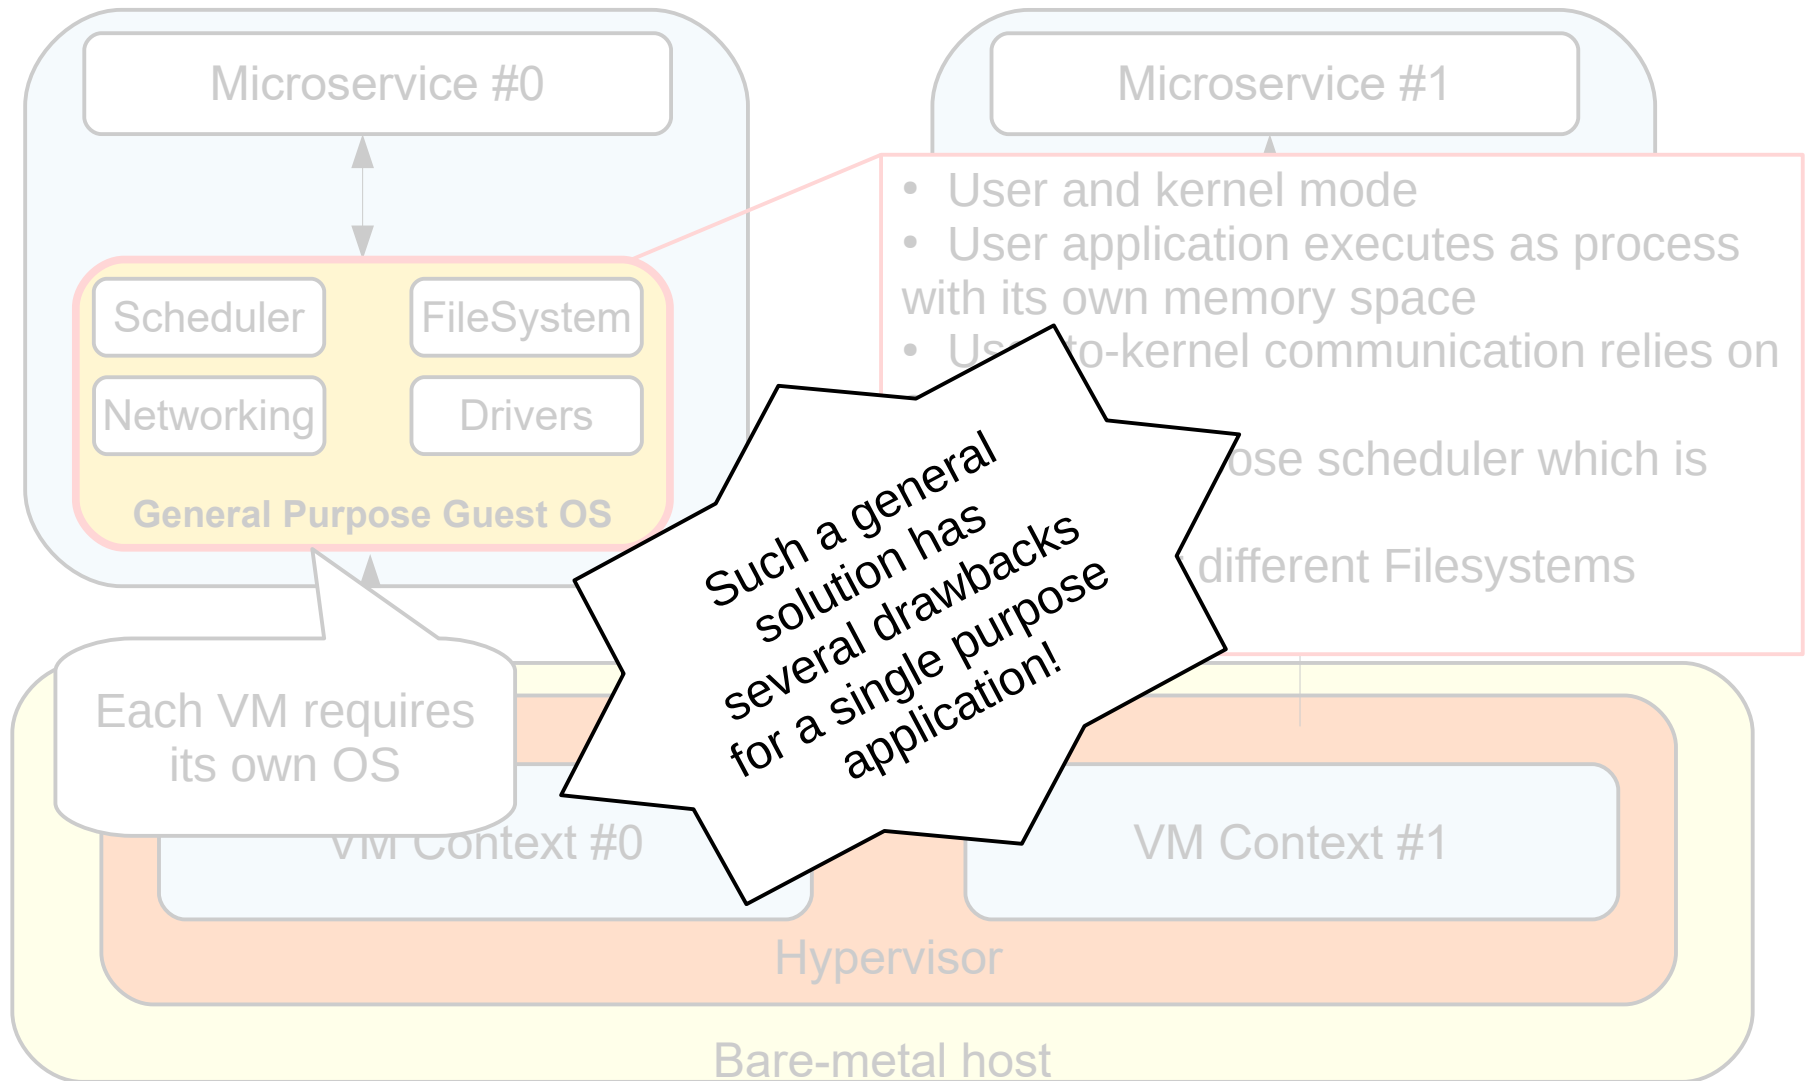


# How are microservices deployed?

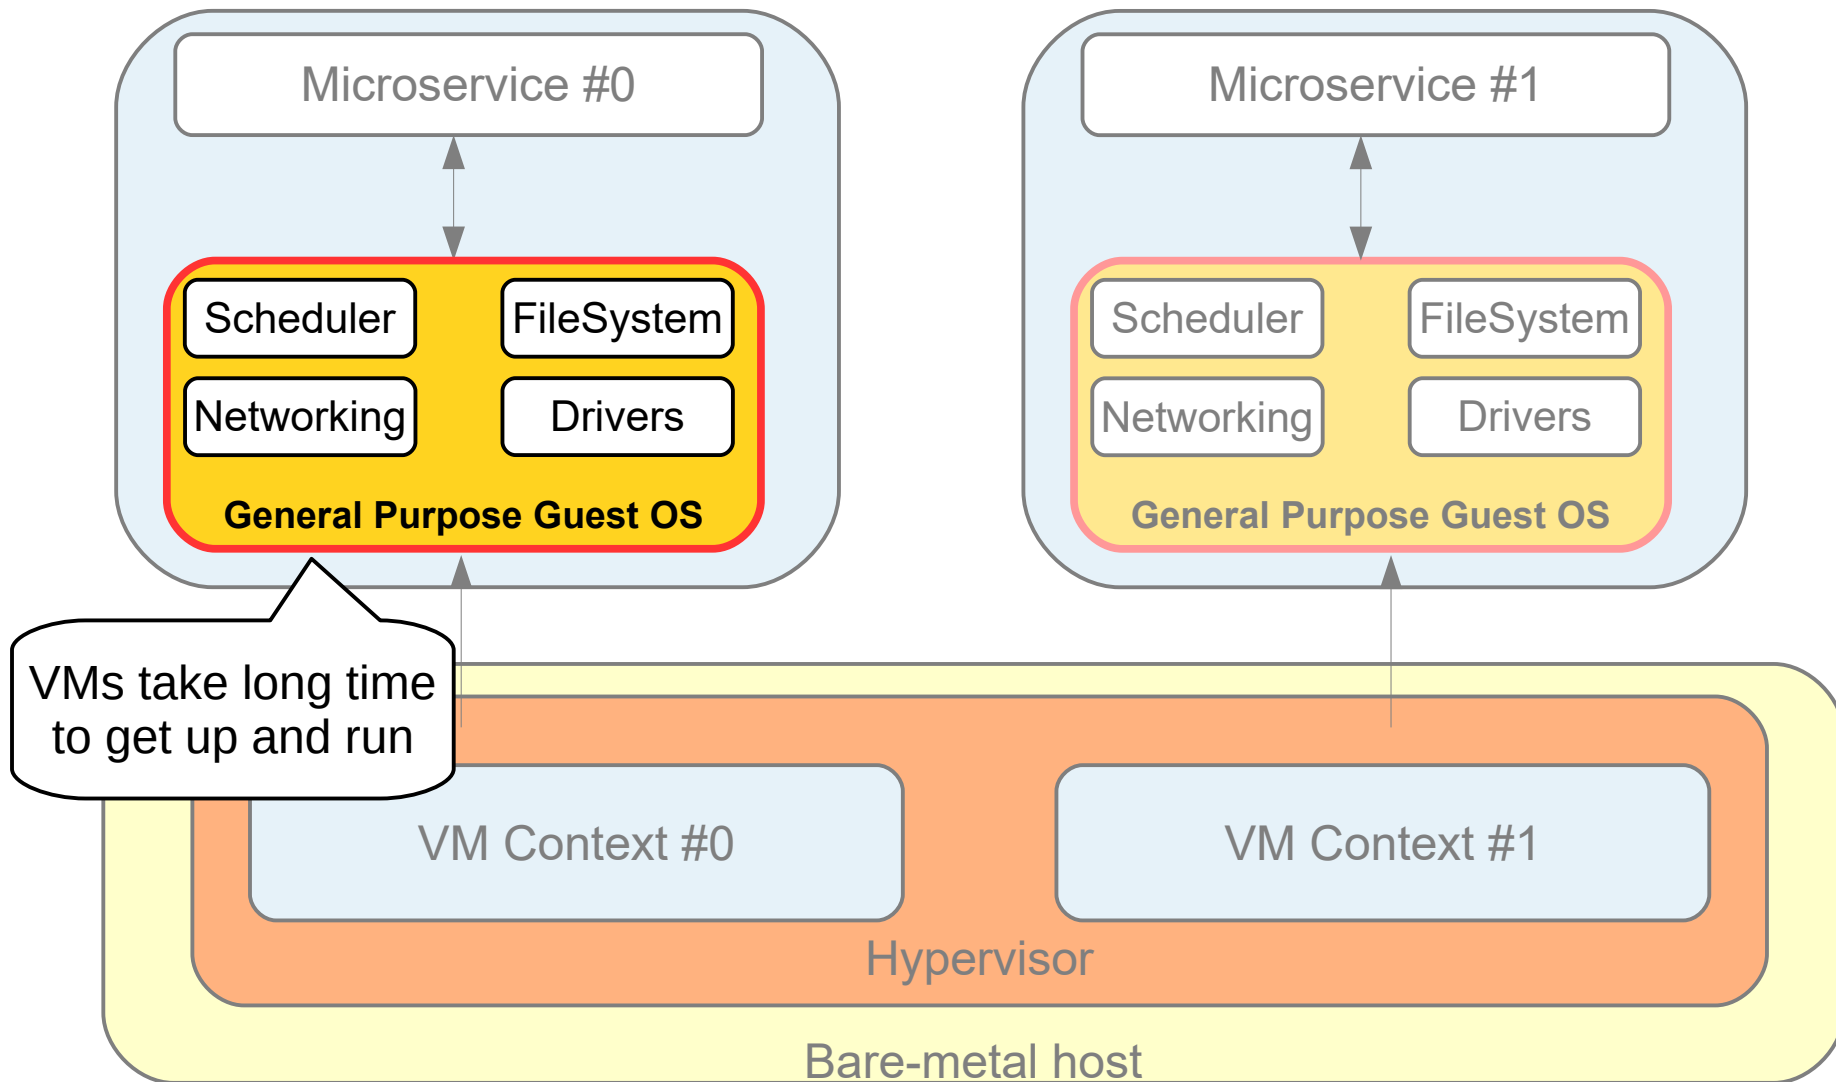


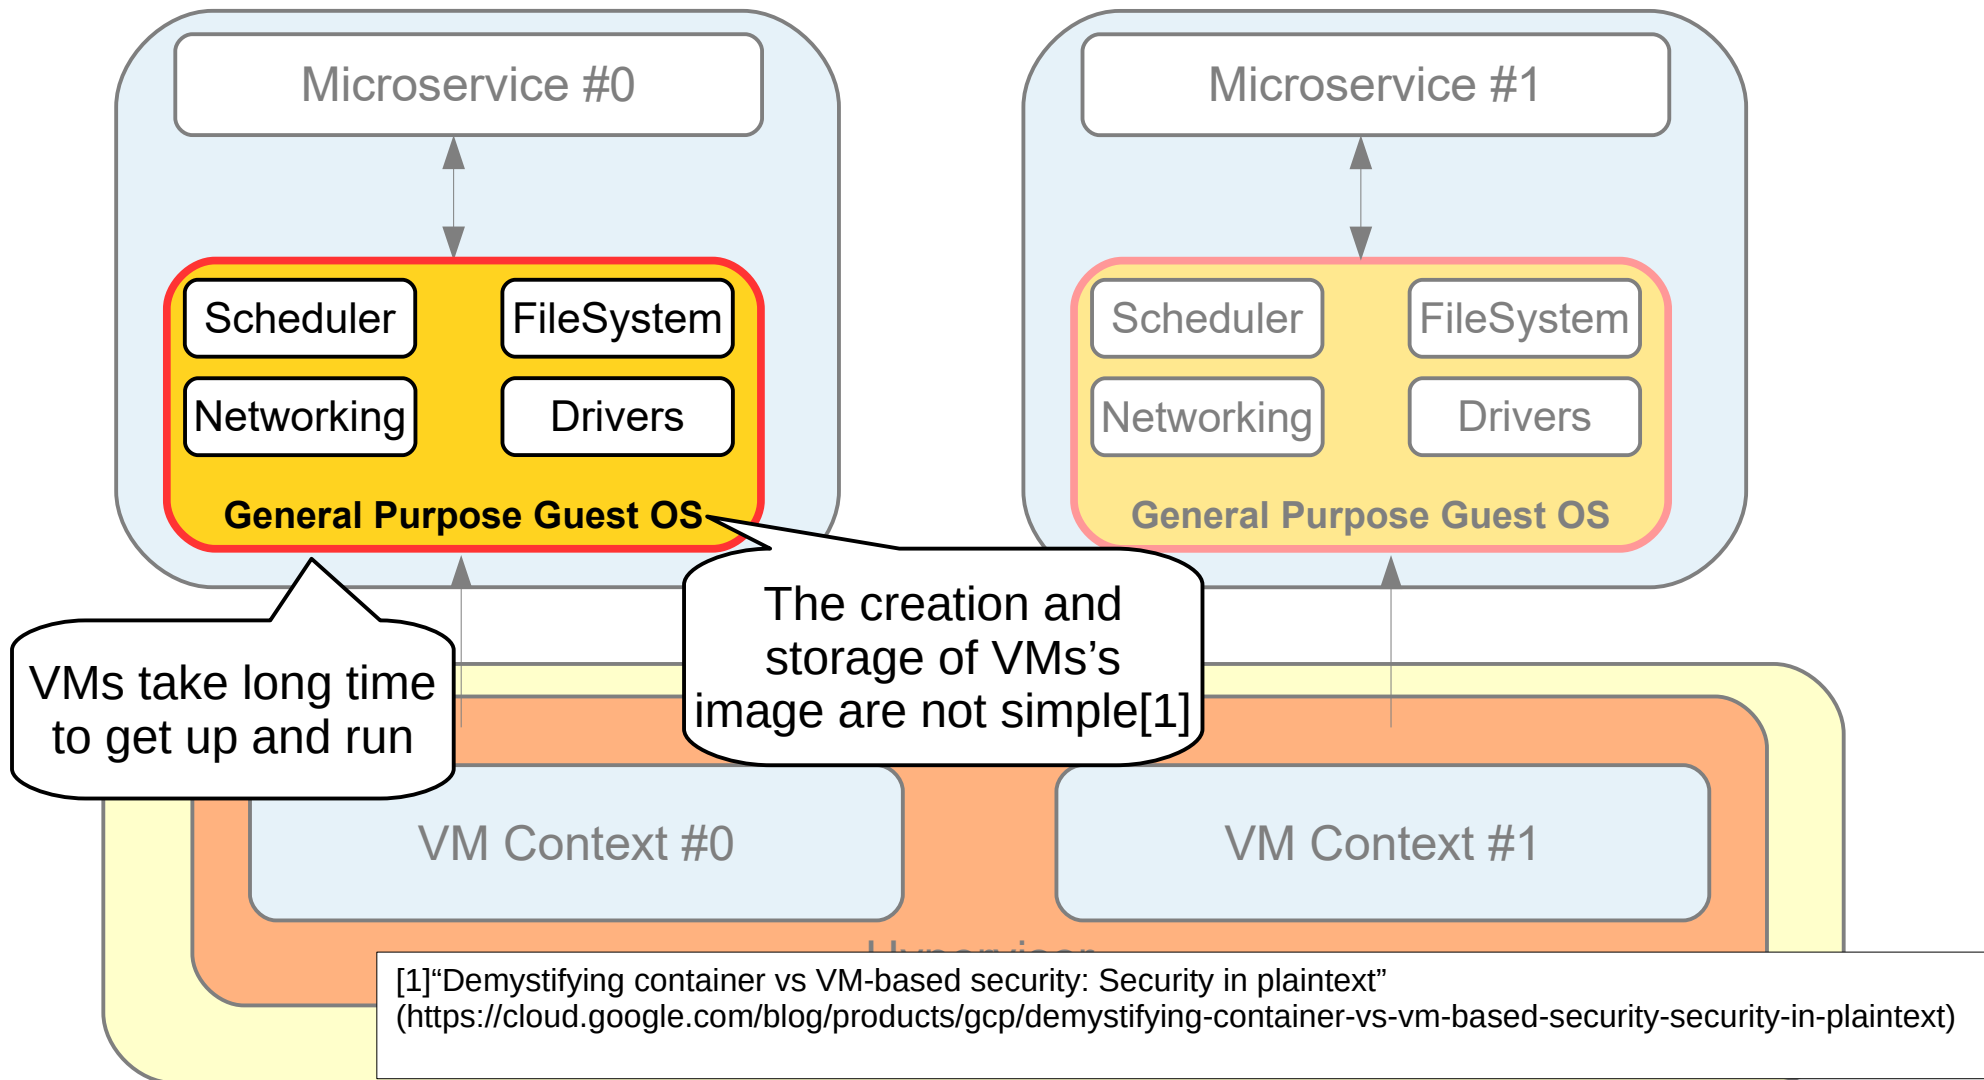


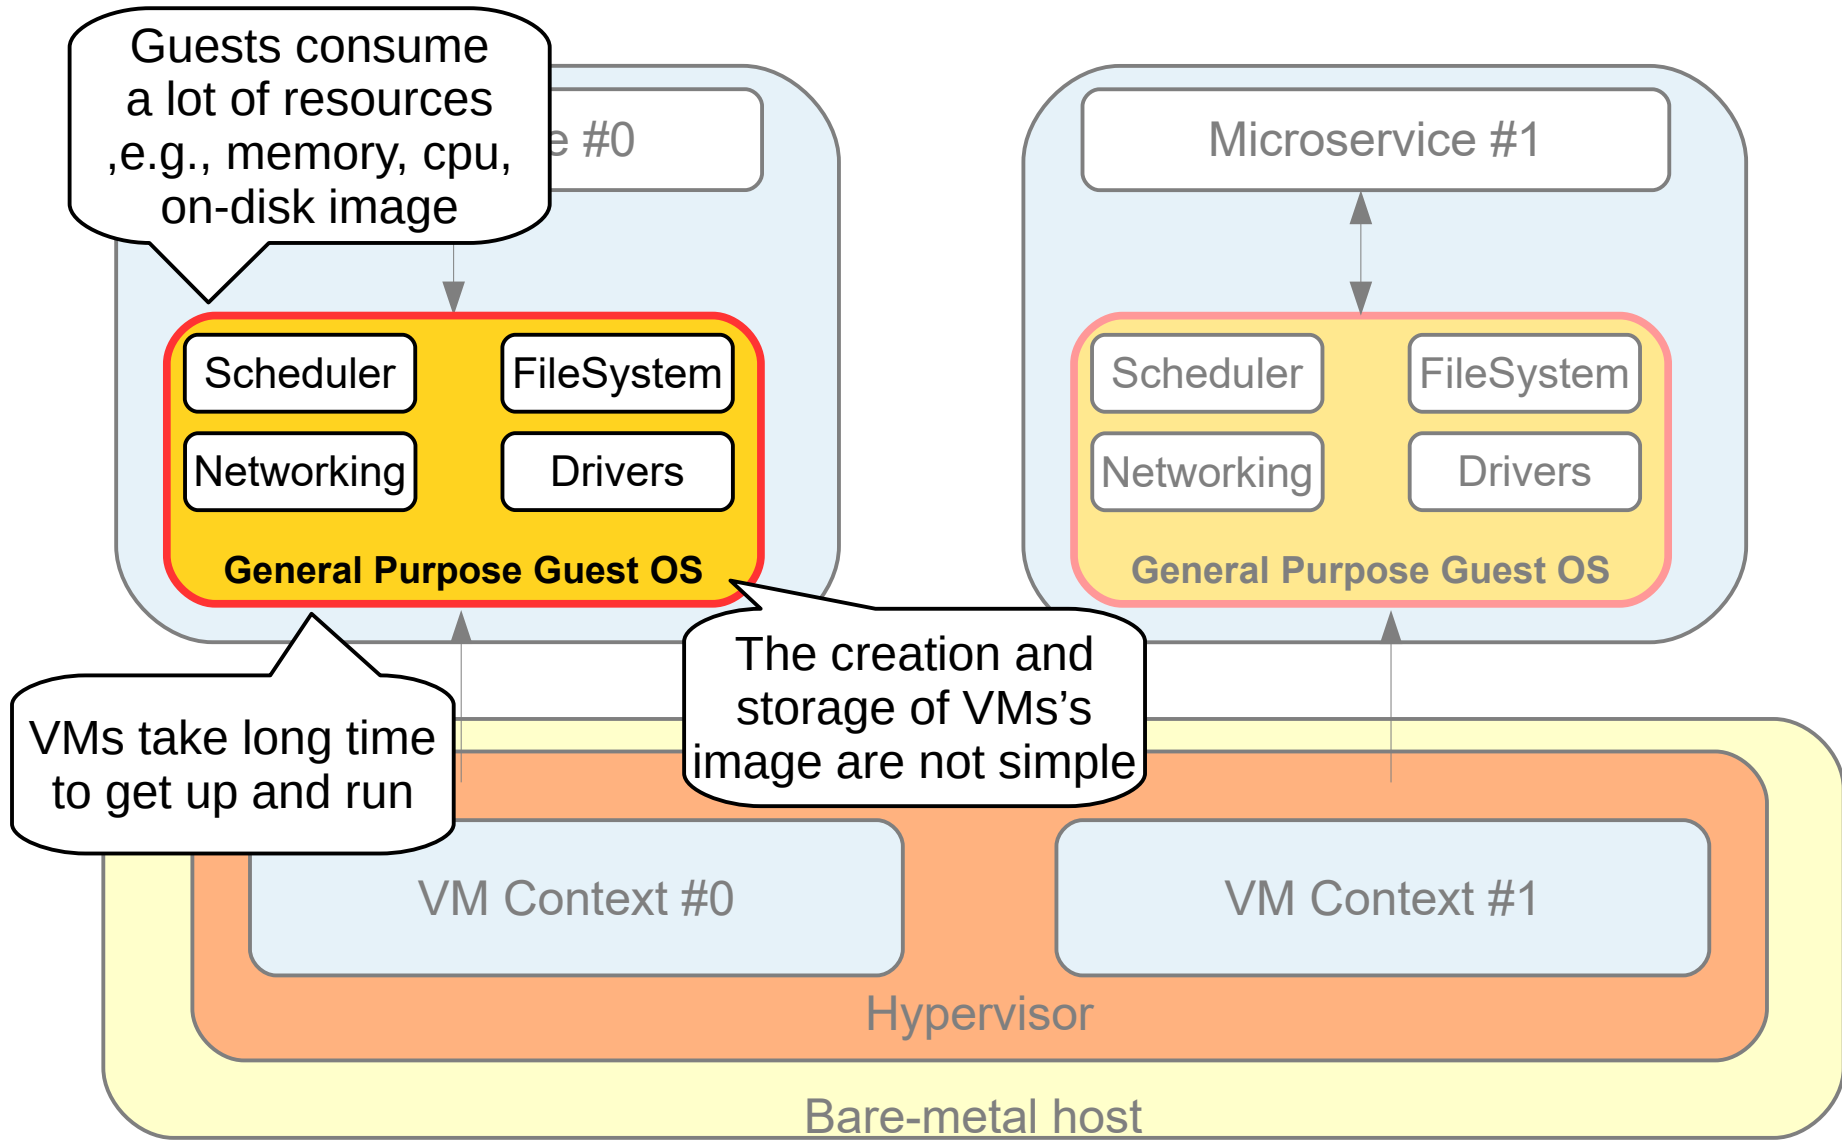


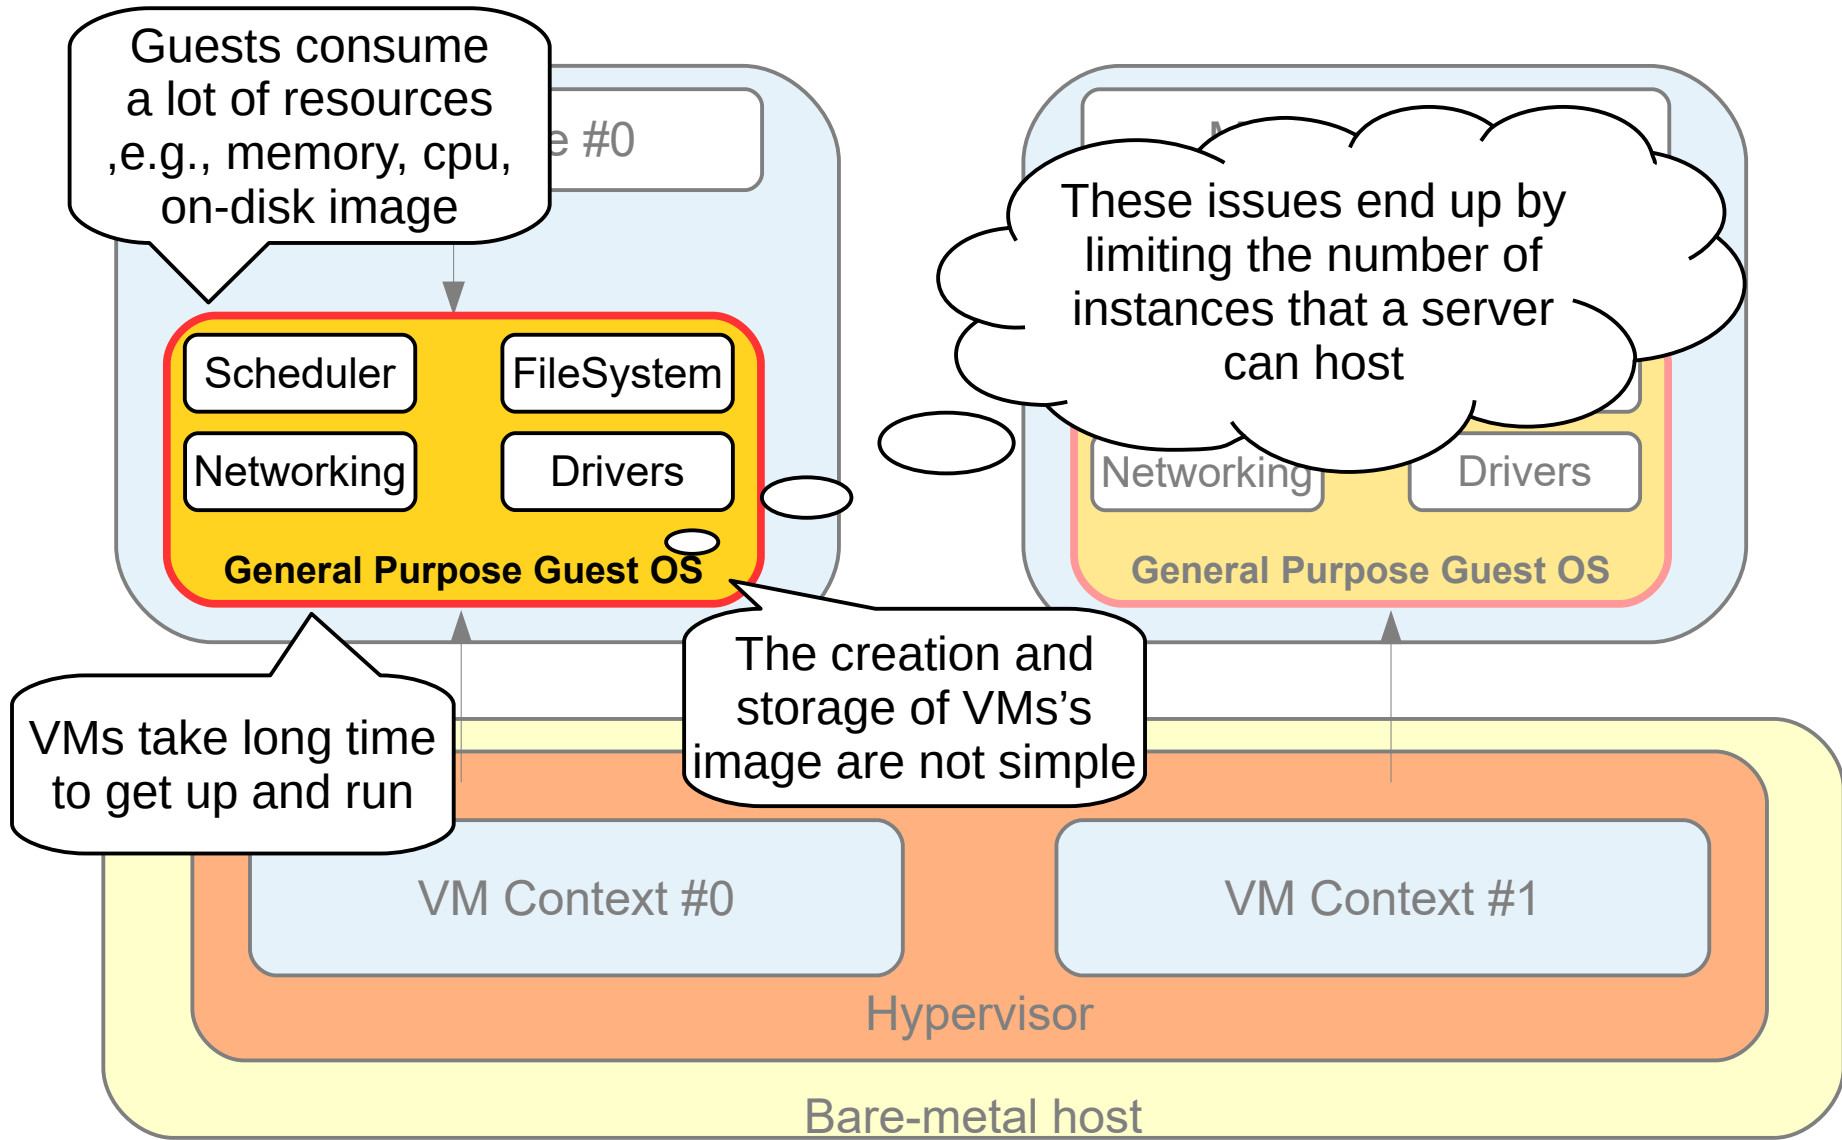


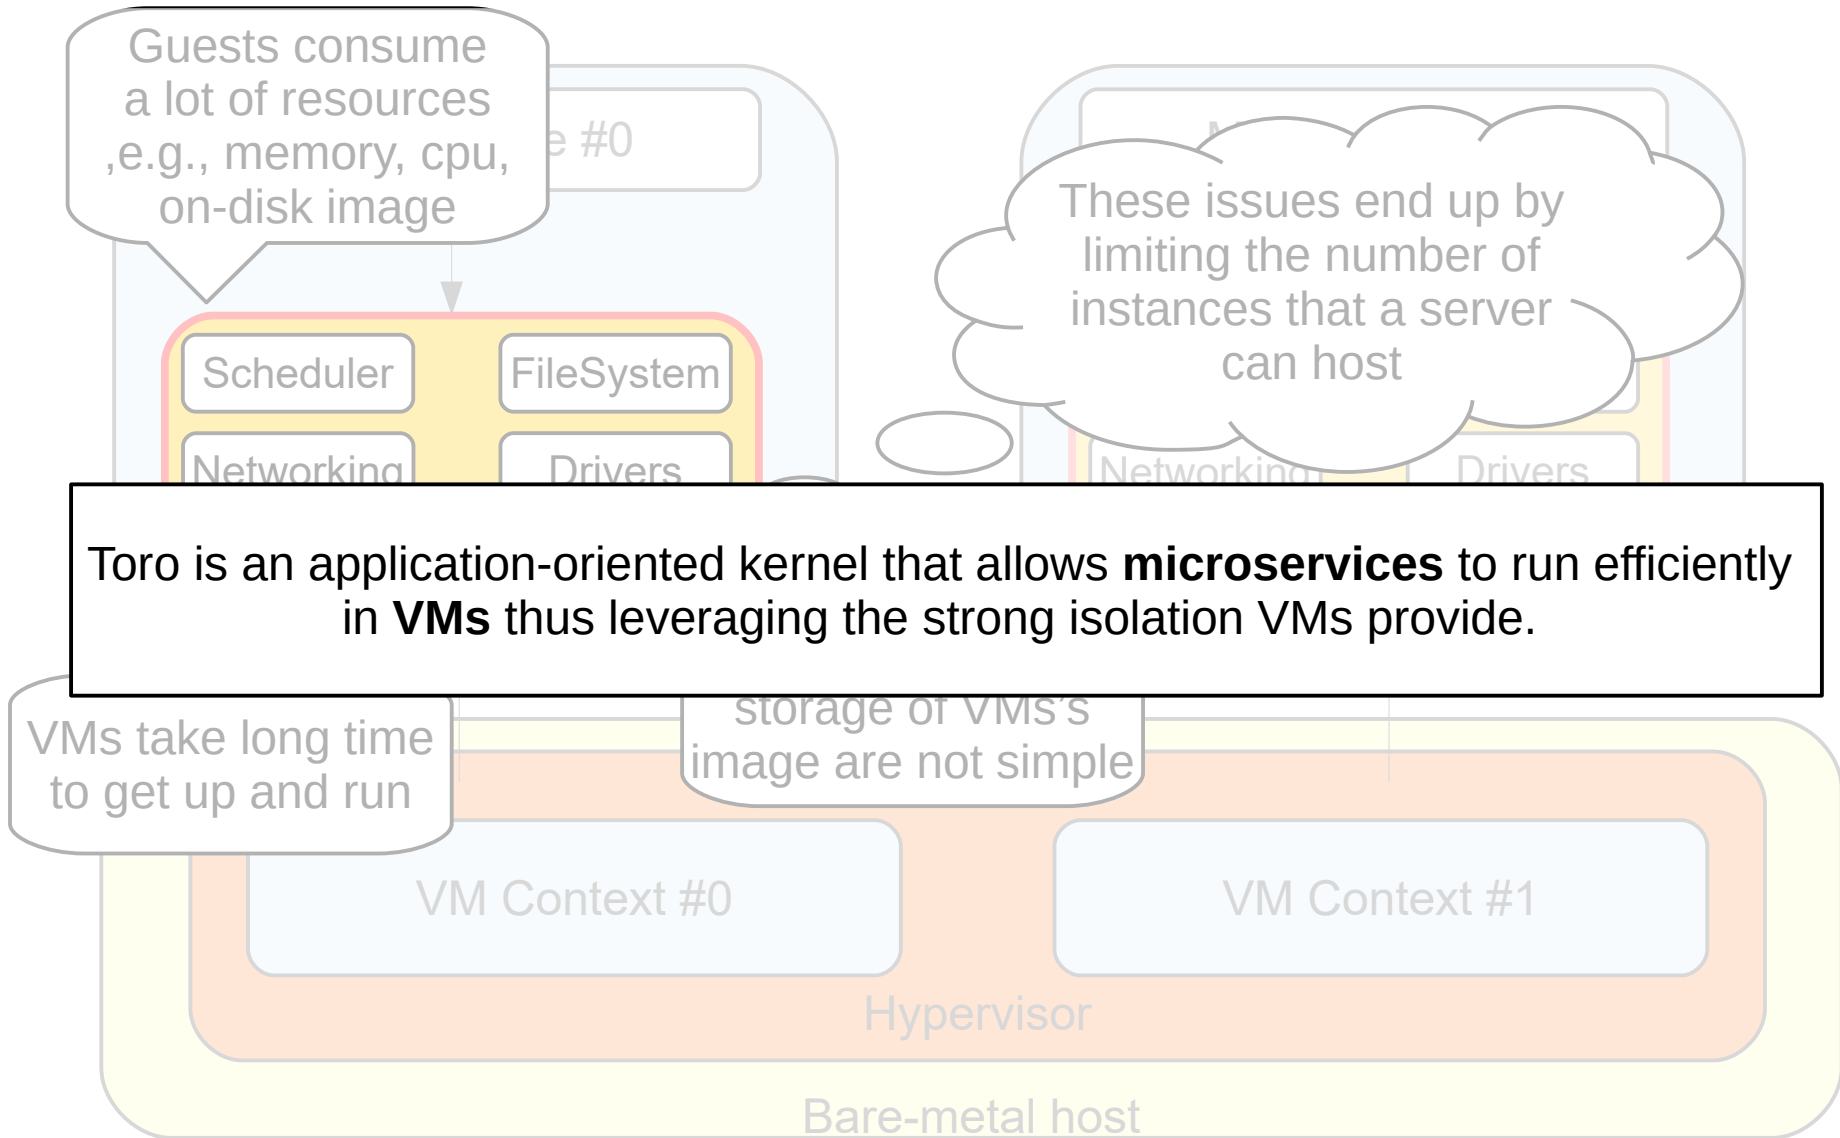













# Ingredients for ToroKernel

- Application-oriented Kernel
- Simple Scheduler
- Dedicated Resources
- Microservice-oriented Networking

# Application-oriented Kernel

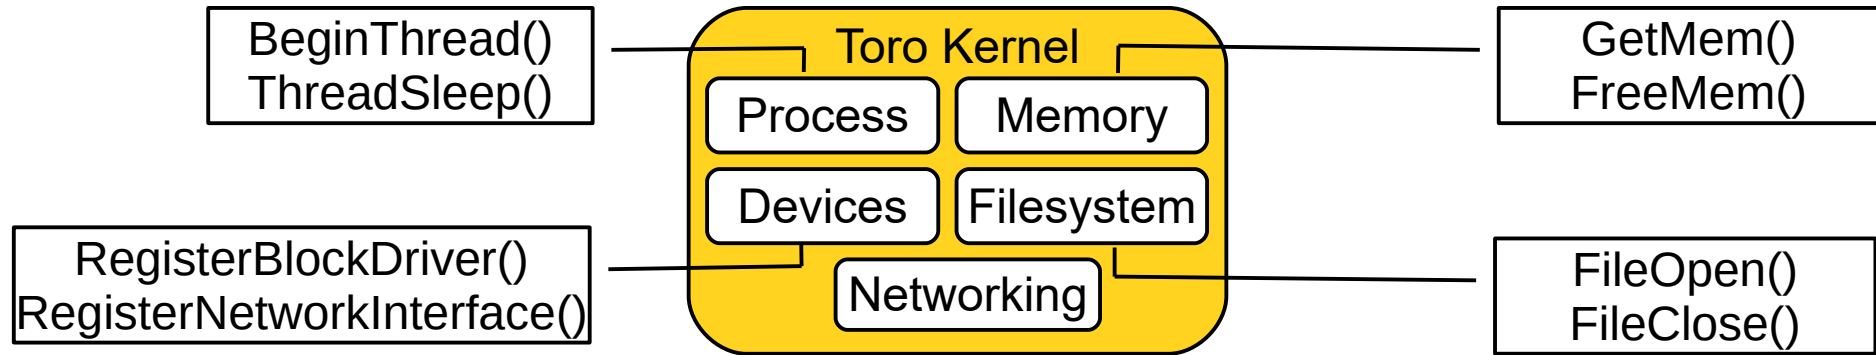


Toro is an embedded kernel including five units:

- Process
- Memory
- Filesystem
- Networking
- Devices, e.g., Block Device, Network Device

Each unit provides minimalist APIs accessible from the embedded application

# Application-oriented Kernel




Toro is an embedded kernel including five units:

- Process
- Memory
- Filesystem
- Networking
- Devices, e.g., Block Device, Network Device

Each unit provides minimalist APIs accessible from the embedded application




# Application-oriented Kernel



- User application and kernel units are compiled in a single binary
- The application includes only the component required

# Application-oriented Kernel



Uses

Microservice

- User application and kernel units are compiled in a single binary

- The com

```
program HelloWorldMicroservice;
```

```
uses
```

```
Memory,  
Filesystem,  
Ext3,  
E1000;
```

```
Begin
```

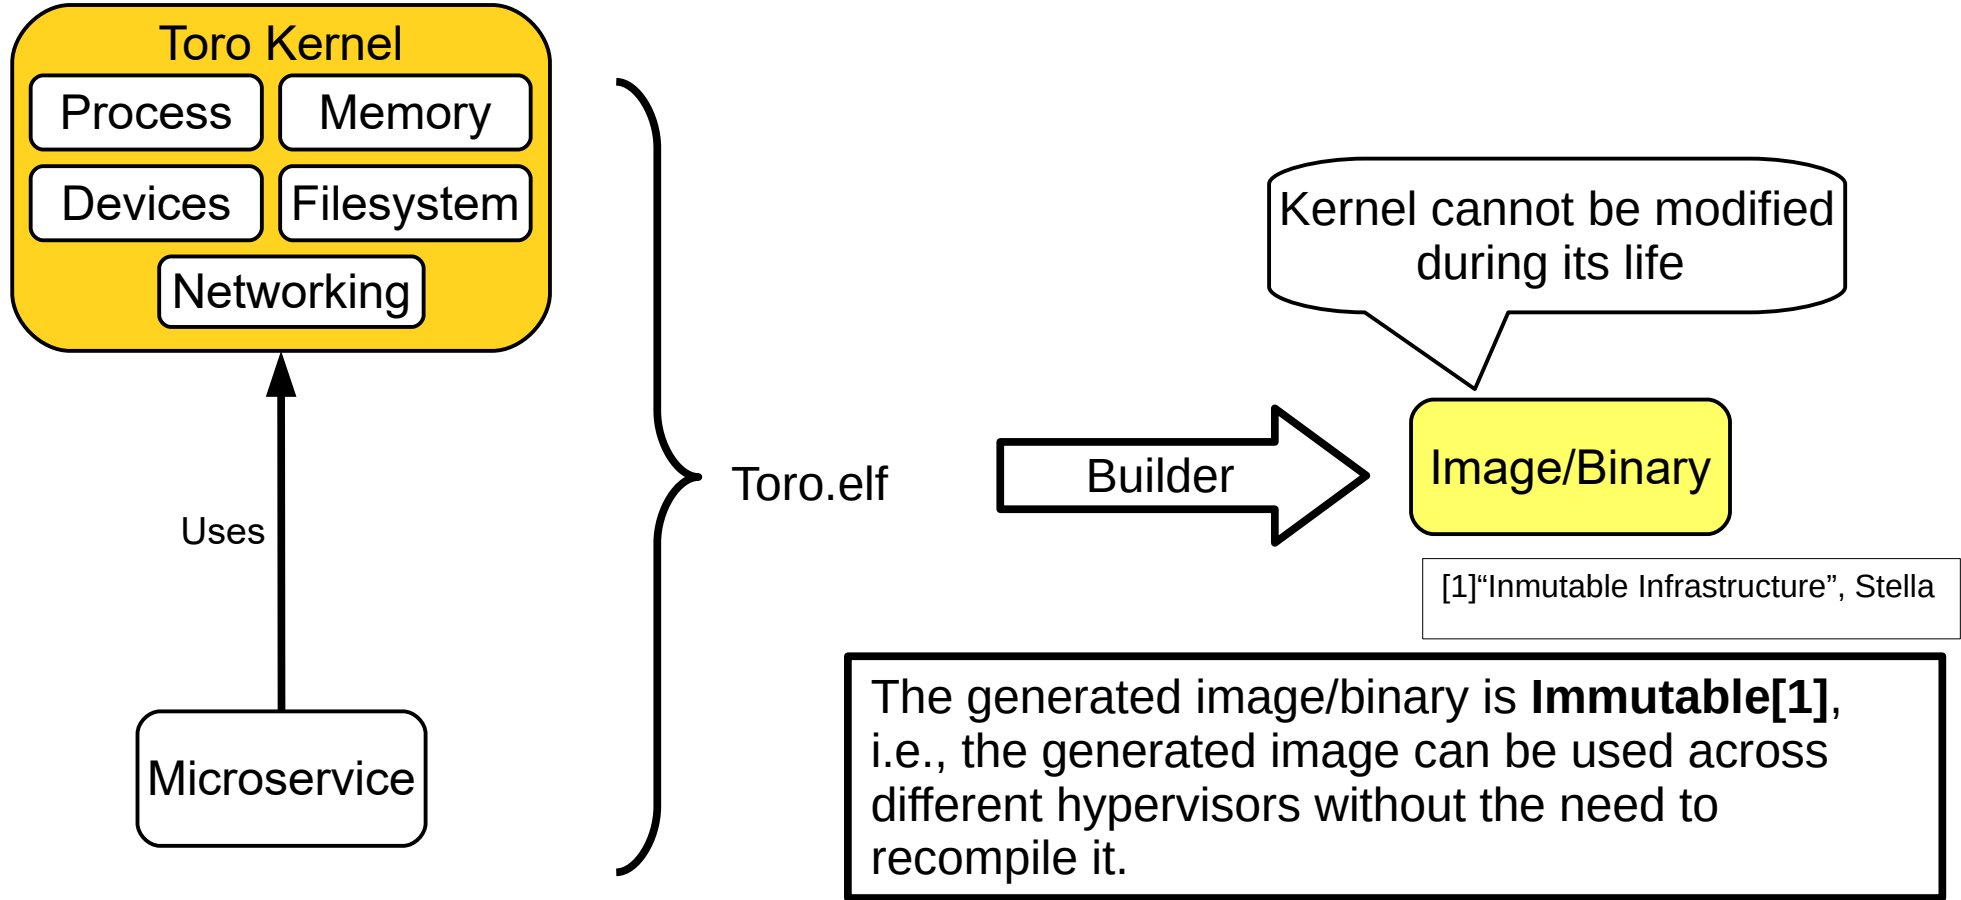
```
//
```

```
// Your Code
```

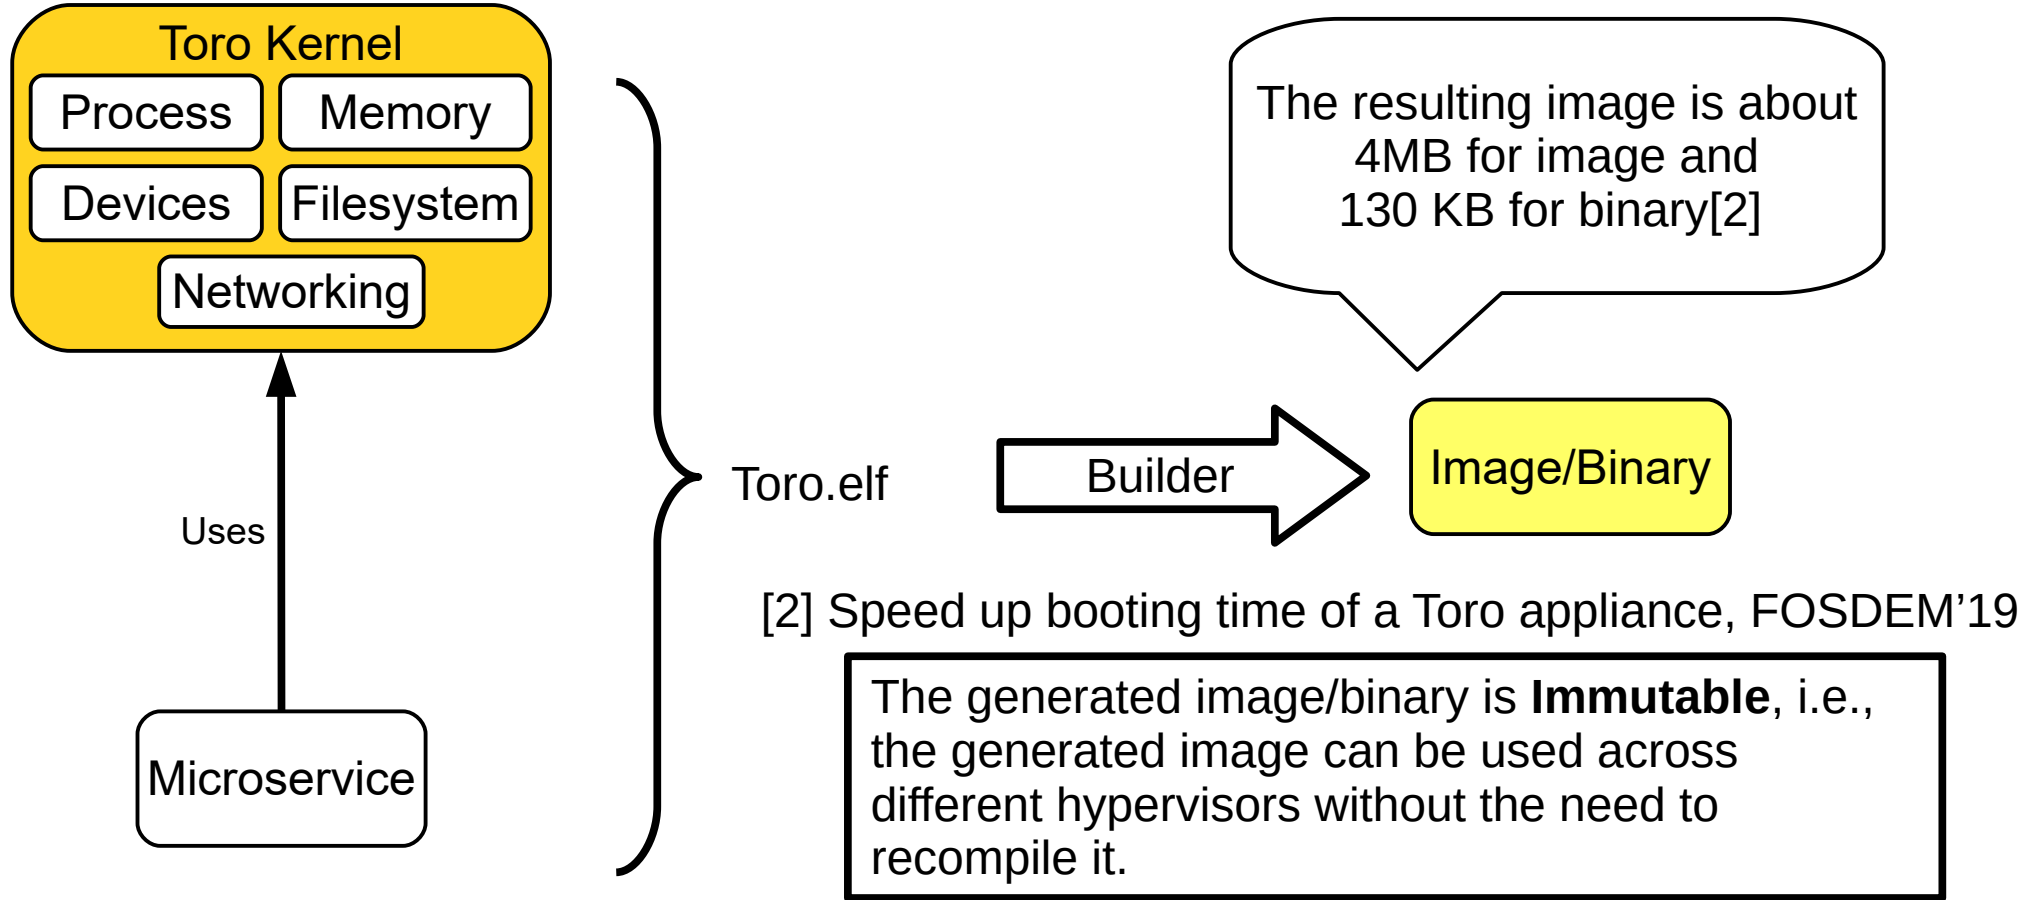
```
//
```

```
End.
```

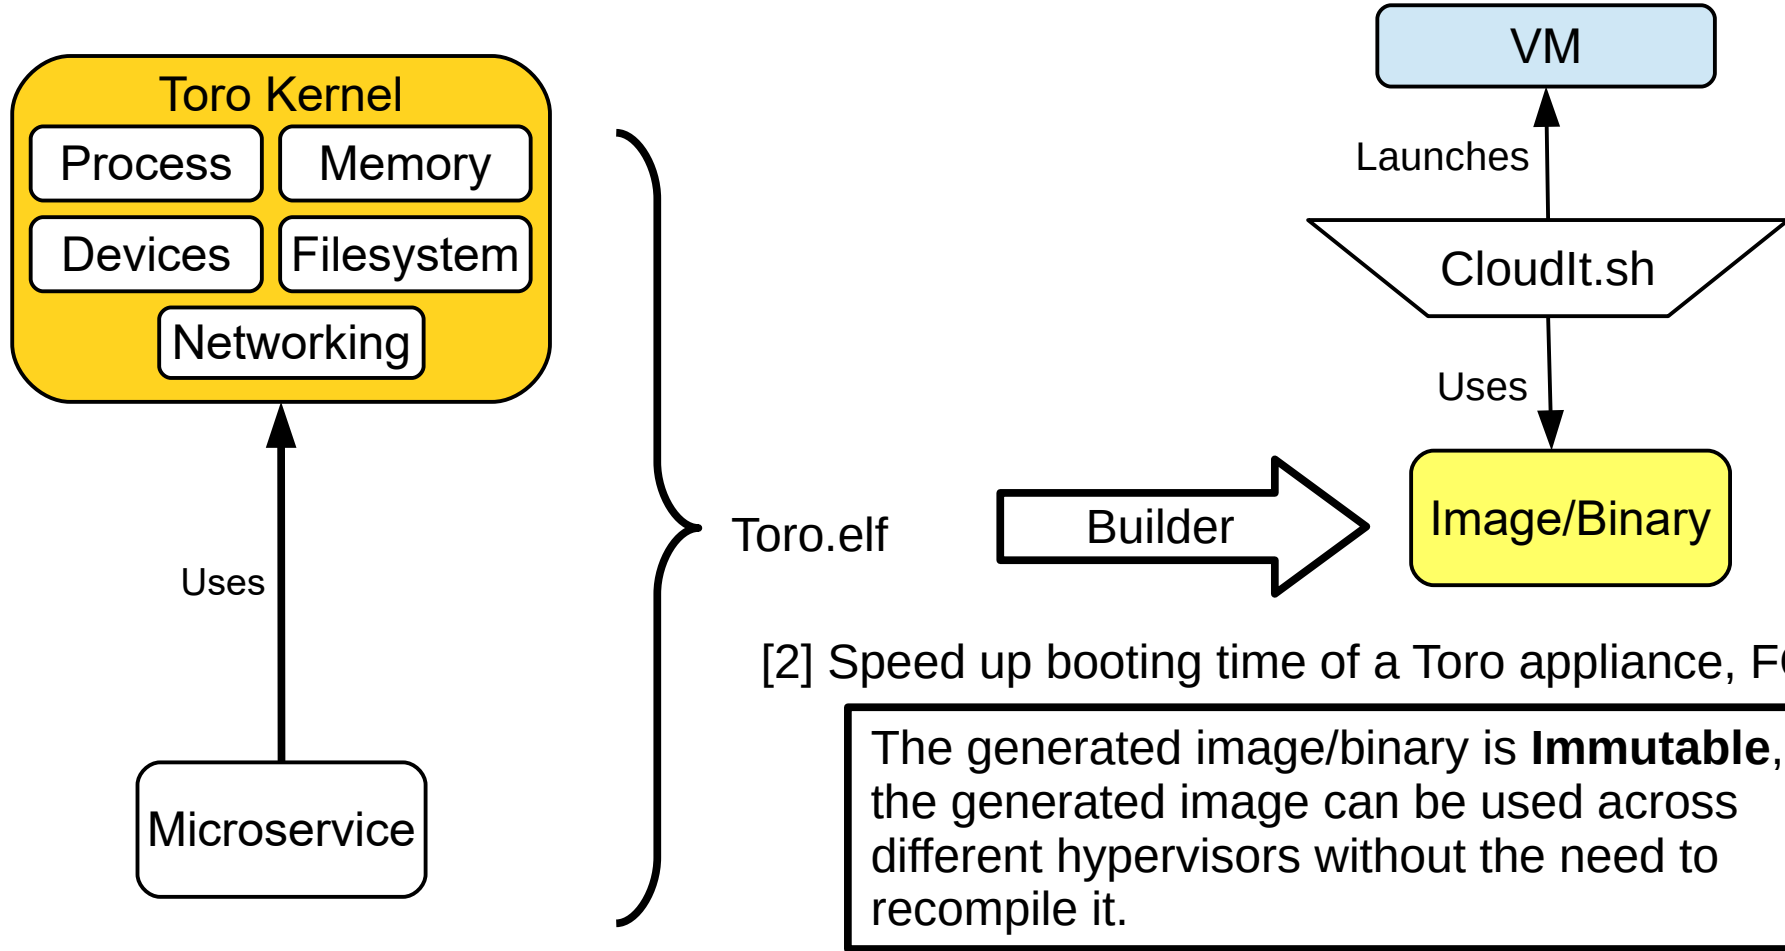
# Application-oriented Kernel



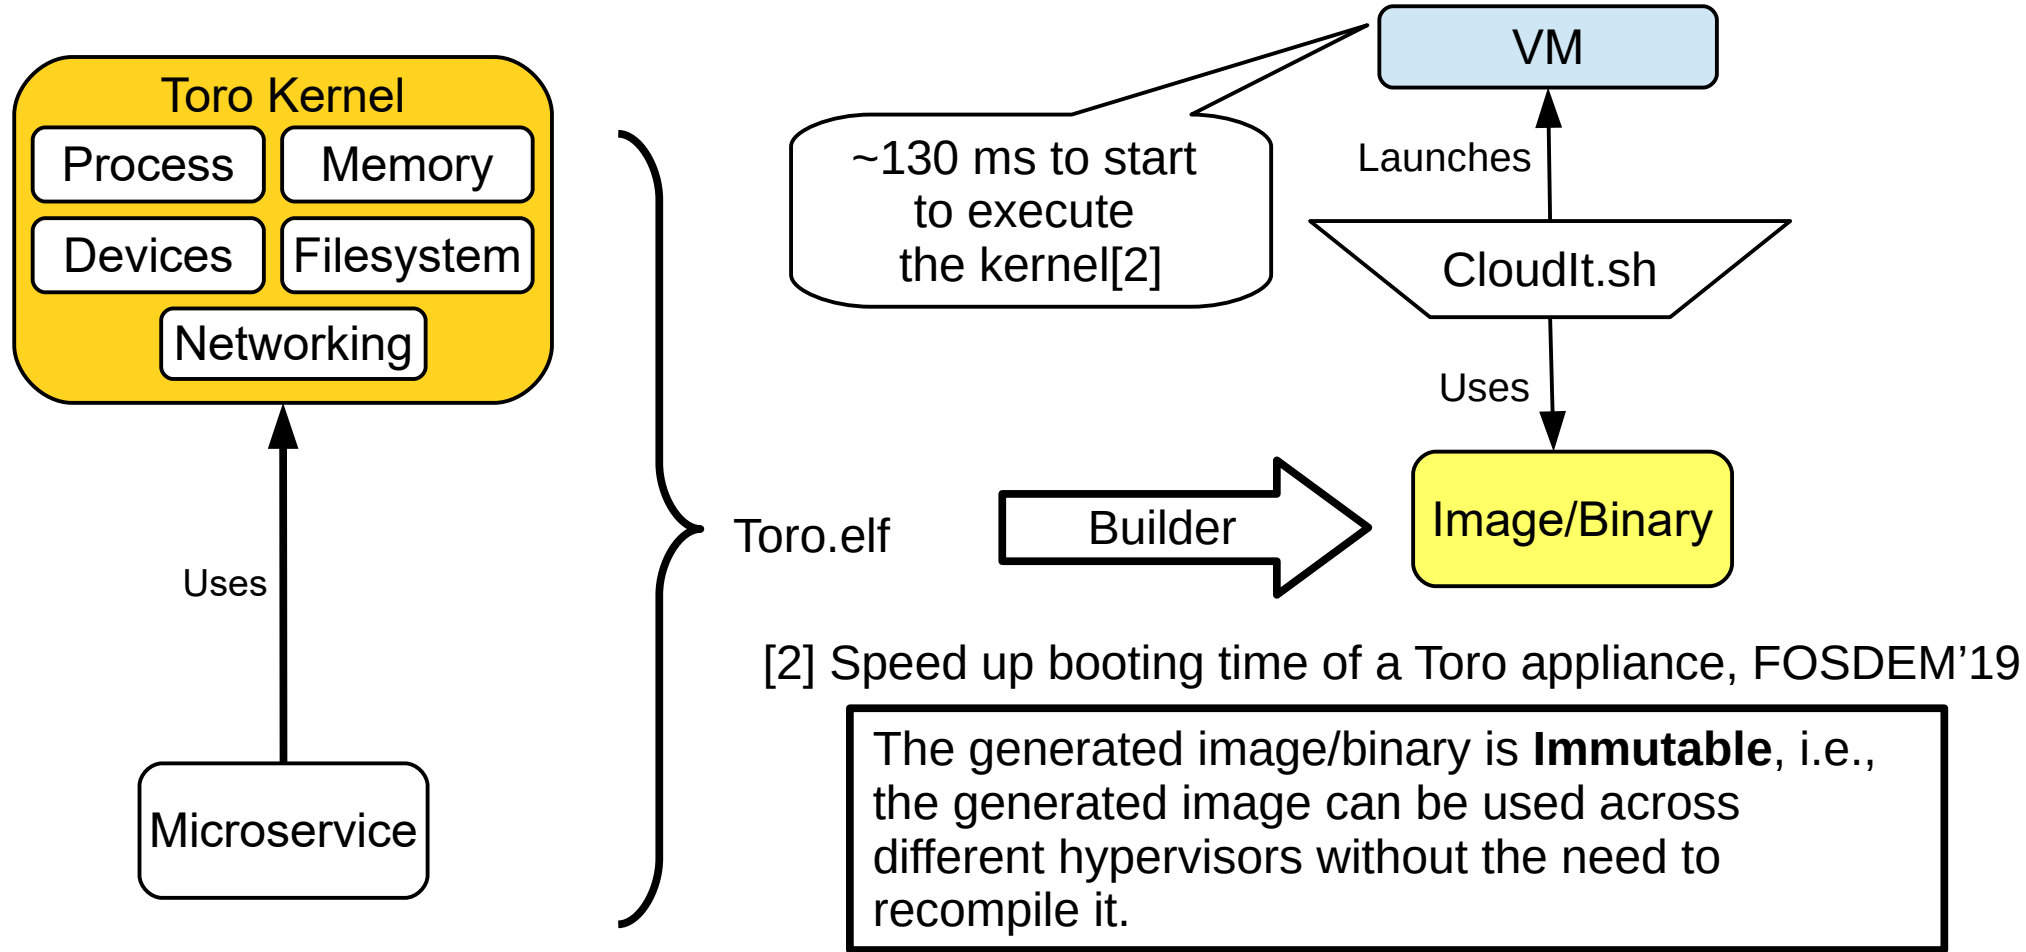
# Application-oriented Kernel



# Application-oriented Kernel



# Application-oriented Kernel



# Benefits of Application-oriented Kernel

- Security is based on the hypervisor which ensures a reduced attack surface
- The kernel size is smaller since it only includes the units required for the application
- The smaller kernel size reduces the booting time and eases image manipulation, i.e., your kernel is your program!
- No communication overhead since application is using the kernel APIs

# Ingredients for ToroKernel

- Application-oriented Kernel
- Simple Scheduler
- Dedicated Resources
- Microservice-oriented Networking



# Simple Scheduler Requirements


- Scale with the number of cores
- Simple scheduler's algorithm, i.e., Cooperative threading model
- Minimalist context switches

# Simple Scheduler

In Toro, there are only threads

Thread 1

Thread 2



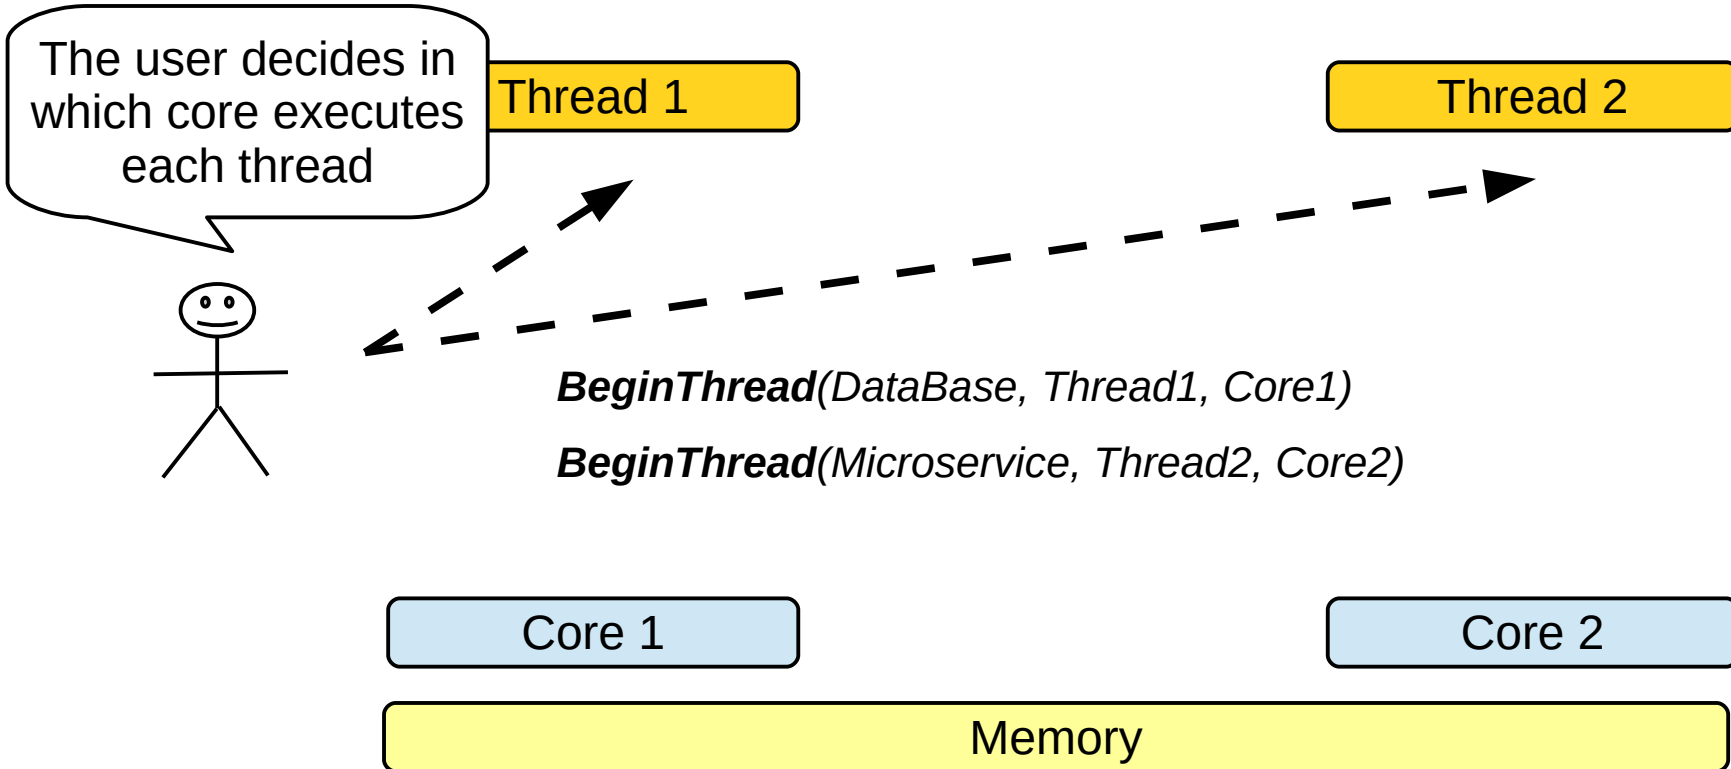
Core 1

Core 2

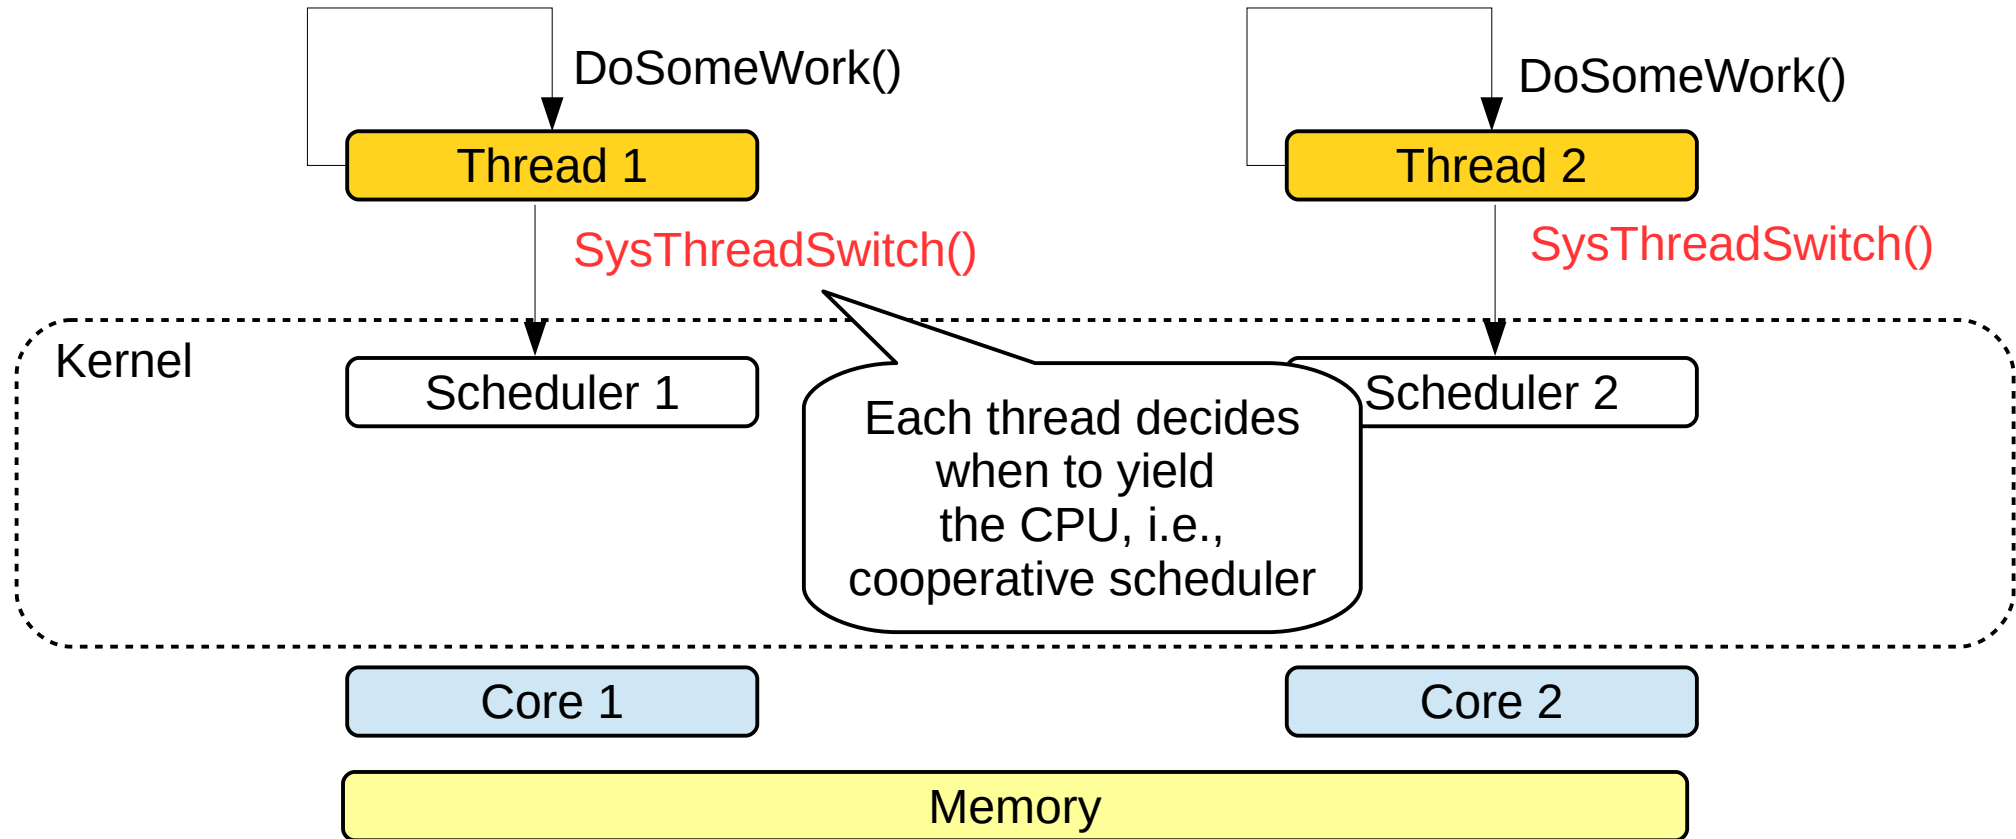
Threads share the memory space

Memory

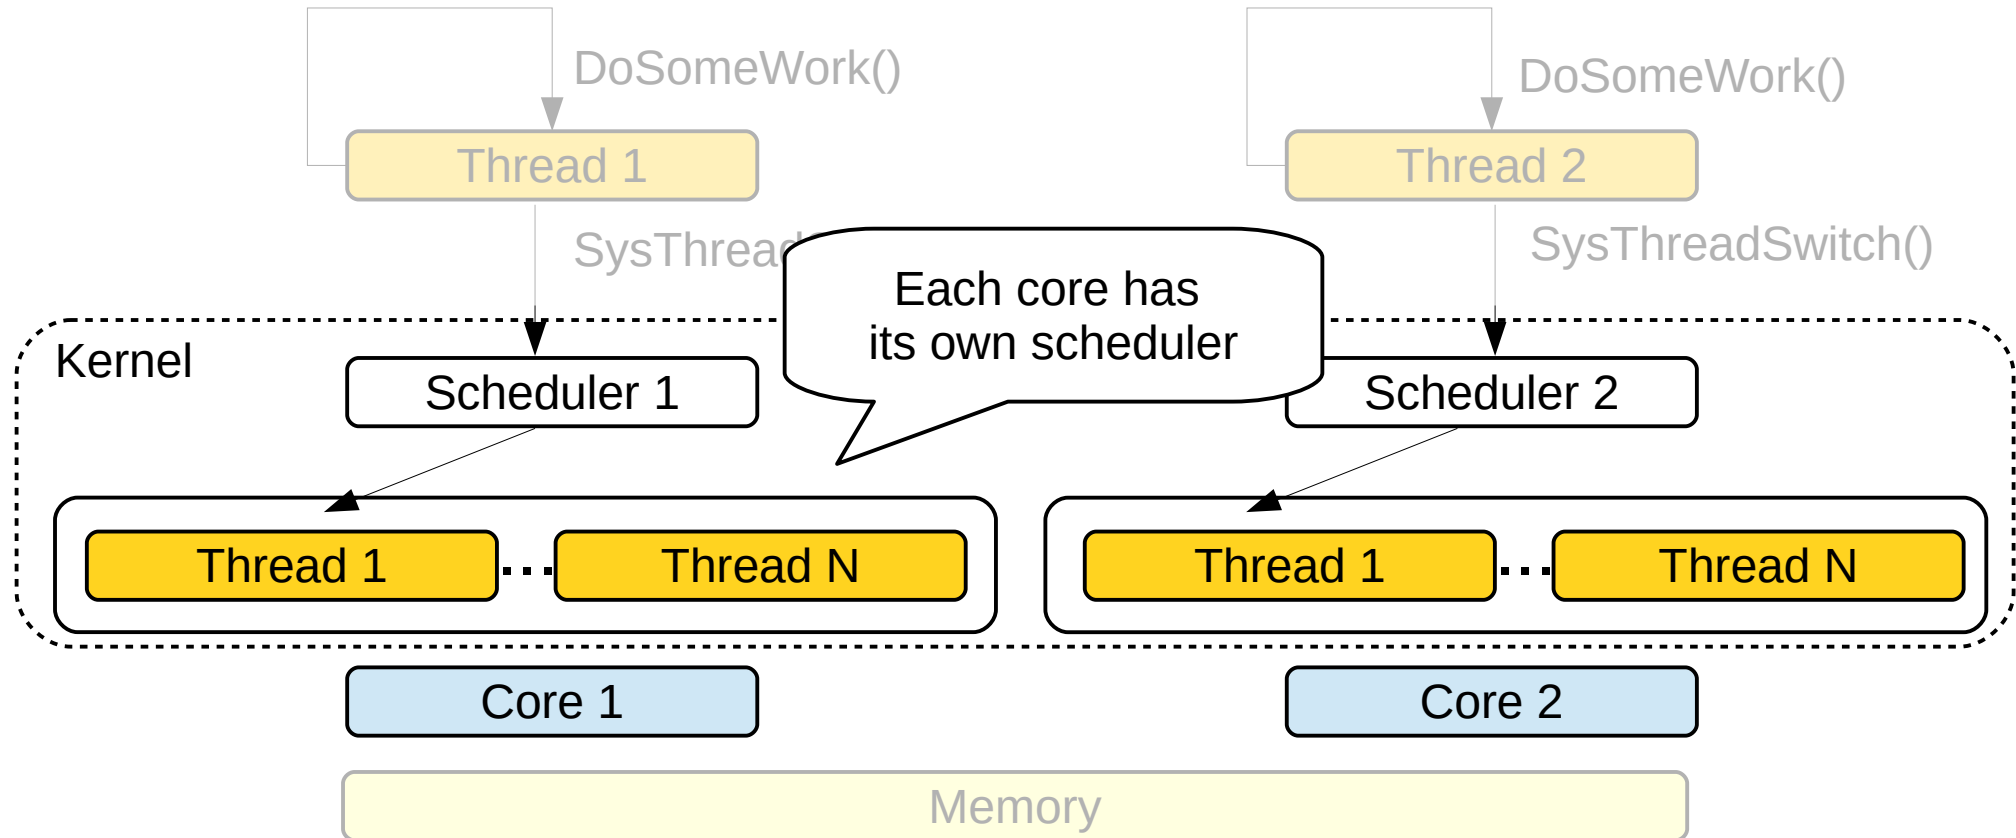
# Simple Scheduler



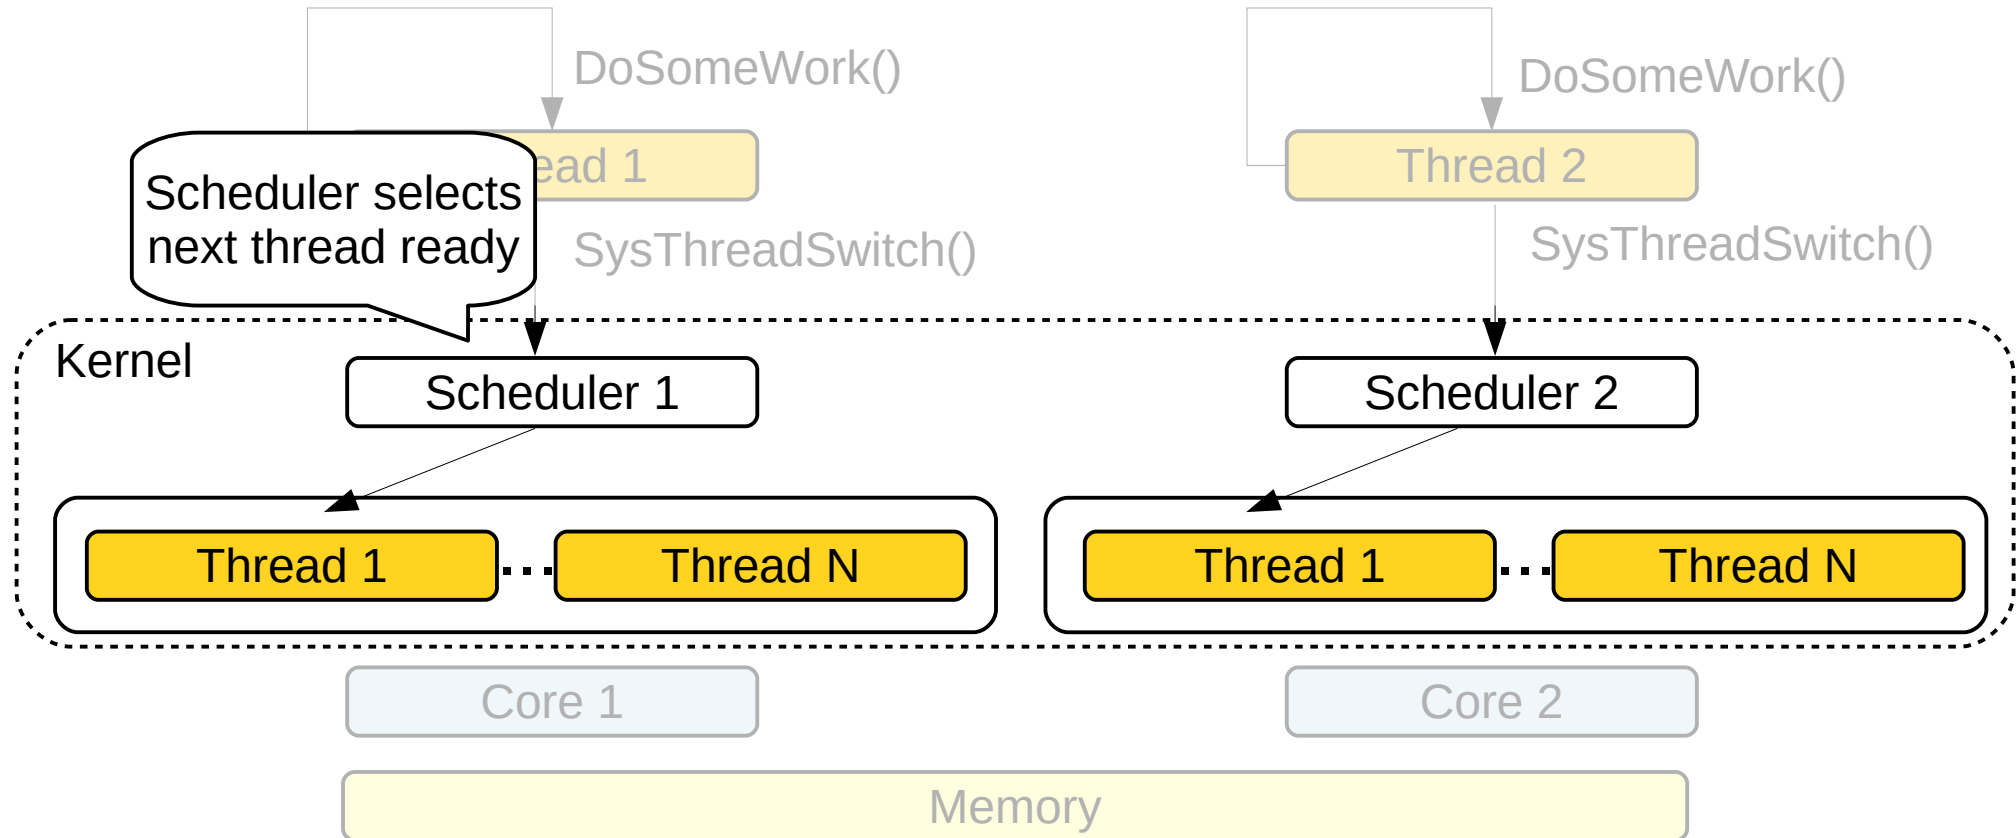
# Simple Scheduler



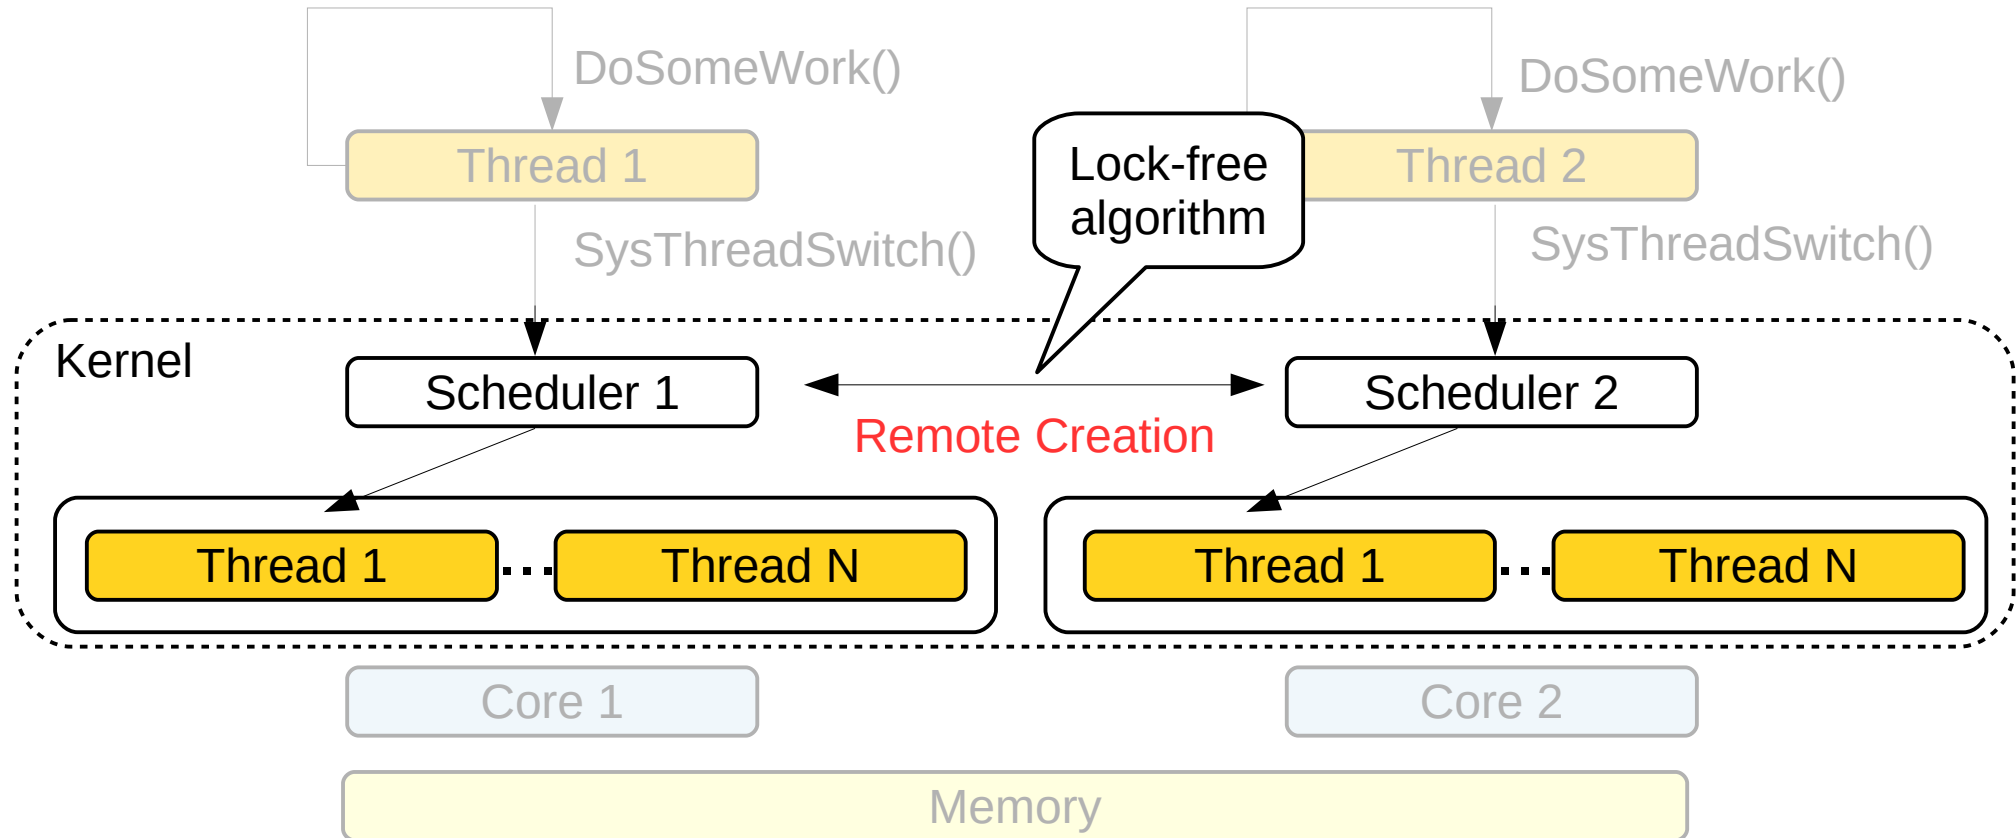
# Simple Scheduler



# Simple Scheduler



# Simple Scheduler



# Benefits of Simple Scheduler

- The use of threads makes context switching cheaper
- The use of a cooperative scheduler:
  - simplifies user's and kernel's code by avoiding to implement protection inside the scheduler
  - reduces number of context switches since it does not relies on interruptions



# Ingredients for ToroKernel

- Application-oriented Kernel
- Simple Scheduler
- Dedicated Resources
- Microservice-oriented Networking

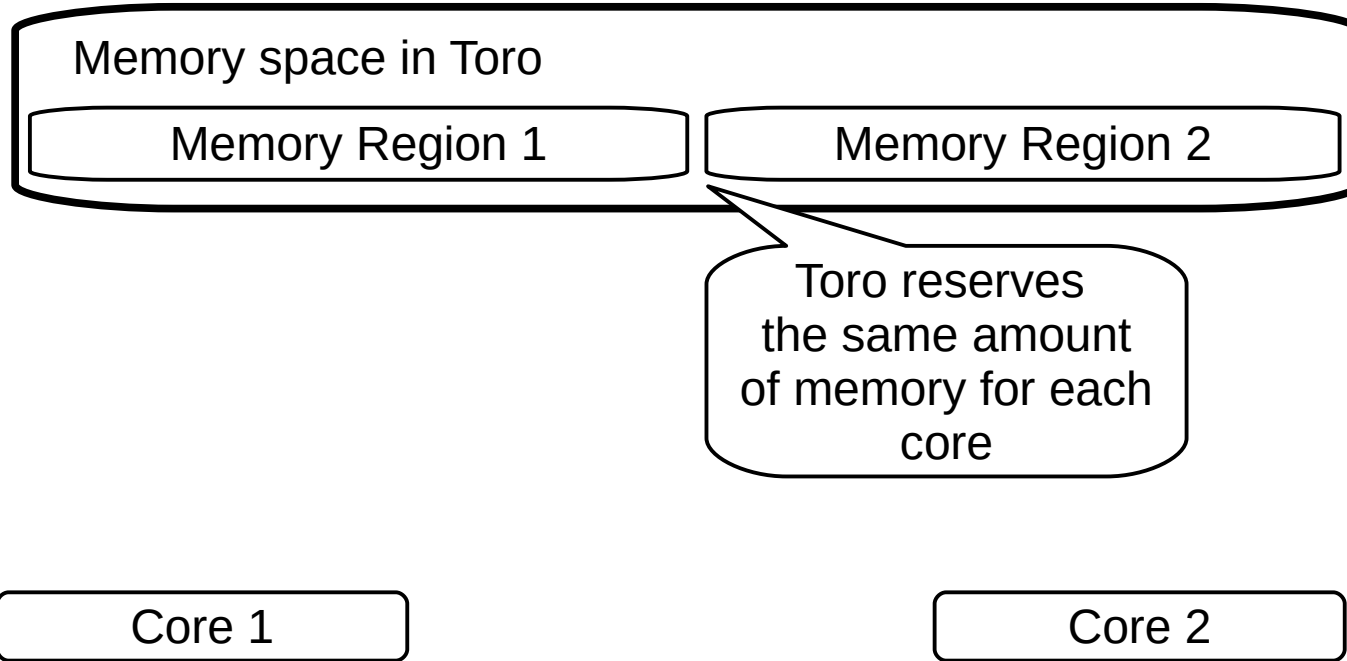
# Dedicated Resources

- In a multicore system, the problematic resource is the shared memory. The use of shared memory causes:
  - Overhead in the memory bus
  - Overhead in the cache to keep it coherent
  - Overhead to guaranty mutual exclusion when spinlocks are used

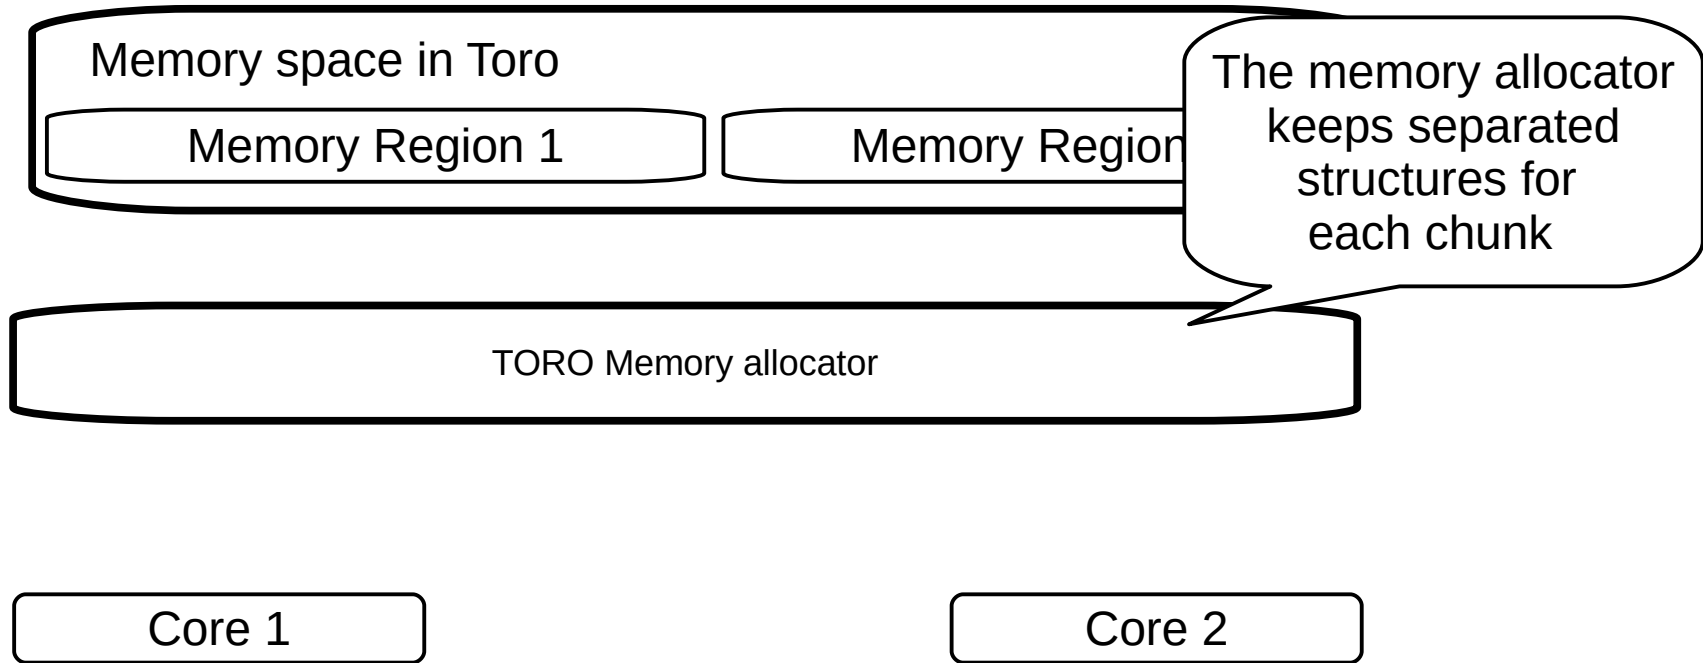
# Dedicated Resources Proposal

- Toro improves memory access by
  - keeping the resources locals:
    - The memory is dedicated per core
    - The kernel data structures are dedicated per core
    - The access to kernel data structures is lock free
  - leveraging NUMA technologies, e.g., Hypertransport, Intel QuickPath

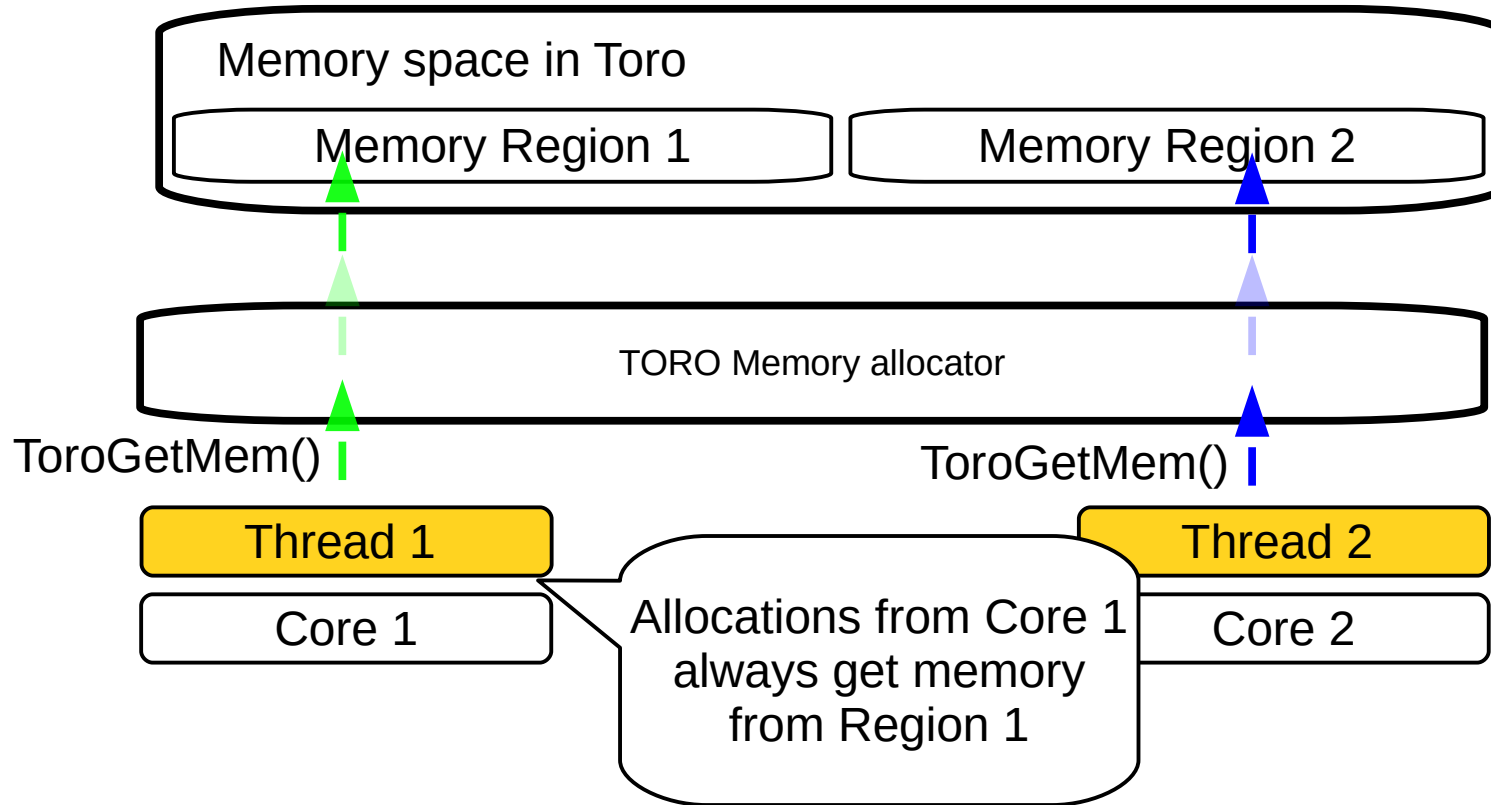
# Dedicated Resources Memory



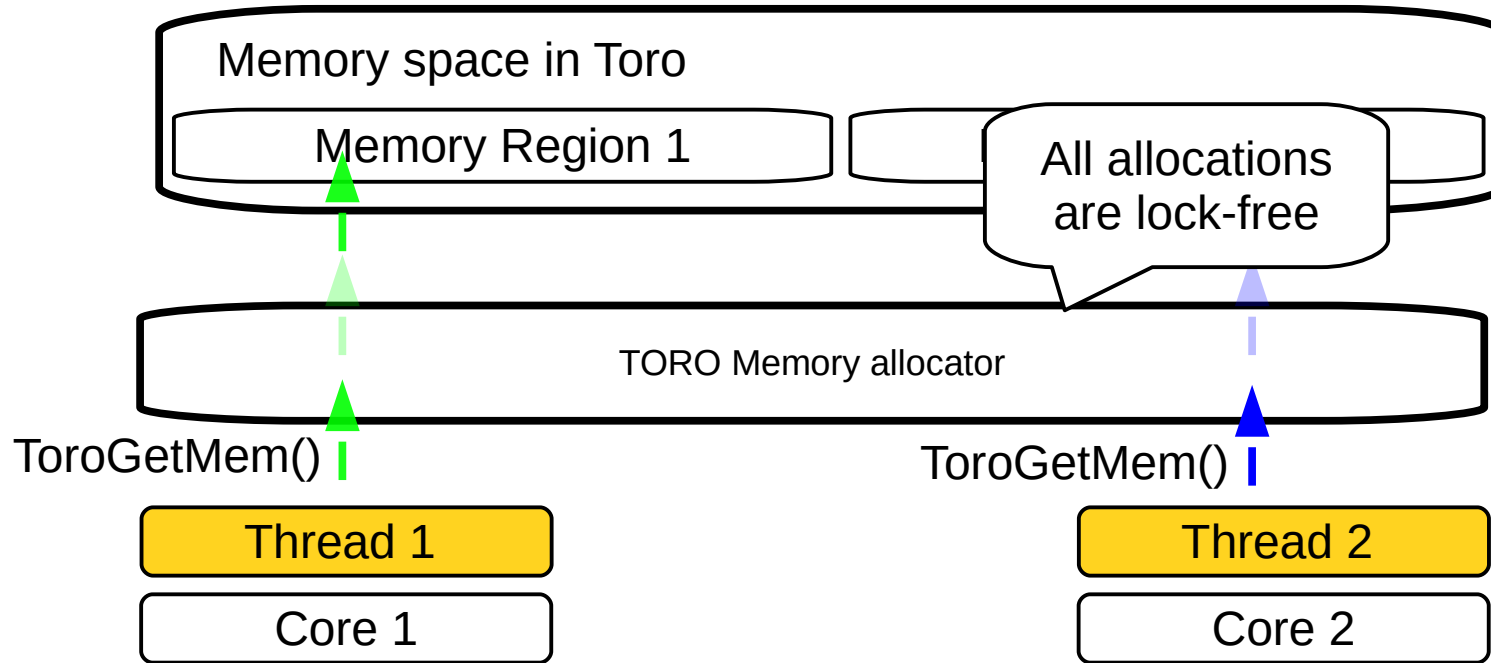
# Dedicated Resources Memory



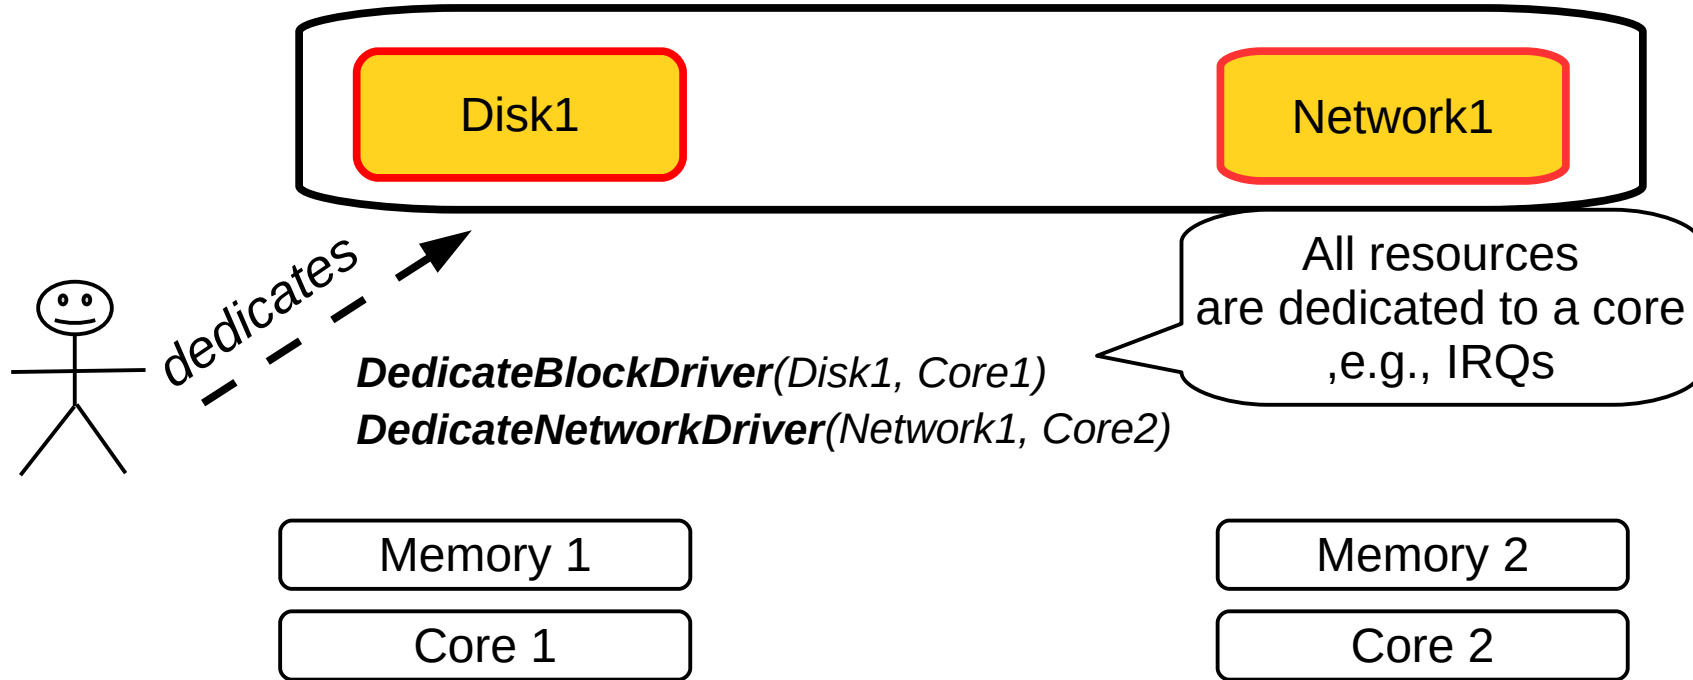
# Dedicated Resources Memory



# Dedicated Resources Memory

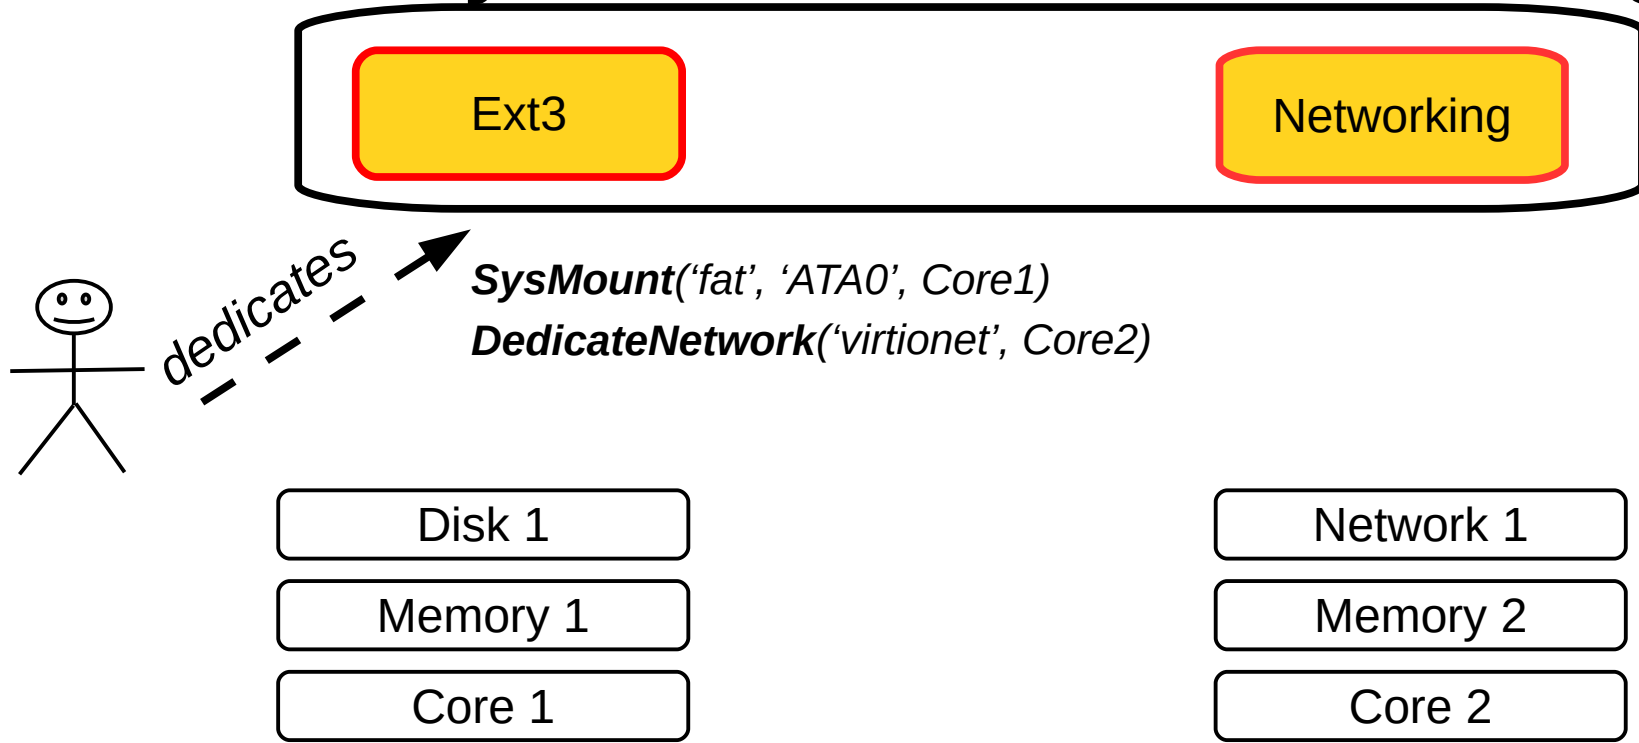


# Dedicated Resources Devices



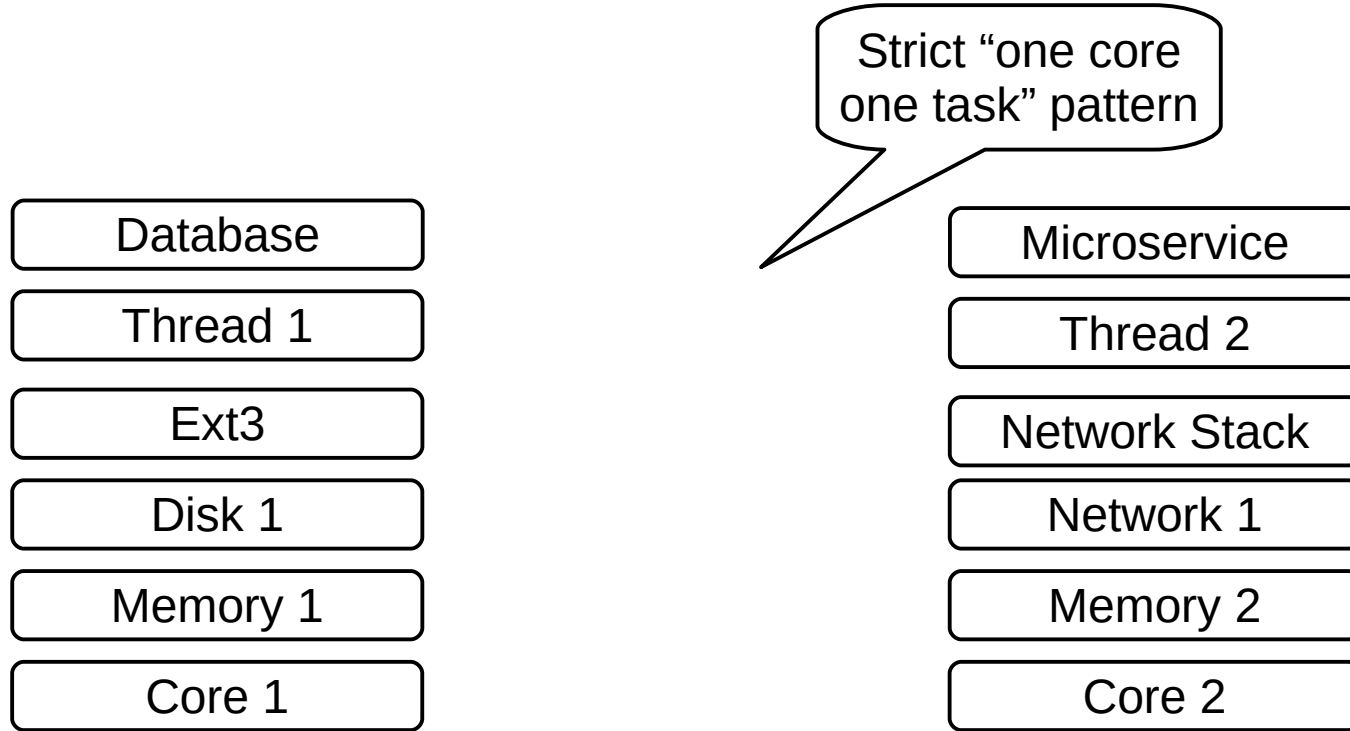


# Dedicated Resources Filesystems and Networking



# Dedicated Resources

## Filesystems and Networking



# Ingredients for ToroKernel

- Application-oriented Kernel
- Simple Scheduler
- Dedicated of Resources
- **Microservice-oriented Networking**

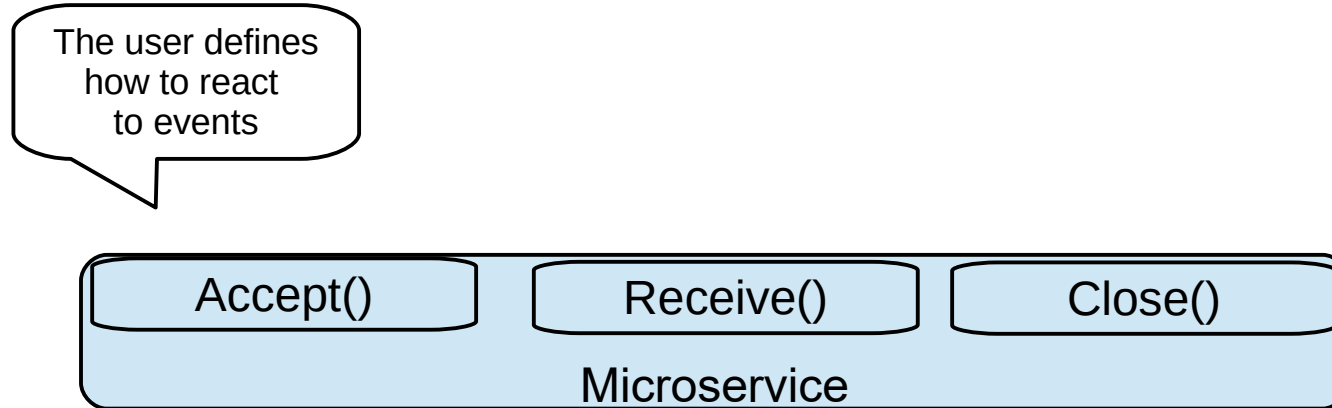
# Microservice-oriented Networking

- The kernel provides a dedicated API to create microservices
- Toro proposes two ways to develop microservices:
  - Single thread event loop model[2] for non-blocking microservices
  - Classical Berkeley sockets for those microservices that do IO
- Event loop has been optimized to reduce CPU's usage[3]

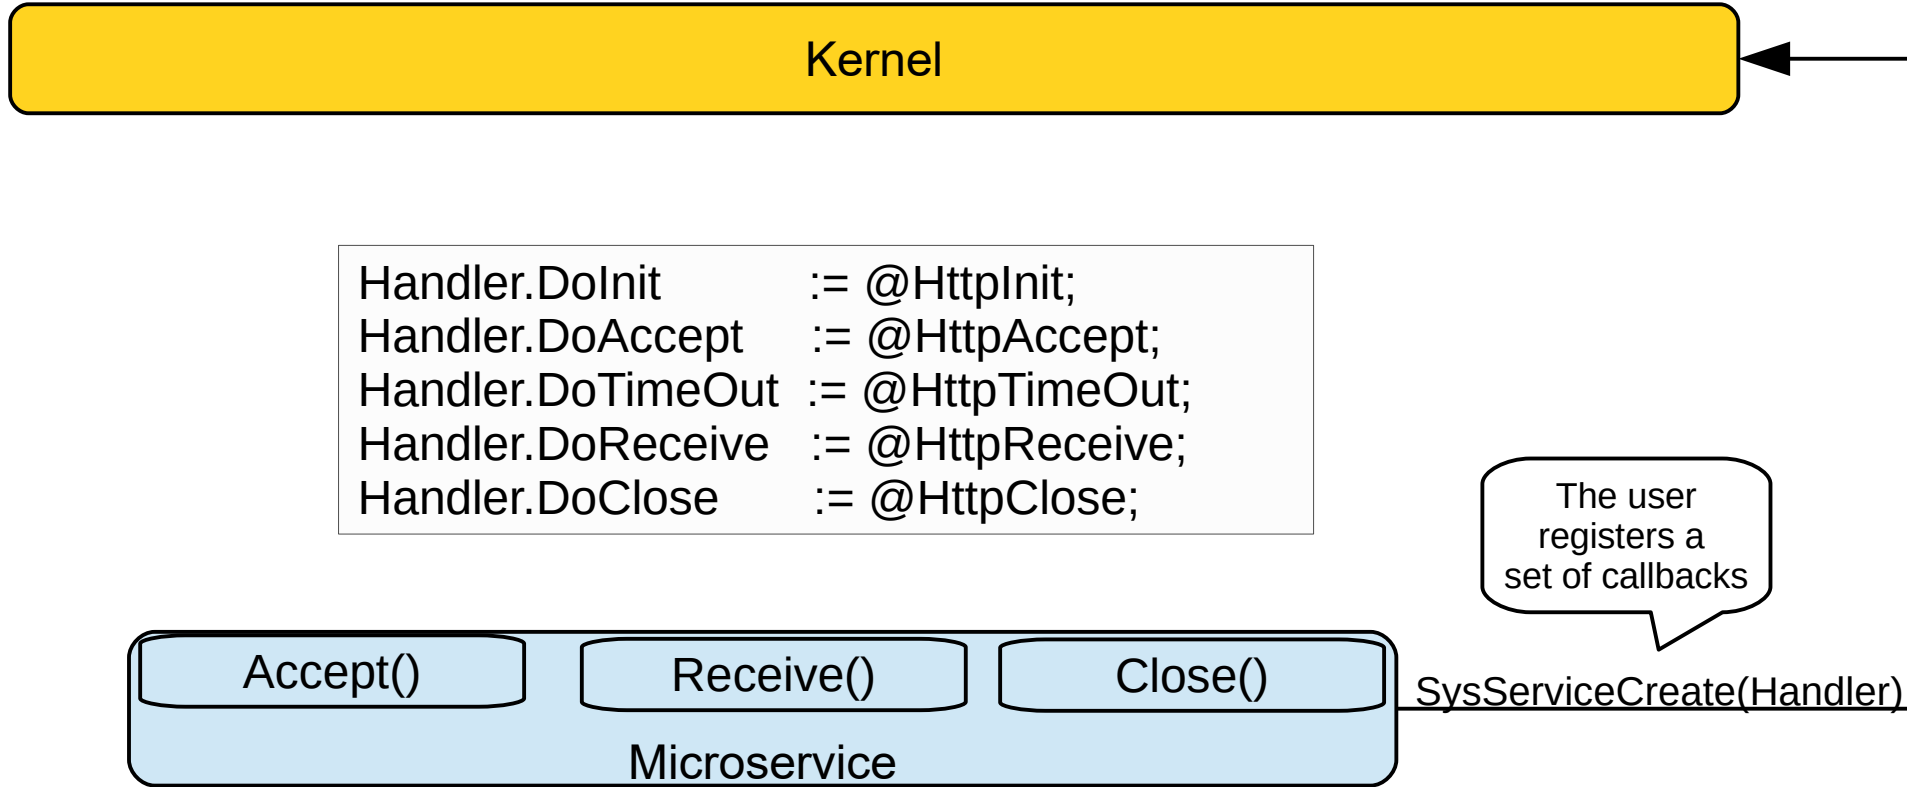
[2] Node JS Architecture – Single Threaded Event Loop

[3] Reducing CPU usage of a Toro Appliance, FOSDEM'18

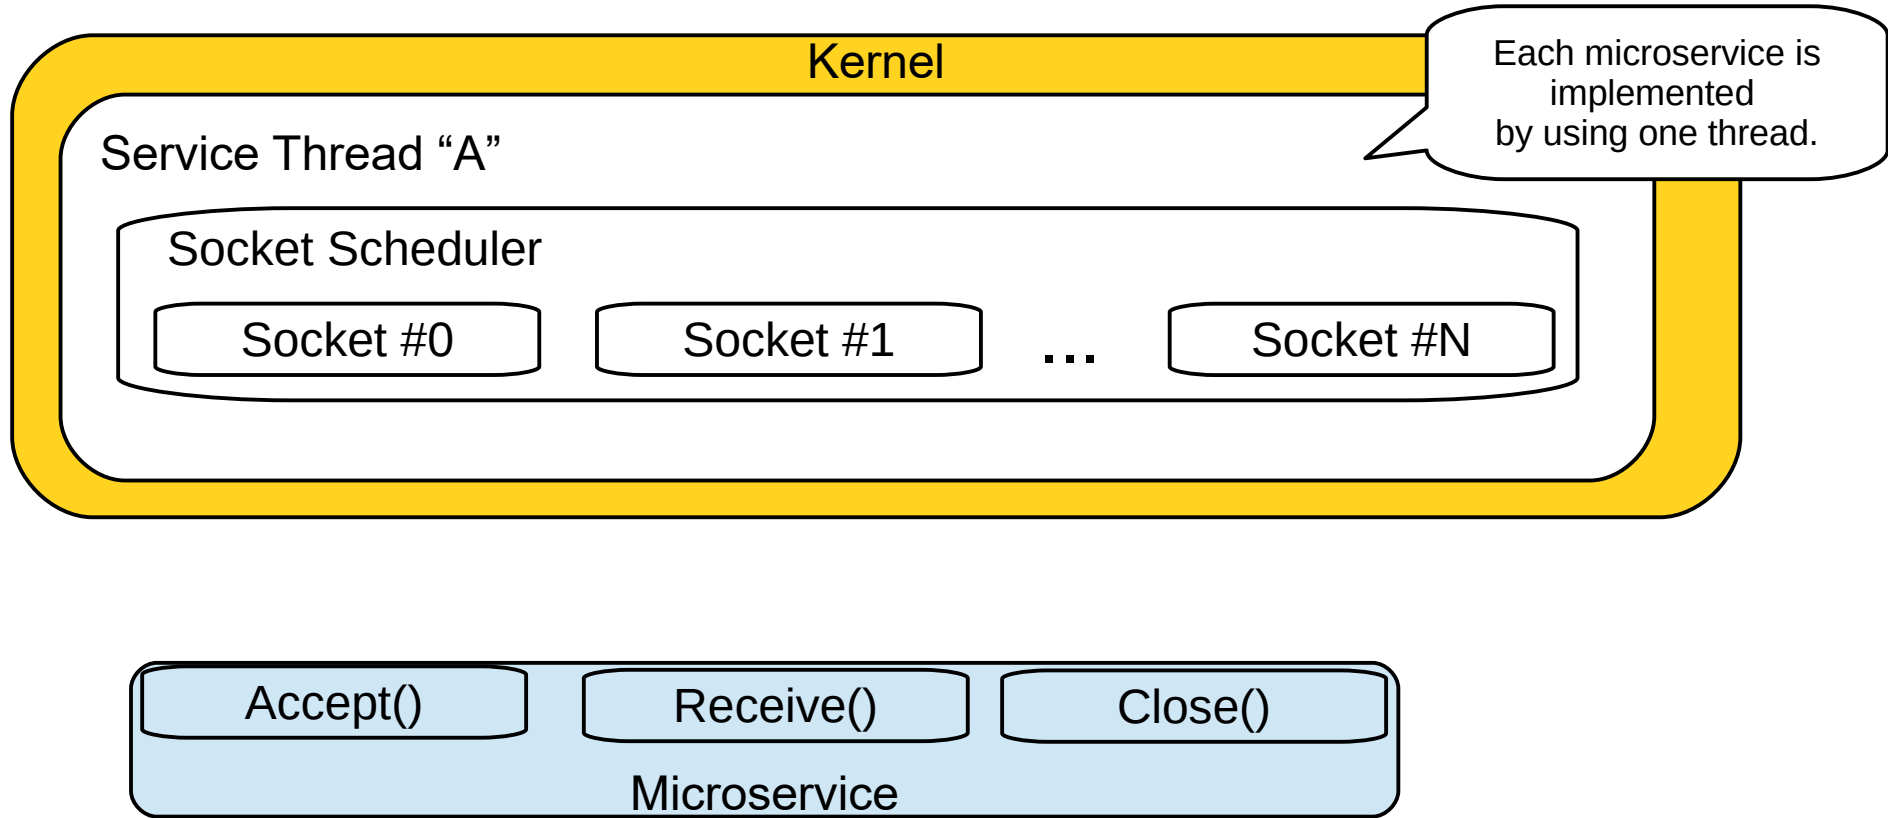
# Non-Blocking Microservices



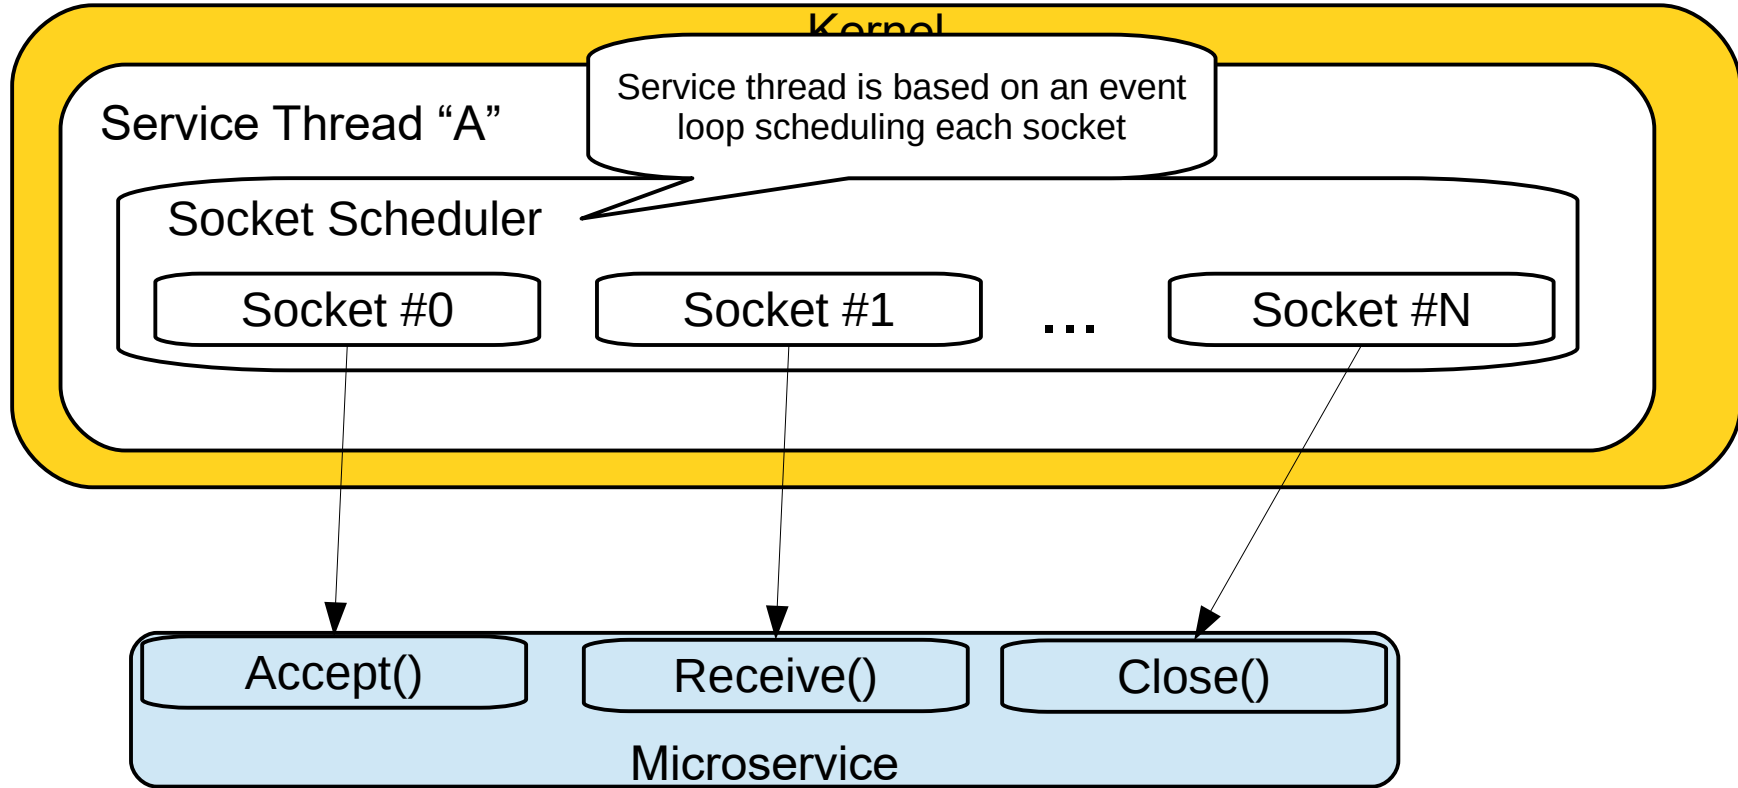
# Non-Blocking Microservices



# Non-Blocking Microservices

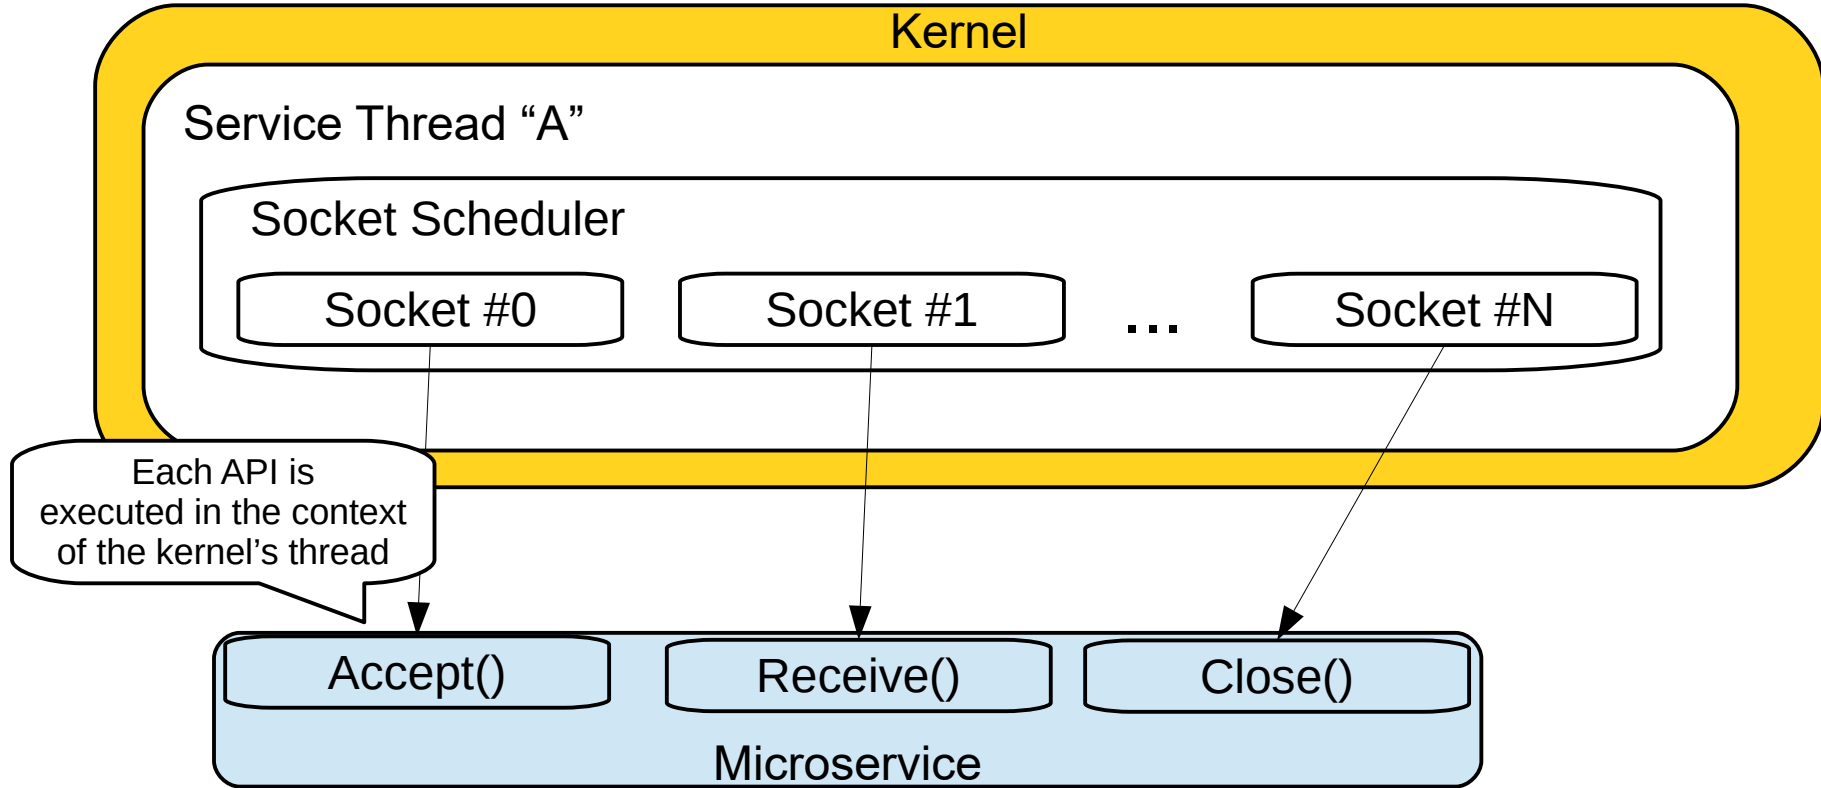


# Non-Blocking Microservices

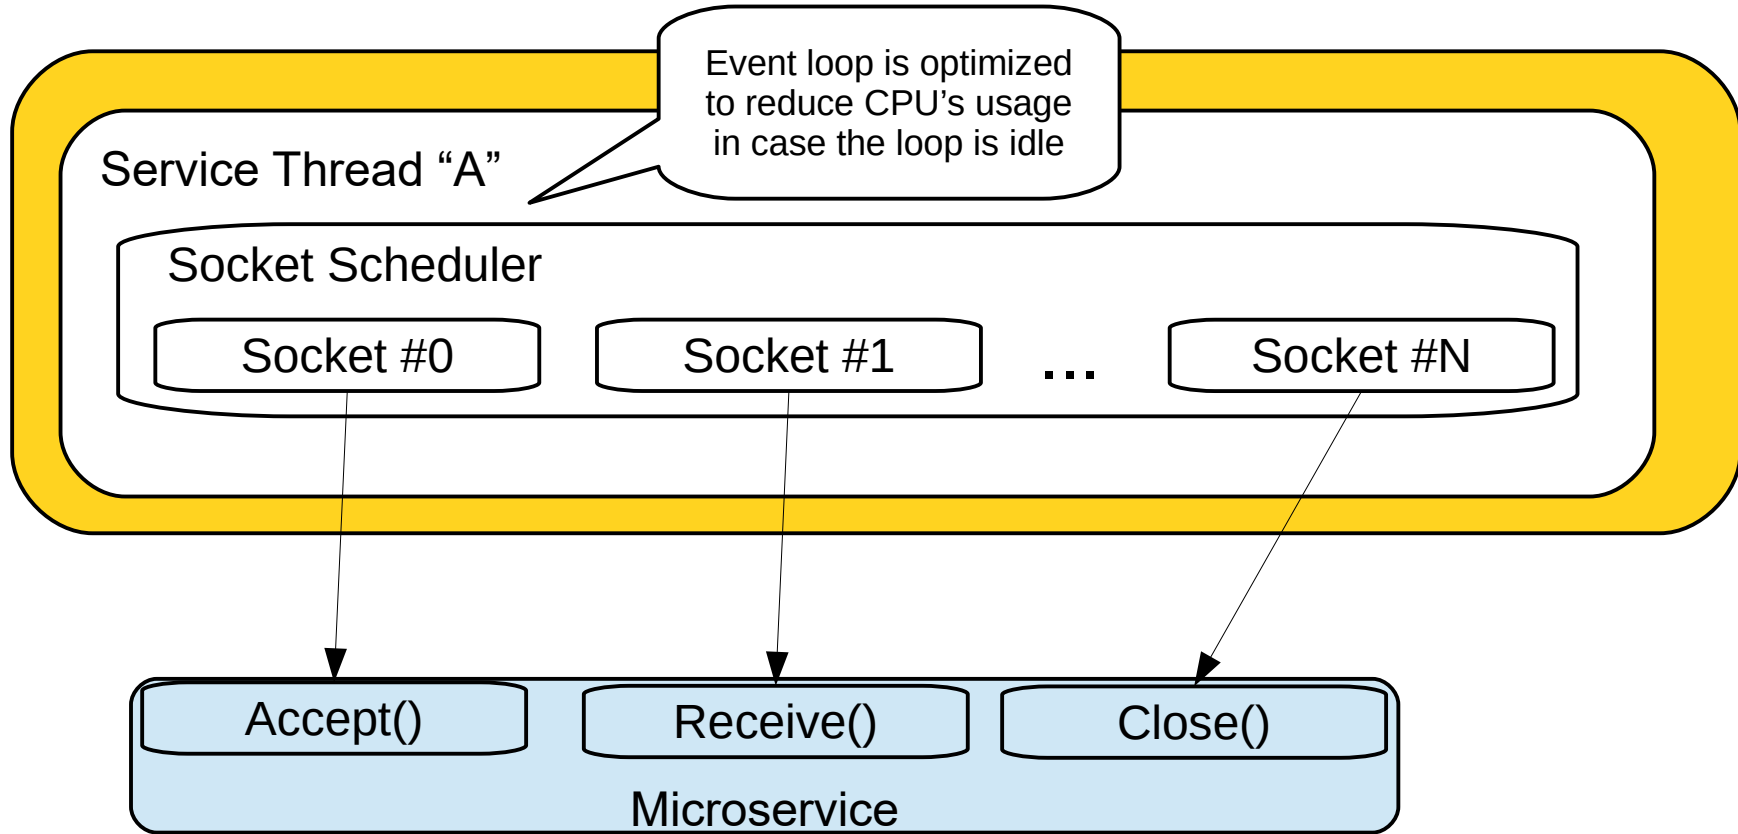




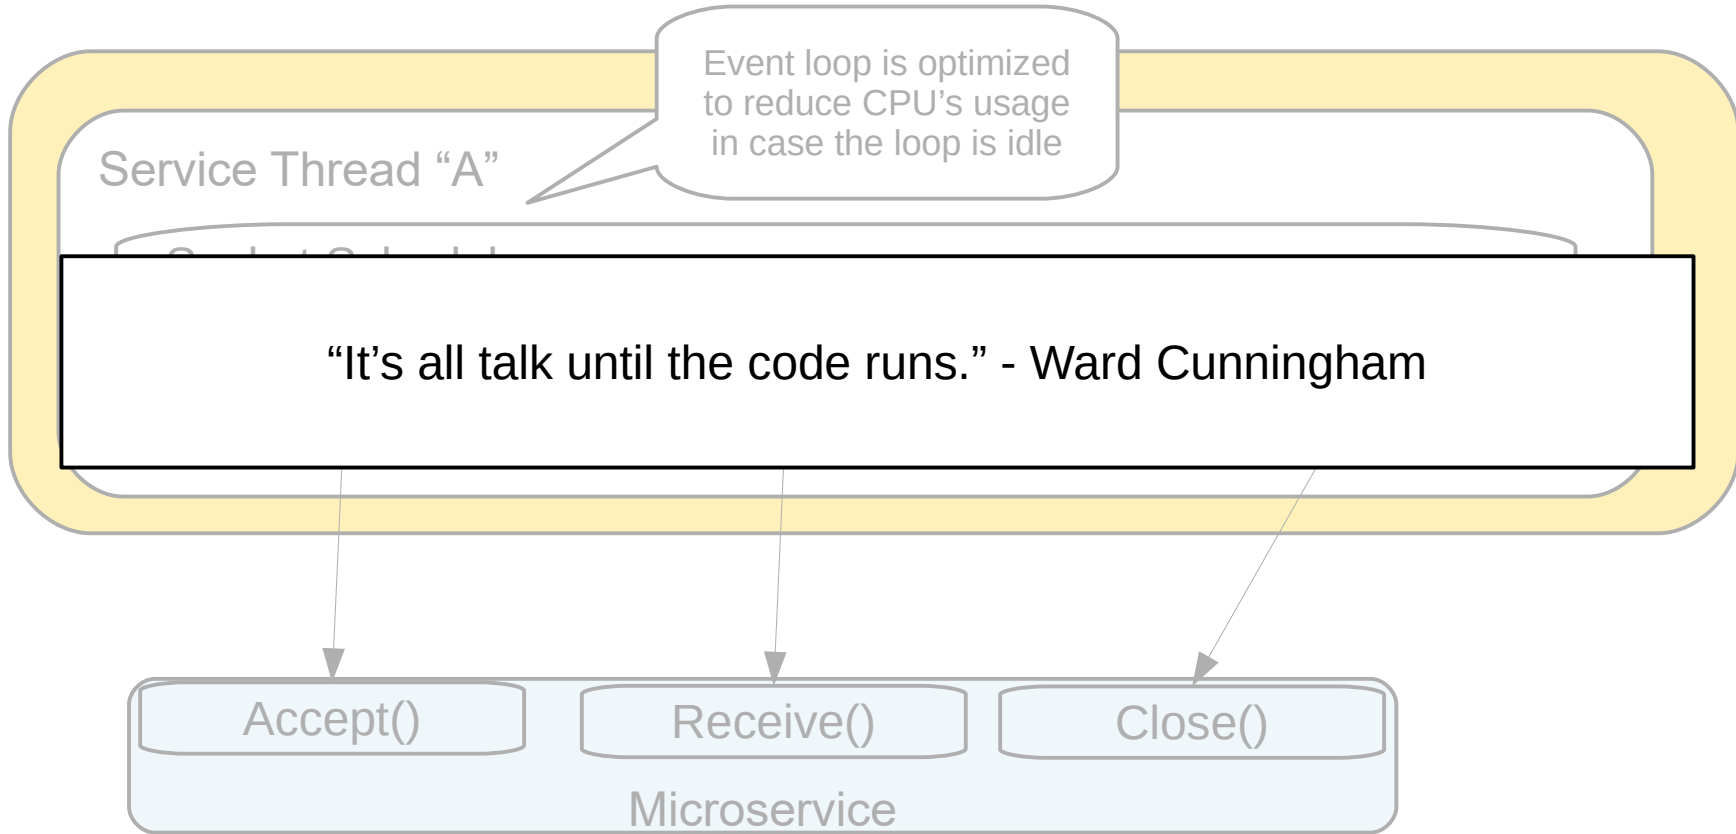
# Non-Blocking Microservices



# Non-Blocking Microservices




# Non-Blocking Microservices



# Evaluation

- We implement a “Hello World” microservice[4] (see <https://github.com/MatiasVara/torokernel/tree/master/examples/HelloWorldMicroservice>)
- We deploy it in Docker, Ubuntu (KVM Guest) and Toro (KVM Guest)
- We chose the following technology:
  - NGINX, UWSGI (4 processes) and Flask
- Memory is limited in all the platforms
- We compare these approaches in term of:
  - Deploying Time
  - Bootstrap Time
  - Image Size
  - CPU usage
  - Time per Request




Microservice's footprint

[4] “Toro, a Dedicated Kernel for Microservices”, OSSummit'18


# Deploying Time

- Time required to build an image within the microservice




# Bootstrap Time

- Time to boot and to serve the first request



# Bootstrap Time


- Time to boot and to serve the first request



[3] Speed up Booting time of a Toro Appliance, FOSDEM'19

# Image Size


- Size of the image that contains the microservice and its dependencies.






# Image Size

- Size of the image that contains the microservice and its dependencies.




# CPU Usage

TORO CPU Usage



Docker CPU Usage



# End-User Delay

- Benchmark with ab and measuring the Time per Request (mean) [ms]

		Number of Concurrent Request		
Approach	CPUs	200	500	1000
Docker	4	139.980 ms	333.937 ms	801.422
Ubuntu/KVM	1	94.507 ms	238.149 ms	560.513 ms
TORO/KVM	1	120.065 ms	301.736 ms	596.792 ms

# Take away lessons

- Toro image is about 15-times less than NGINX docker image
- 1 sec to re-deploy a microservice
  - Deploy an OS w/similar configuration takes 300 sec, with docker ~50 sec
- Comparable time per request with cutting edge technology (NGINX)
- CPU Usage comparable with Docker
- Safer than Docker

# Demo time

# Summary

- Toro design is improved in five main points:
  - Booting time and building time
  - Kernel API with zero overhead function call
  - Access to shared memory
  - Networking
  - CPU usage

# Future Work

- Improve tooling to build, test and debug microservices
- Investigate new use-cases
- Port networking applications to provide starter kits: SMTP relay, HTTP proxy, Web tracking
- Investigate new ideas to improve the network stack for microservices, e.g., improve socket scheduling for http, resource allocation algorithm
- Investigate the use of microVM technologies, e.g., Firecracker, NEMU.
- Investigate the use of OpenStack to ease the deployment of Toro appliances



# QA

- <http://www.torokernel.io>
- [torokernel@gmail.com](mailto:torokernel@gmail.com)
- Twitter @torokernel
- Torokernel wiki at github
  - My first Three examples with Toro
- Test Toro in 5 minutes (or less...)
  - torokernel-docker-qemu-webservices at Github







# QA

- <http://www.torokernel.io>
- [torokernel@gmail.com](mailto:torokernel@gmail.com)

That's all folks!



torokernel

ki at github

e examples with Toro

15 minutes (or less...)

nel-docker-qemu-webservices at

torokernel