

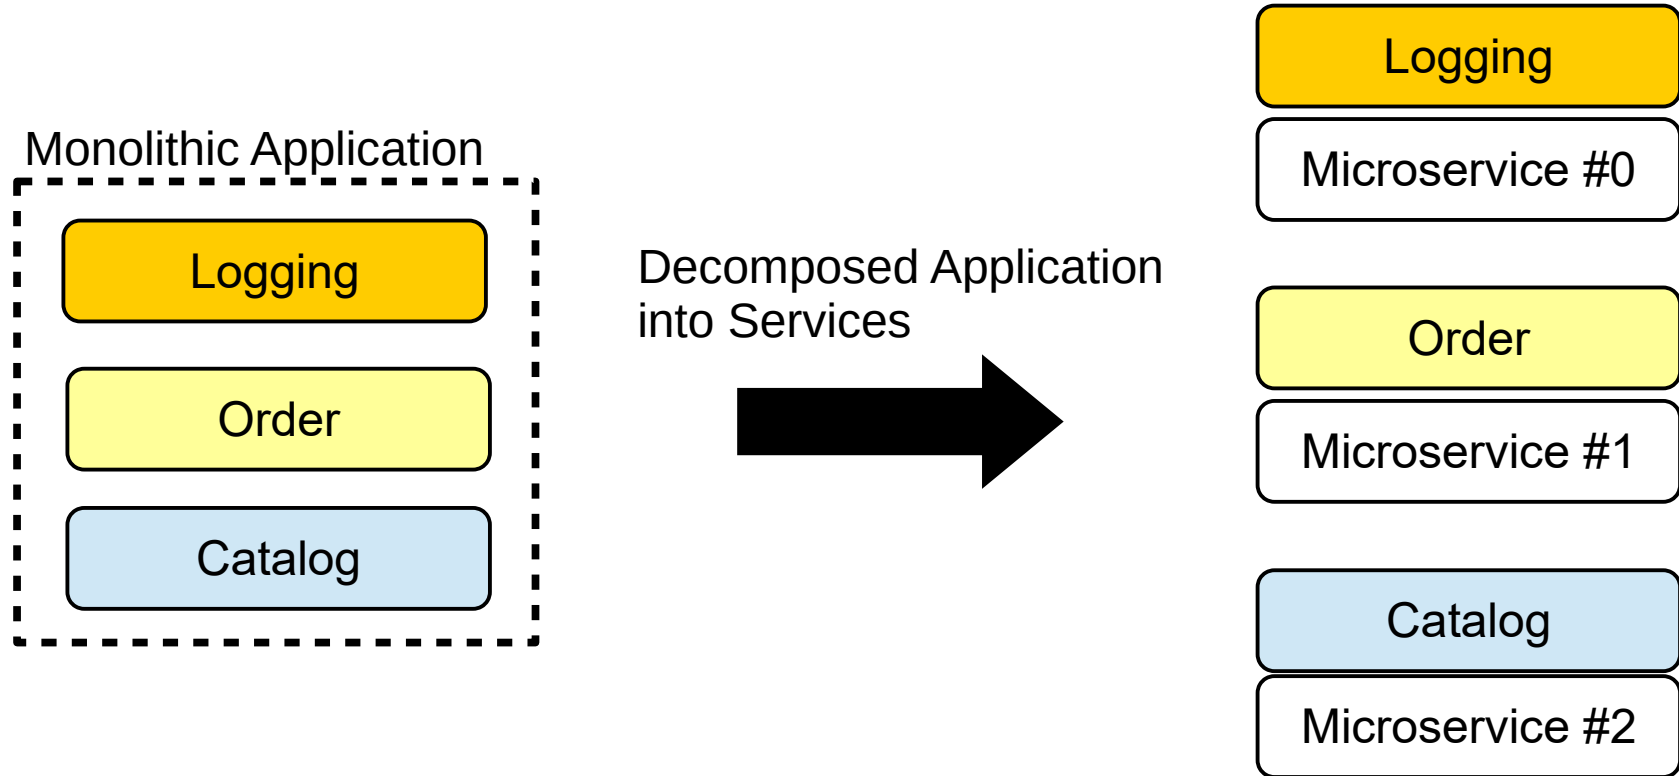


Developing and Deploying Microservices with Toro Kernel

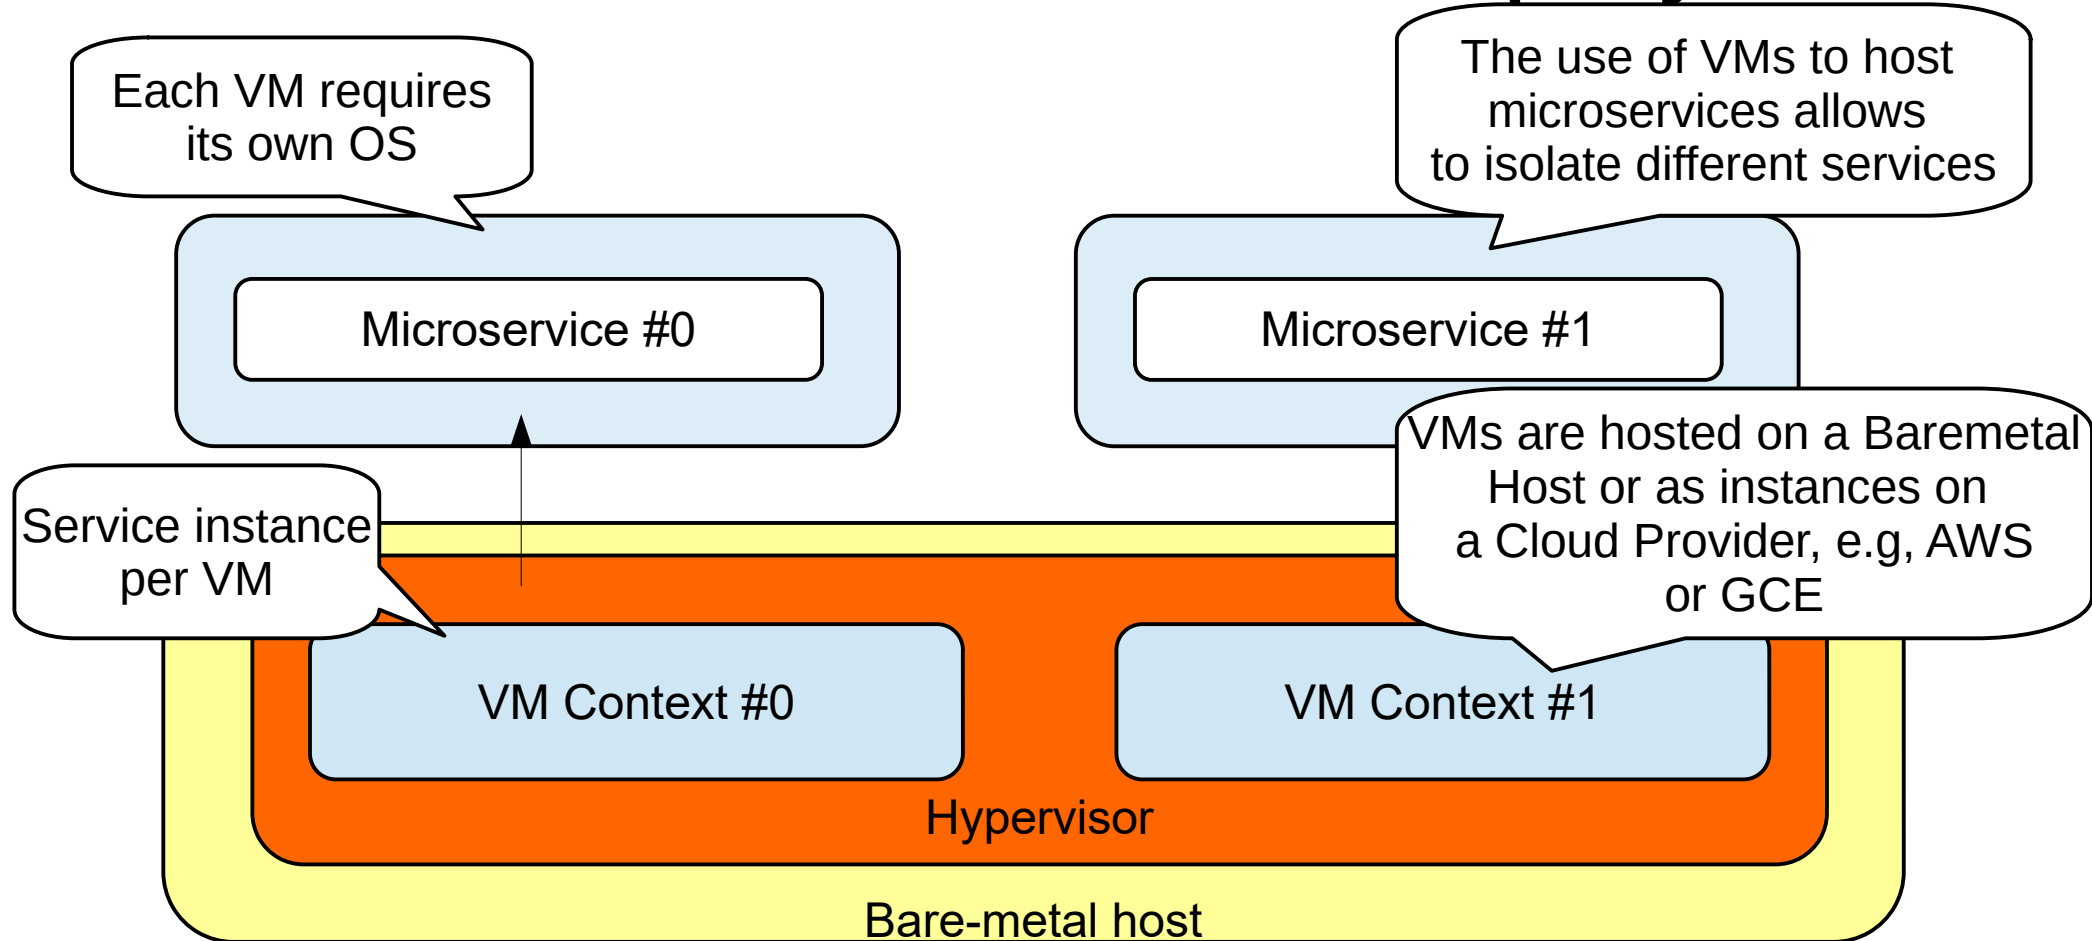
www.torokernel.io

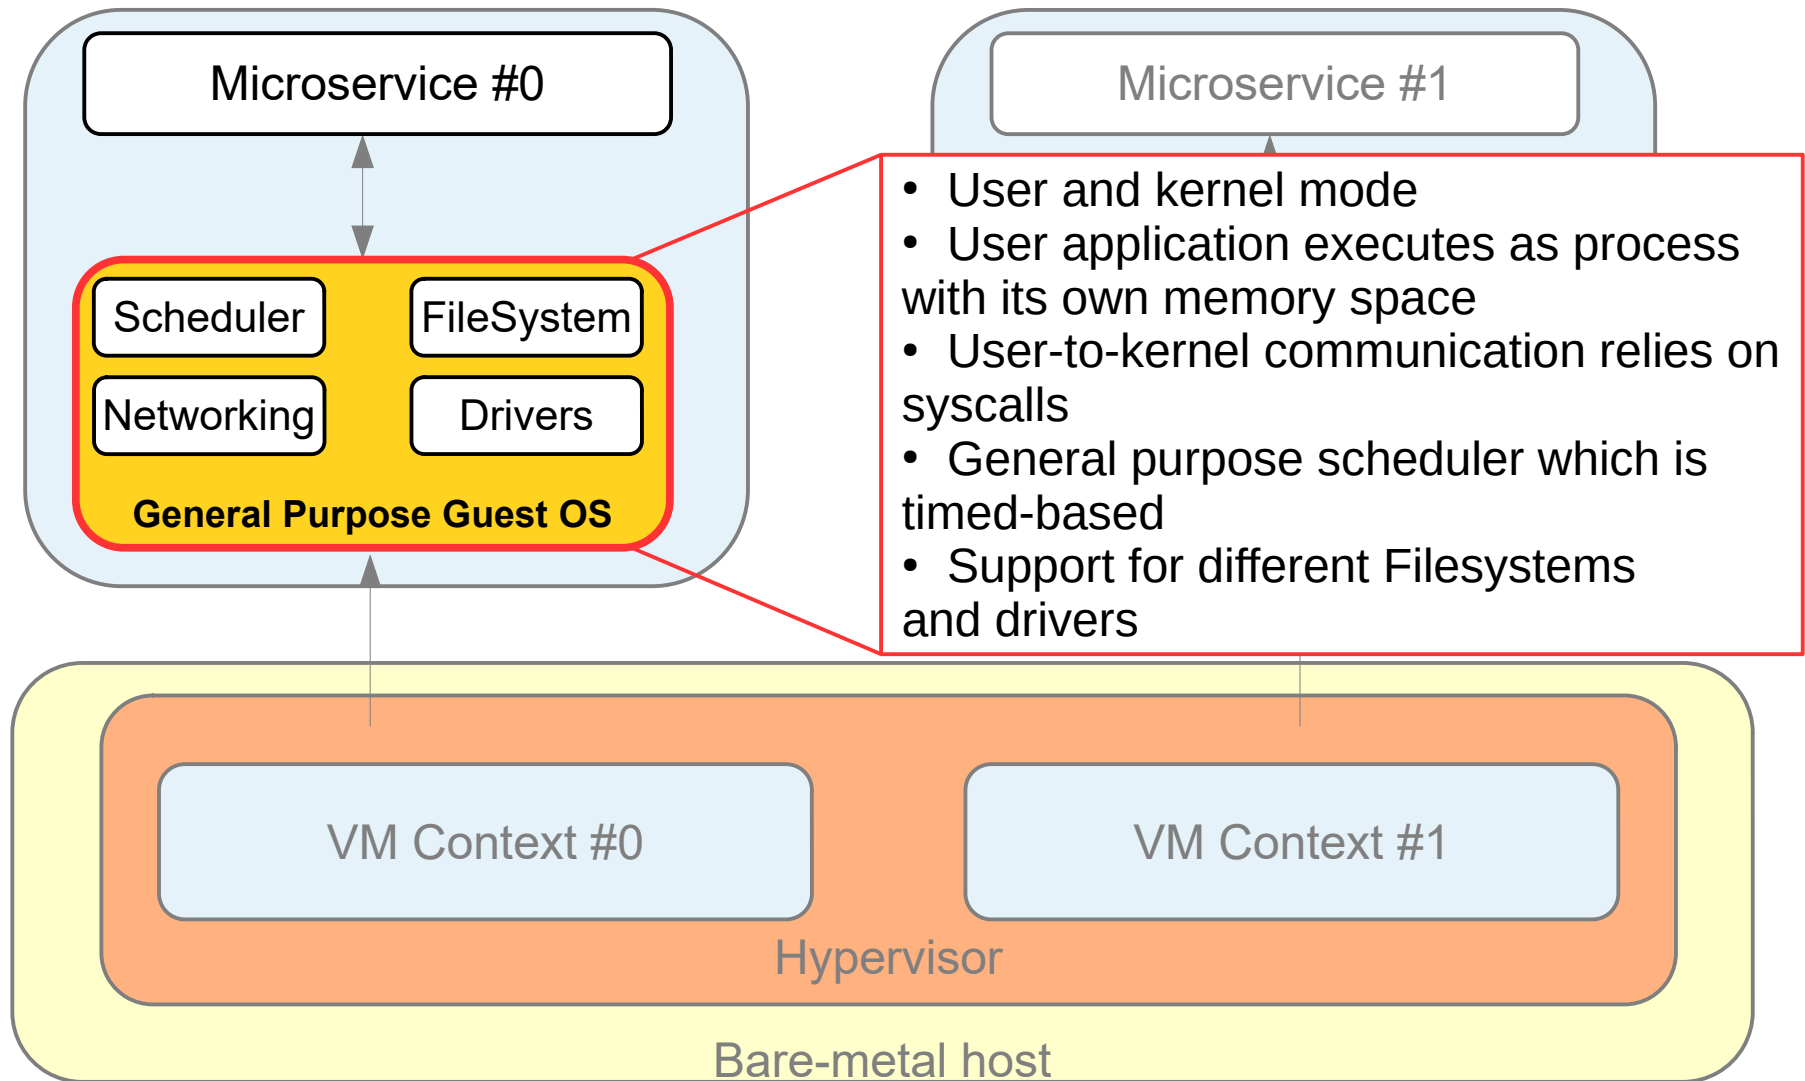
Matias Vara Larsen
matiasevara@gmail.com

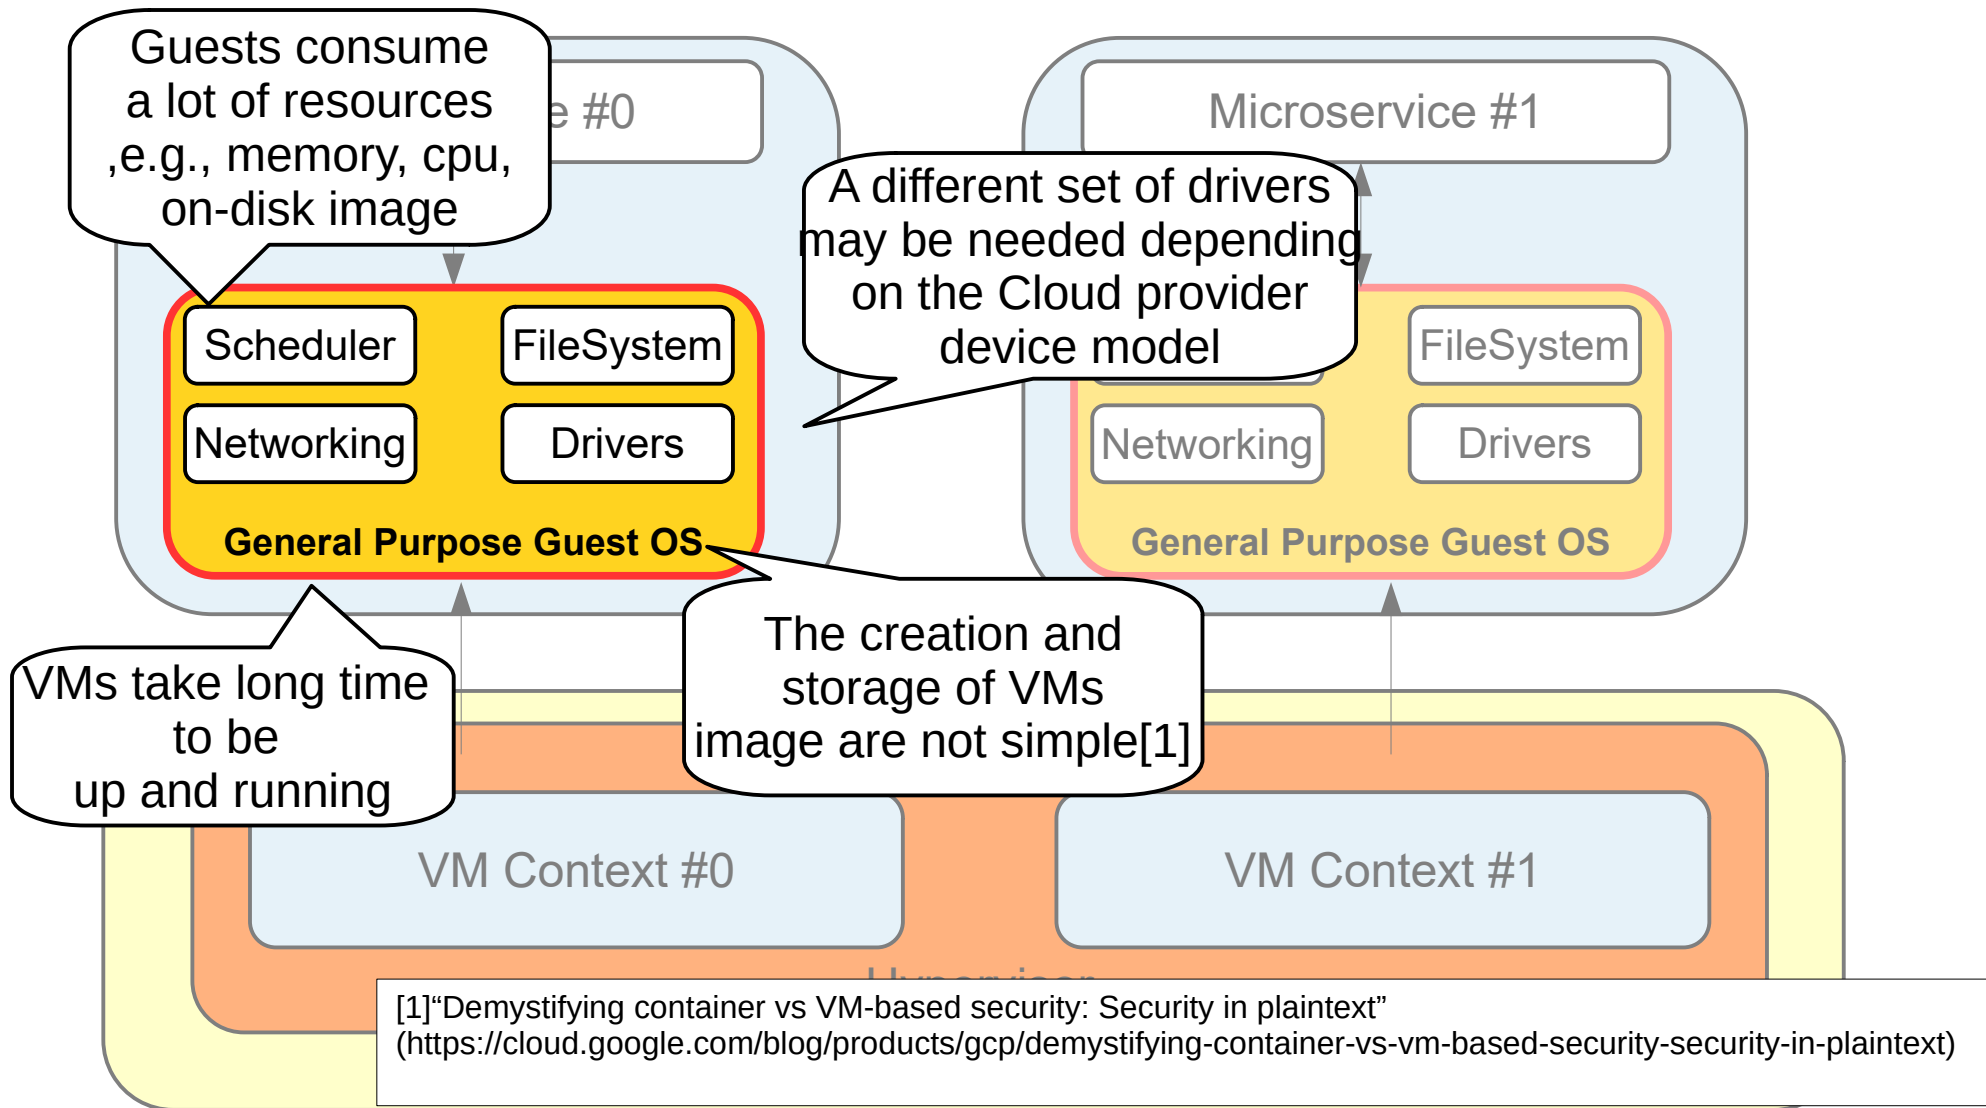
What are microservices?

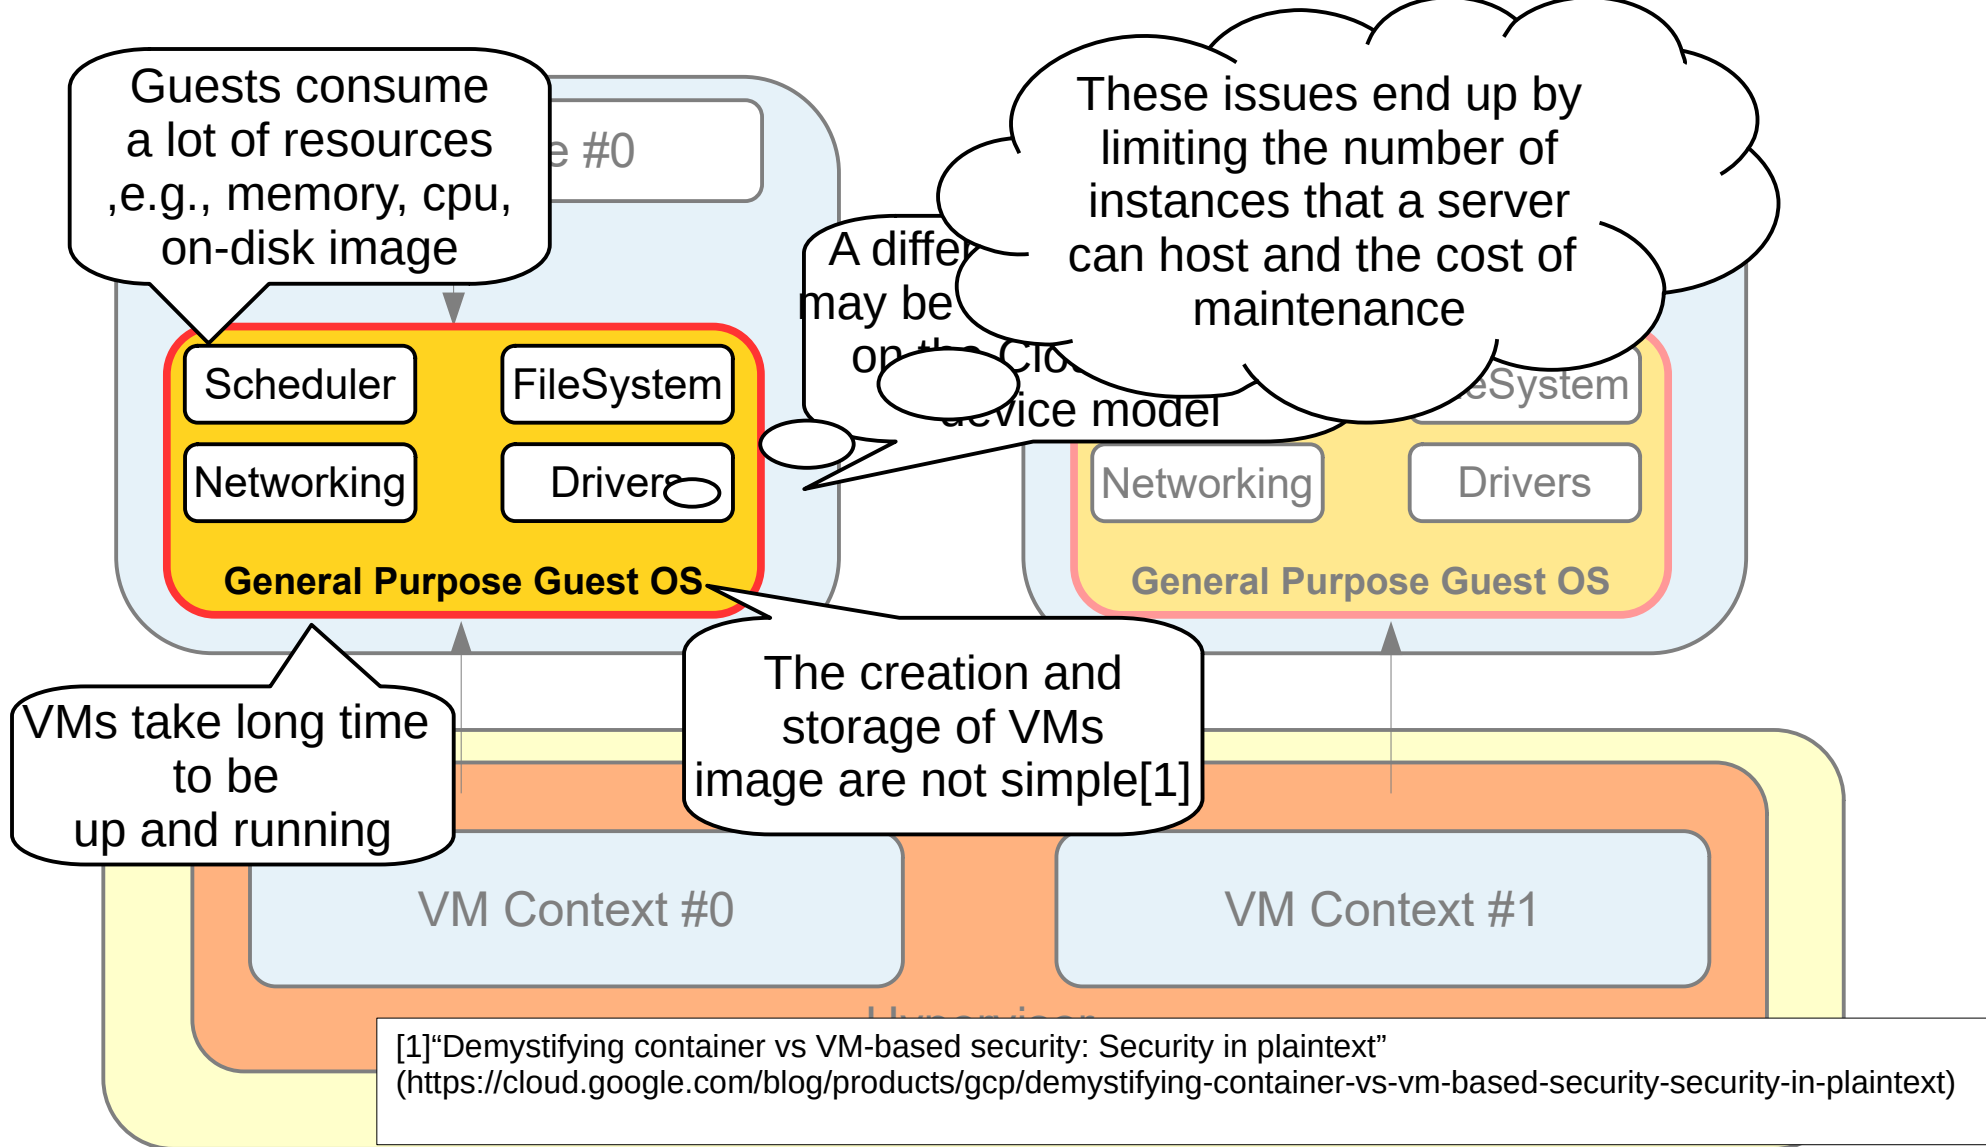


How are microservices deployed?









Guests consume a lot of resources, e.g., memory, cpu, on-disk image

These issues end up by limiting the number of instances that a server can host and the cost of maintenance

A difference may not be on the CIO

Scheduler

FileSystem

Networking

Drivers

Networking

Drivers

Toro is an application-oriented unikernel that allows **microservices** to run efficiently in **VMs** thus leveraging the strong isolation VMs provide.

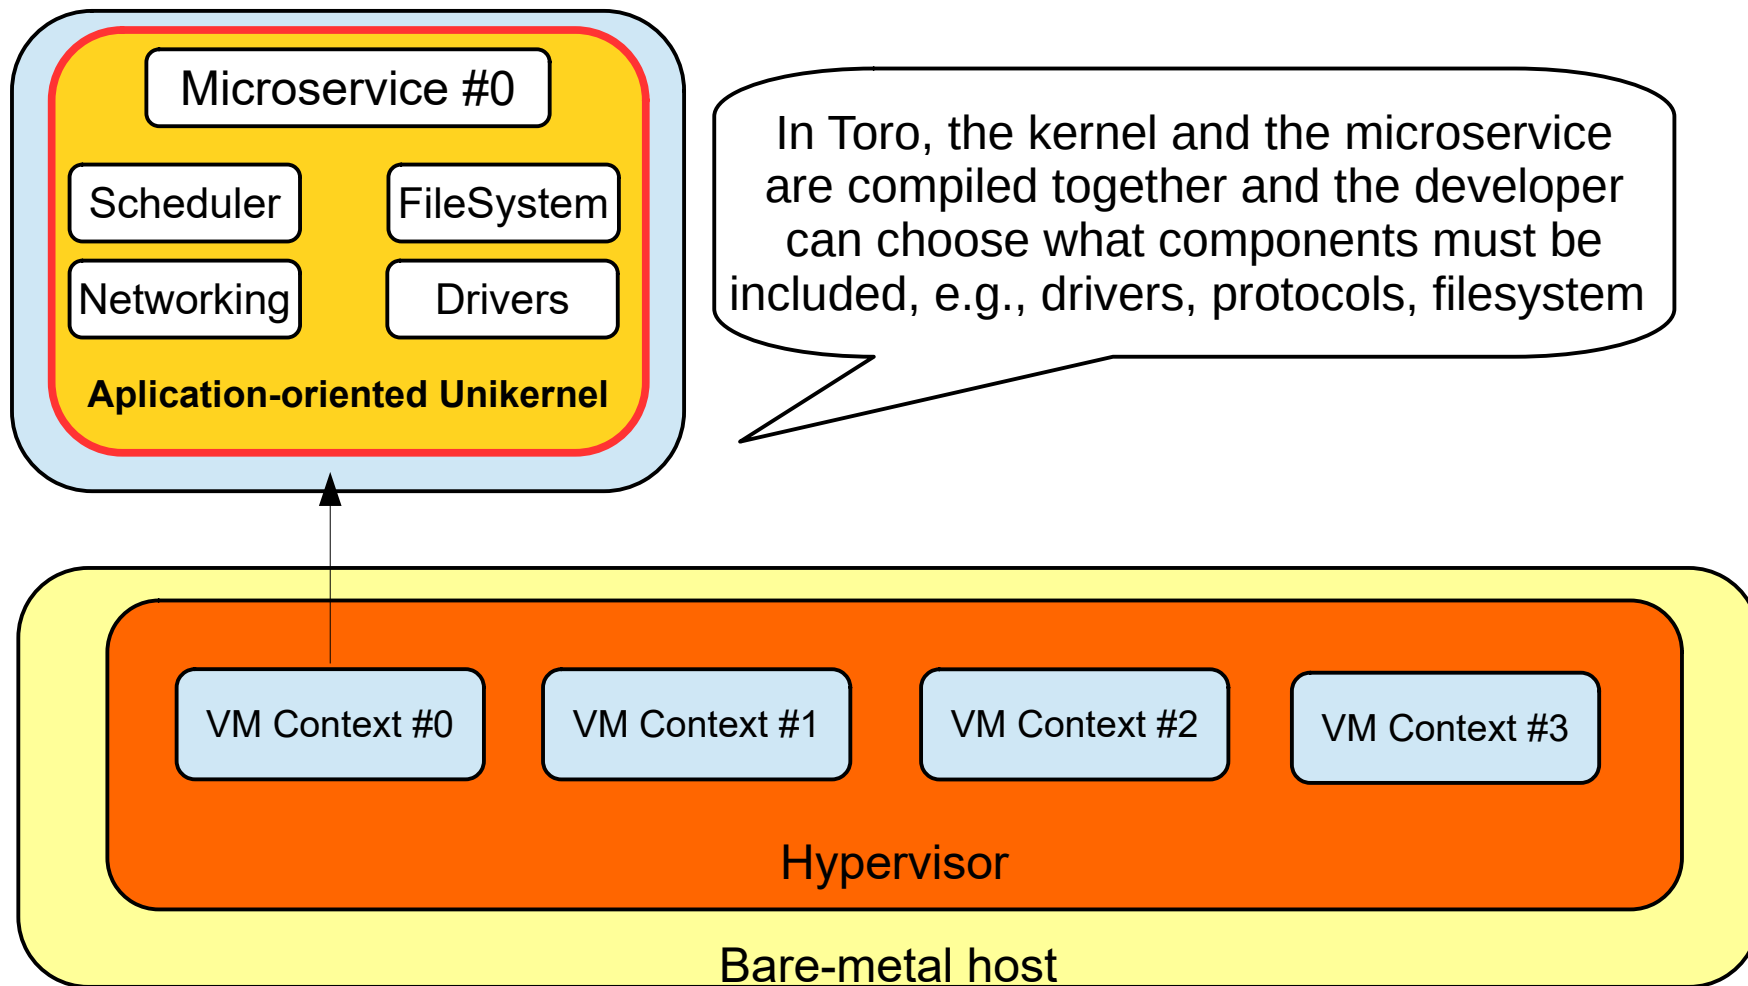
VMs take long time to be up and running

storage of VMs image are not simple[1]

VM Context #0

VM Context #1

[1] "Demystifying container vs VM-based security: Security in plaintext"
(<https://cloud.google.com/blog/products/gcp/demystifying-container-vs-vm-based-security-security-in-plaintext>)



Ingredients for ToroKernel

- A kernel's architecture that improves the execution of a single user application which is based on:
 - Simple Scheduler
 - Dedicated Resources
- A kernel's architecture that improves the execution of a single user application as a guest by providing:
 - Fast instantiation of user application
 - Reduced guest's footprint
 - Improved networking and filesystem by relying on VirtIO devices

Ingredients for ToroKernel

- A kernel's architecture that improves the execution of a single user application which is based on:
 - Simple Scheduler
 - Dedicated Resources
- A kernel's architecture that improves the execution of a single user application as a guest by providing:
 - Fast instantiation of user application
 - Reduced guest's footprint
 - Improved networking and filesystem by relying on VirtIO devices

Simple Scheduler Requirements


- Scale with the number of cores
- Simple scheduler's algorithm, i.e., Cooperative threading model
- Minimalist context switches

Simple Scheduler

In Toro, there are only threads

Thread 1

Thread 2



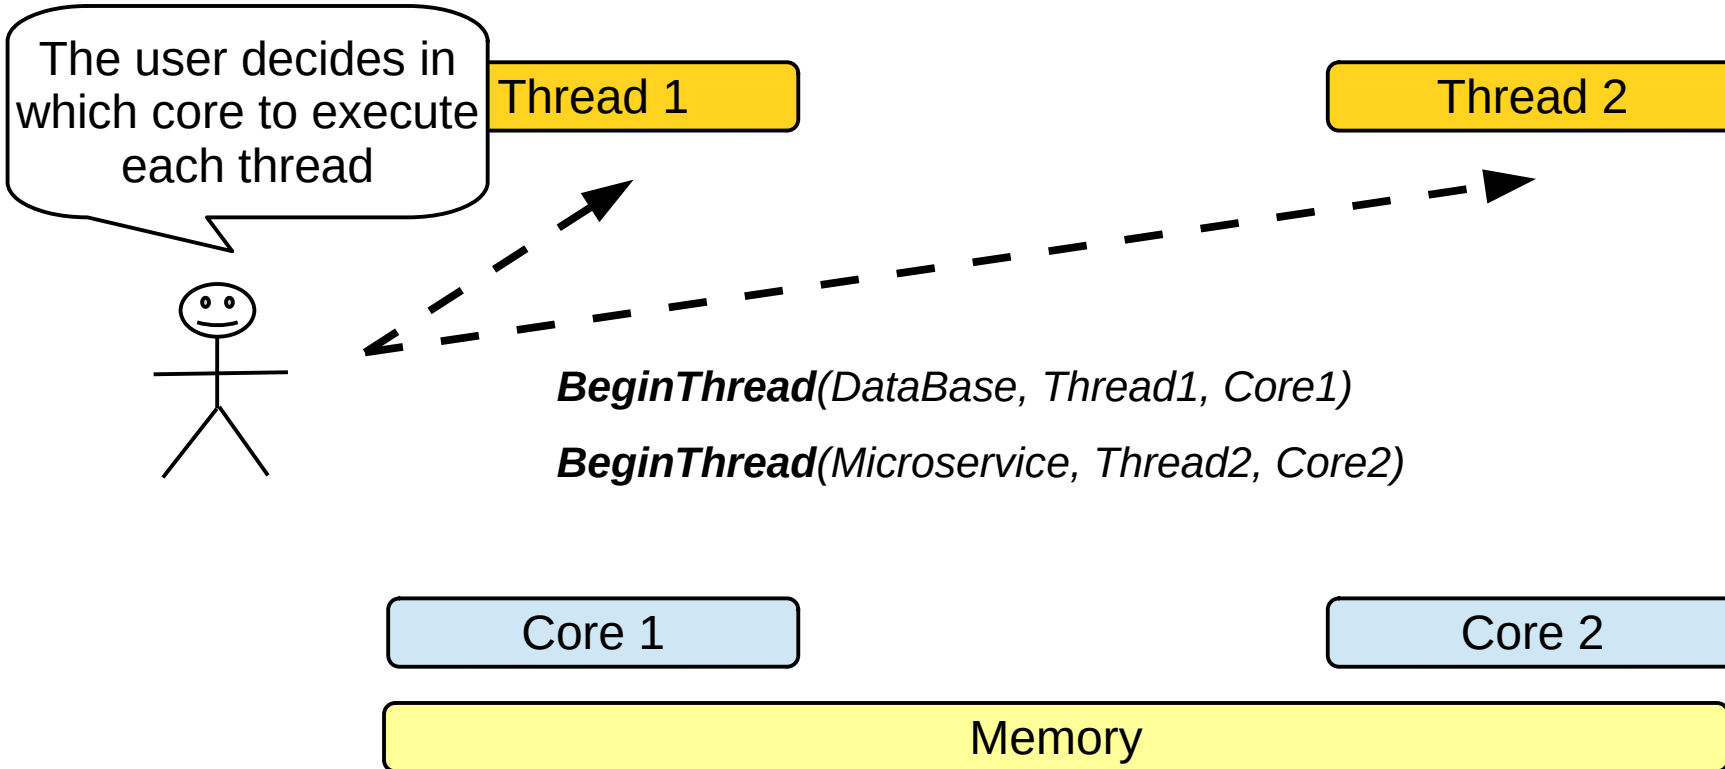
Core 1

Core 2

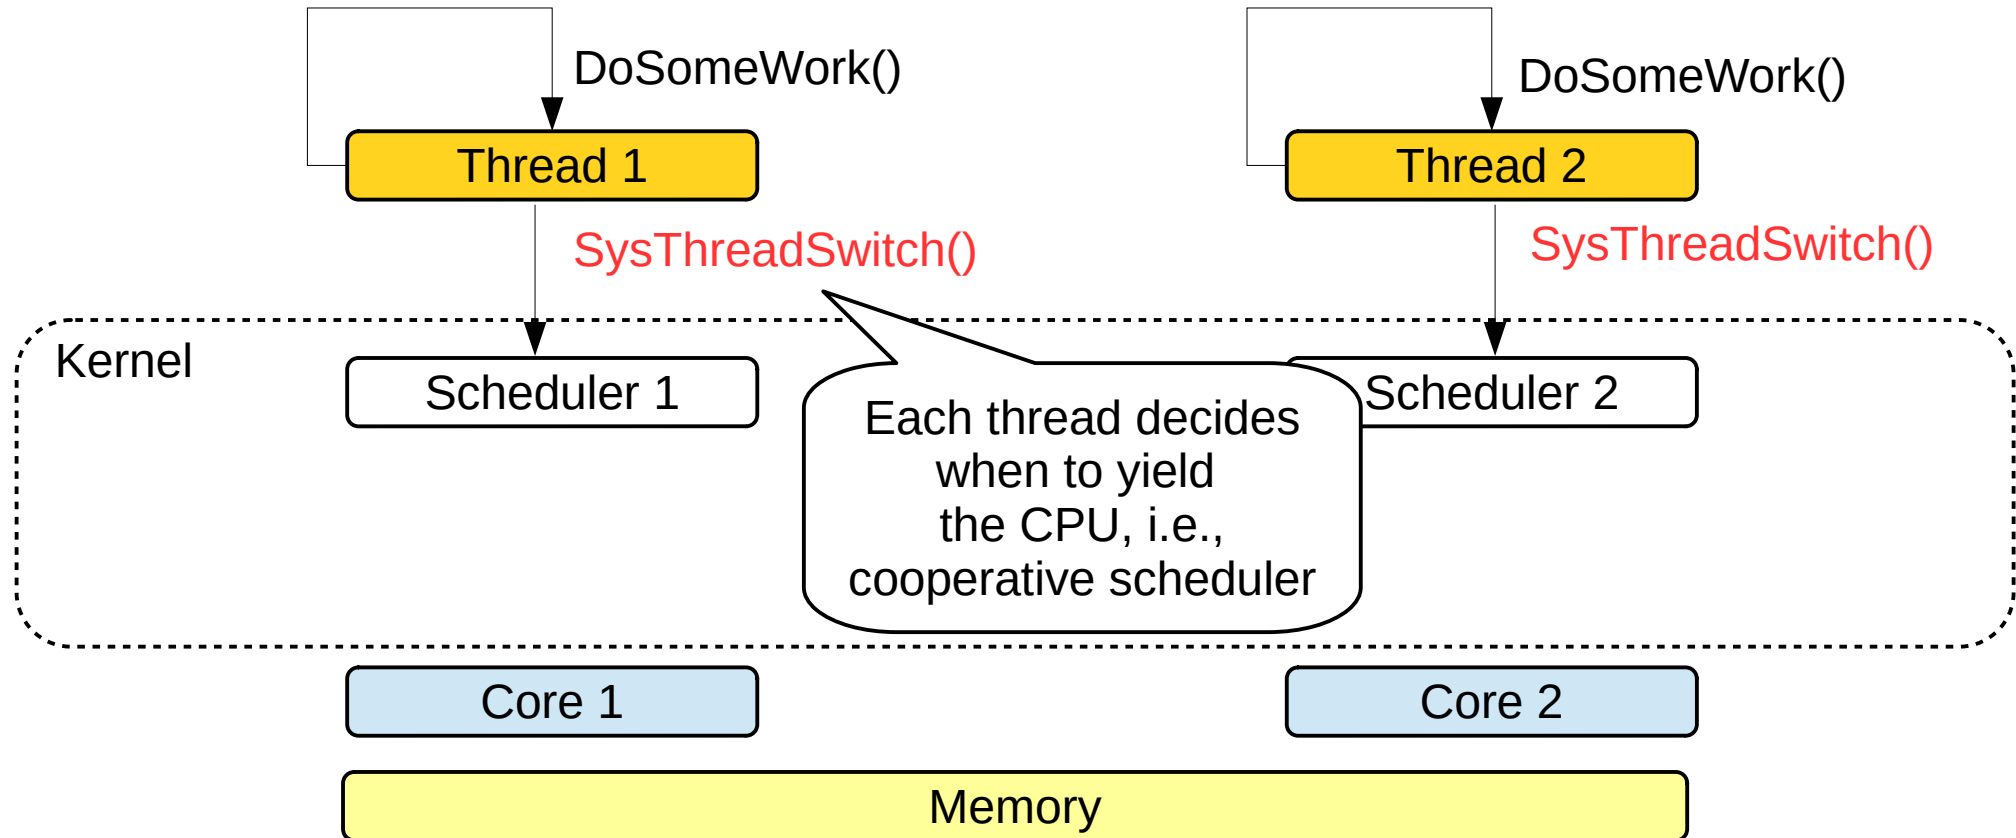
Threads share the memory space

Memory

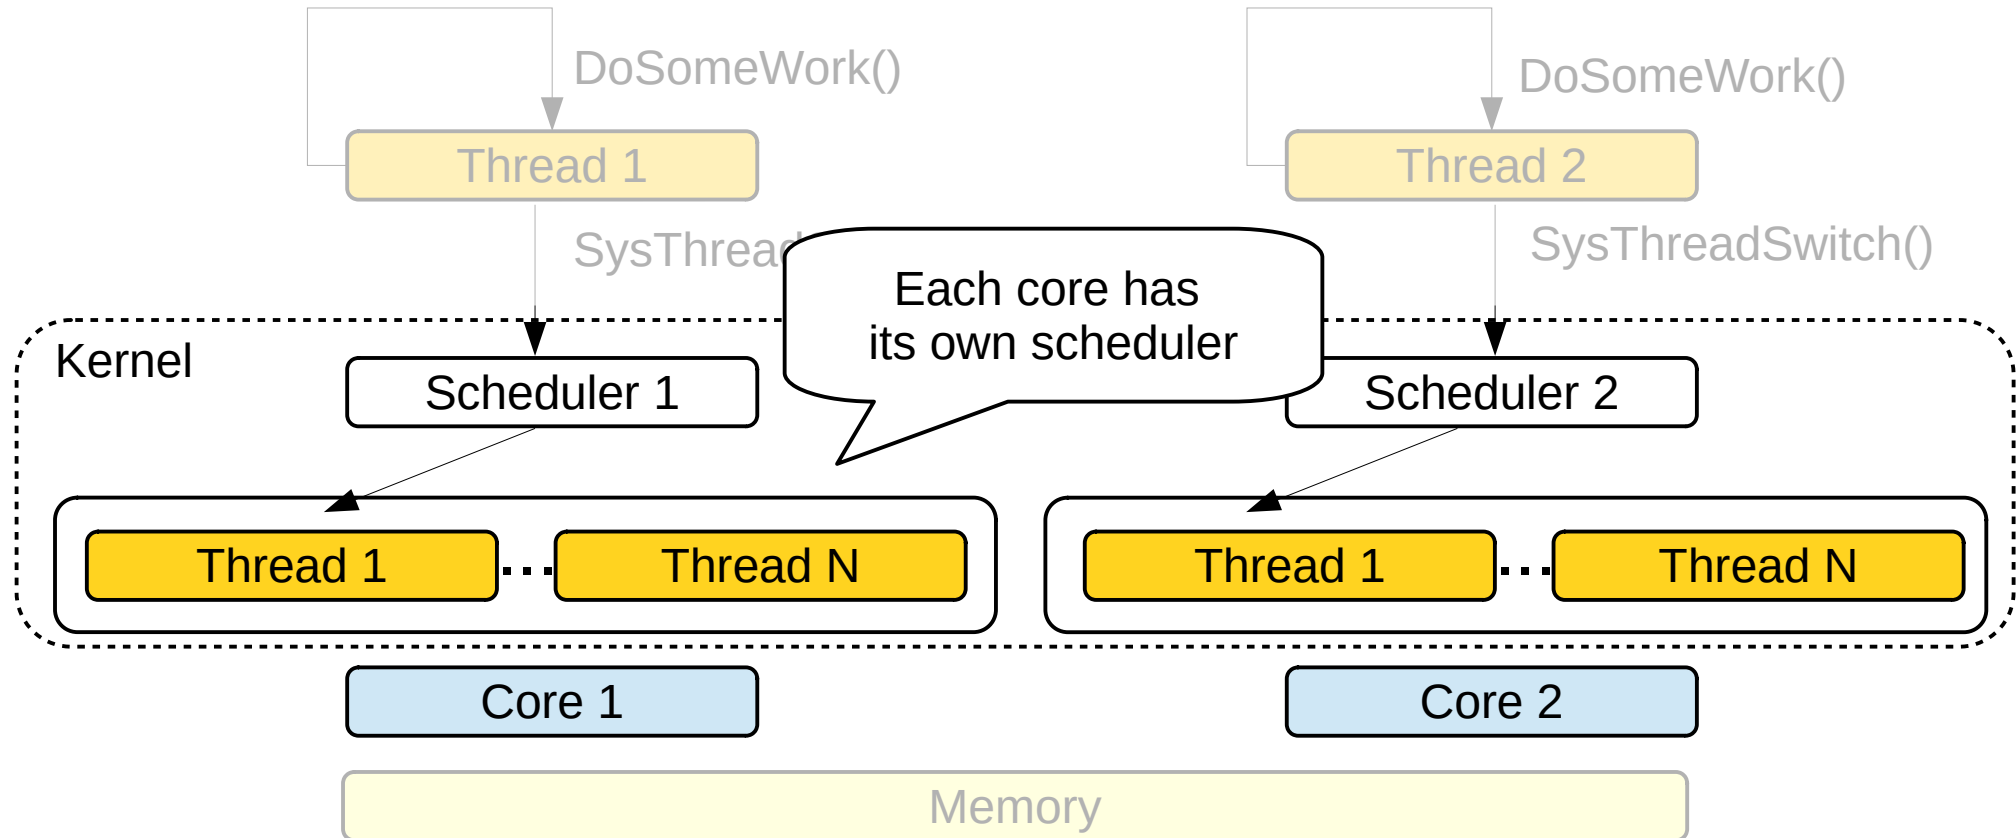
Simple Scheduler



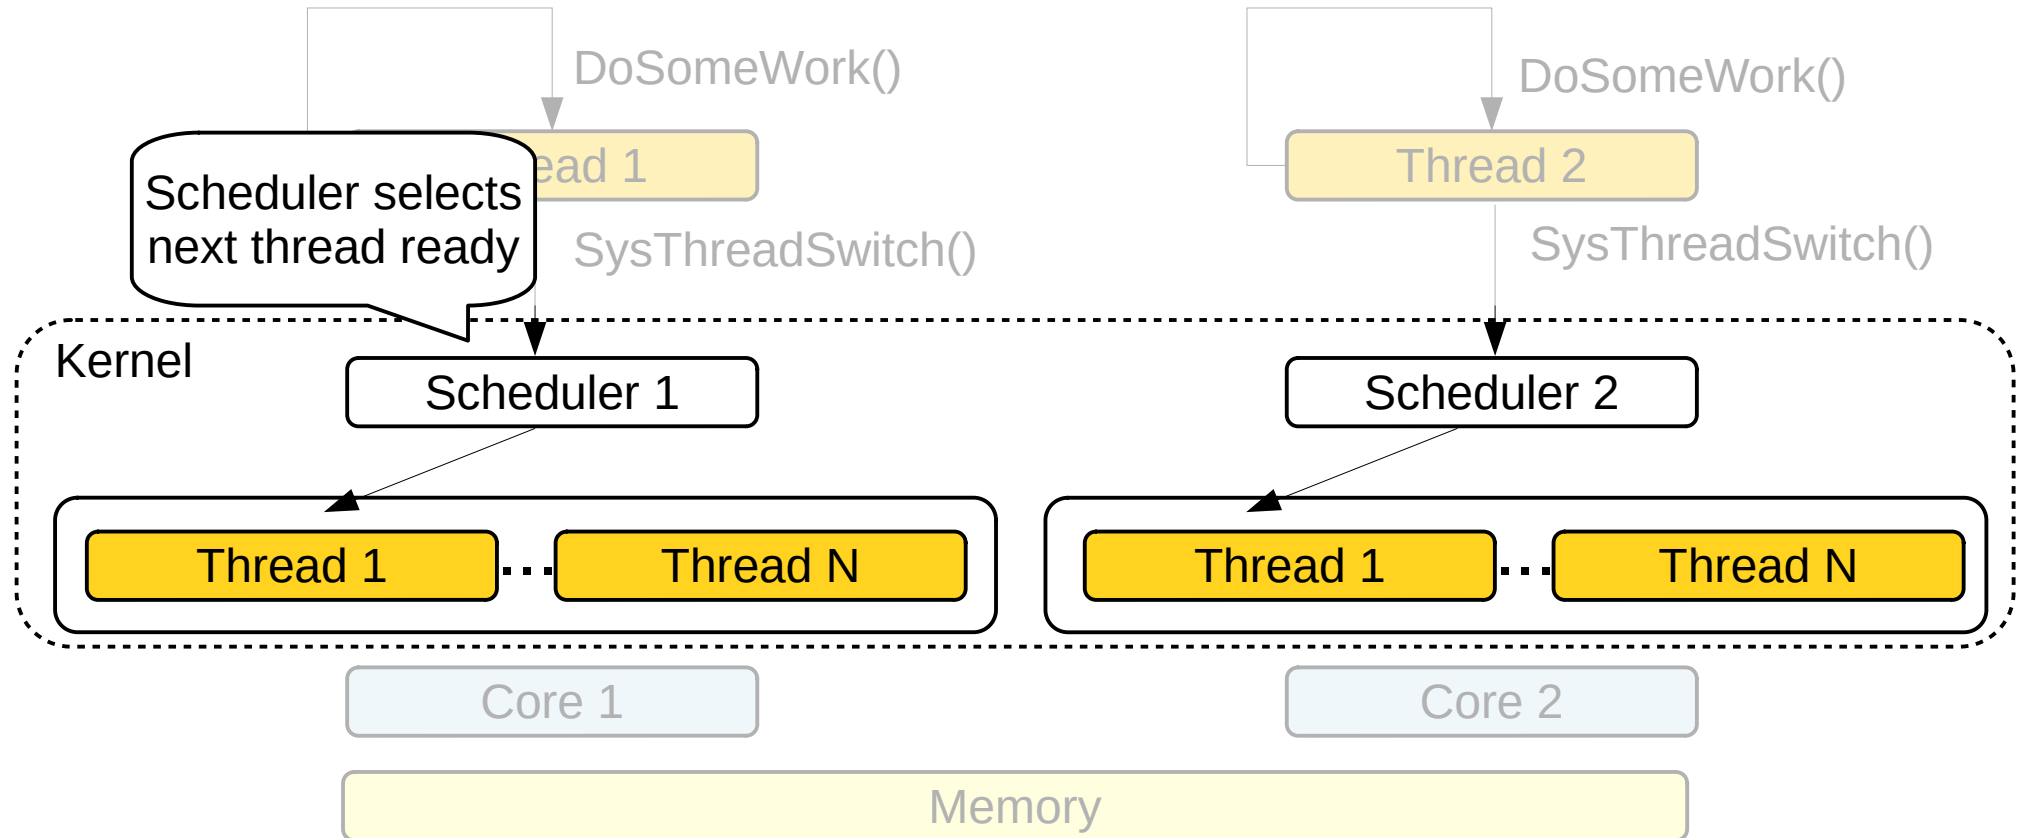
Simple Scheduler



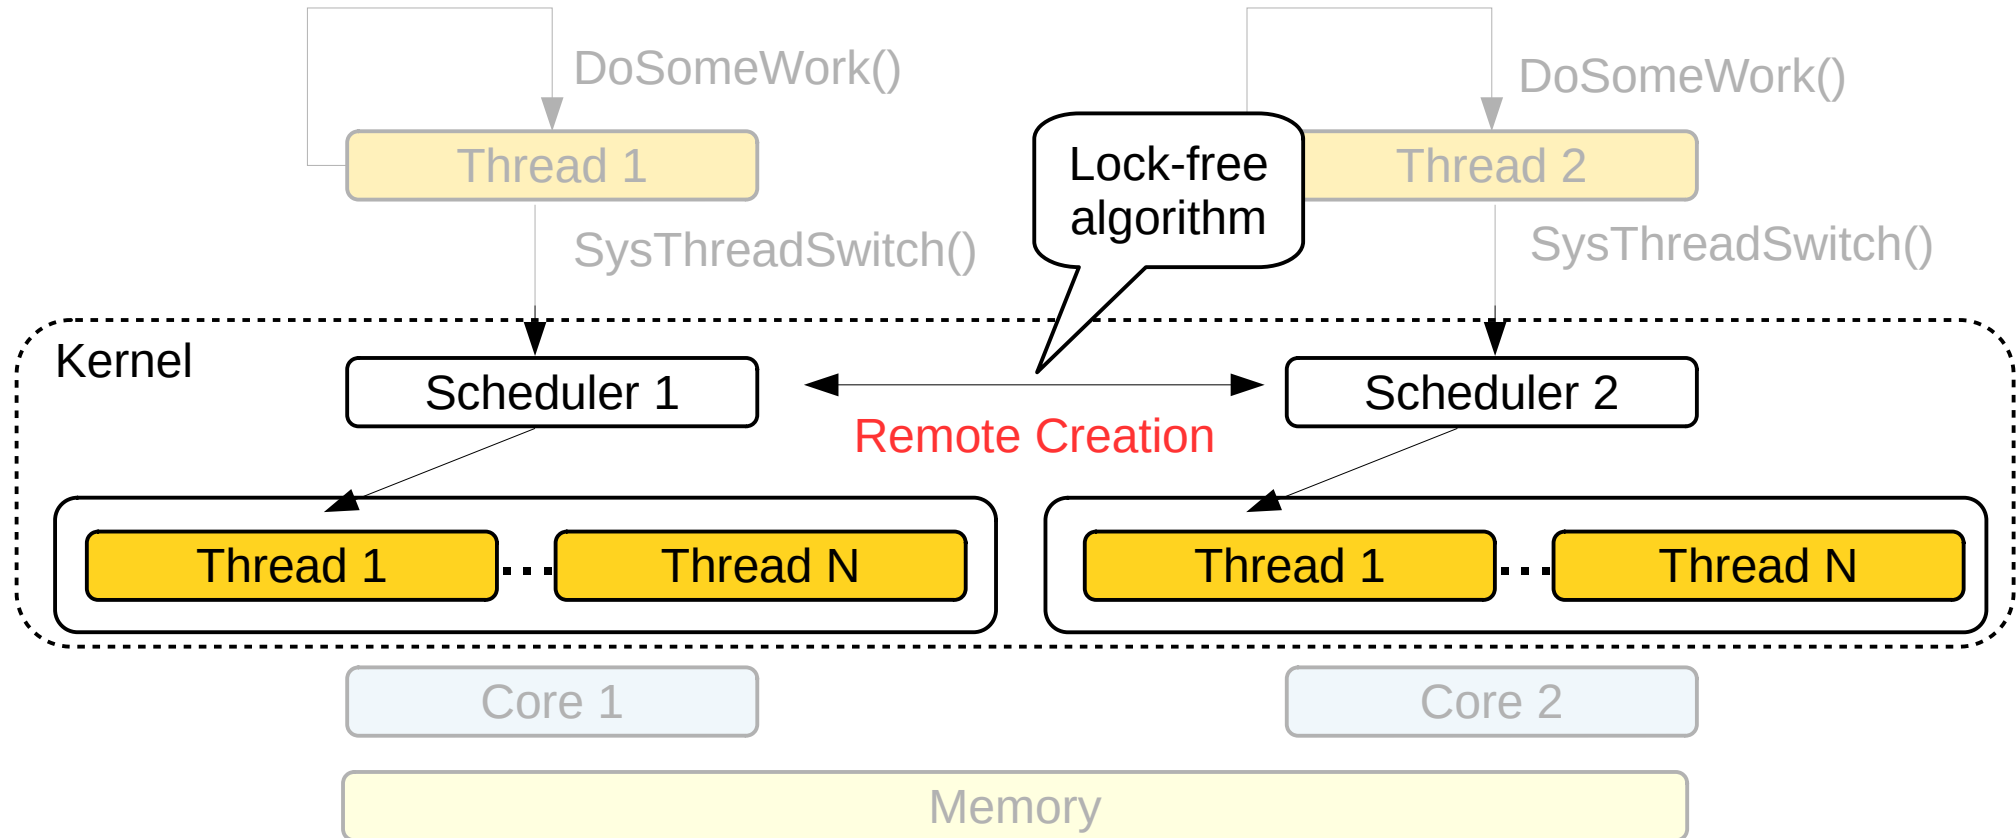
Simple Scheduler



Simple Scheduler



Simple Scheduler



Benefits of Simple Scheduler

- The use of threads makes context switching cheaper
- The use of a cooperative scheduler:
 - simplifies user's and kernel's code by avoiding to implement protection inside the scheduler
 - reduces number of context switches since it does not relies on interruptions

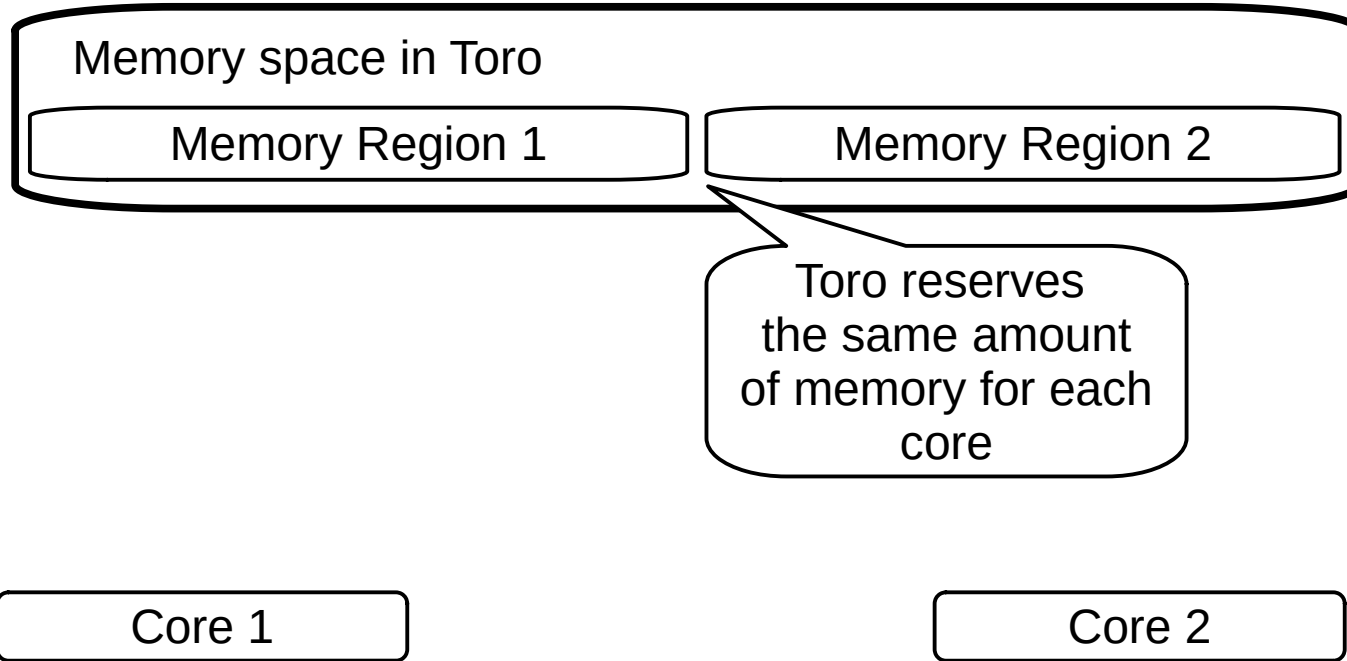
Dedicated Resources

- In a multicore system, the problematic resource is the shared memory. The use of shared memory causes:
 - Overhead in the memory bus
 - Overhead in the cache to keep it coherent
 - Overhead to guaranty mutual exclusion when spinlocks are used

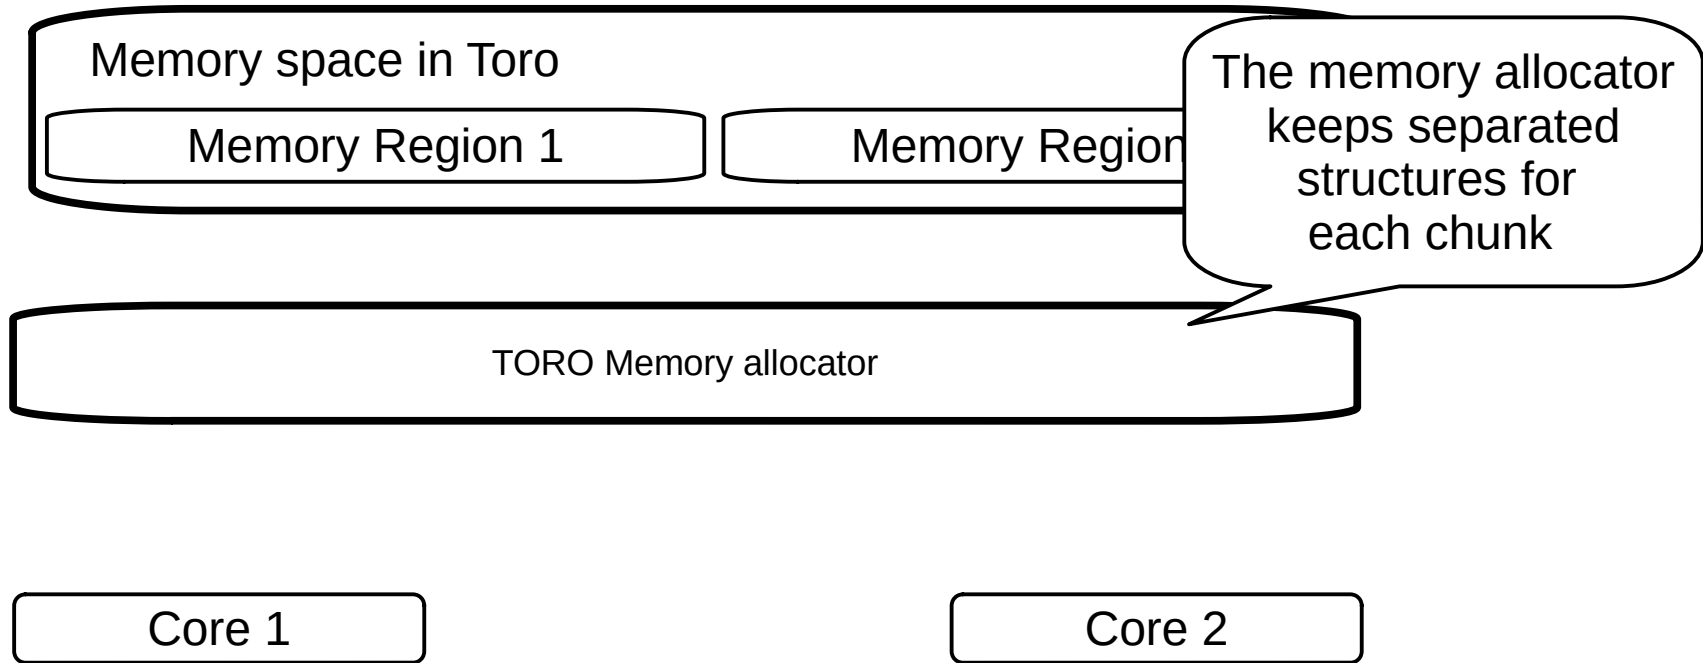
Dedicated Resources Proposal

- Toro improves memory access by
 - keeping the resources locals:
 - The memory is dedicated per core
 - The kernel data structures are dedicated per core
 - The access to kernel data structures is lock free
 - leveraging NUMA technologies, e.g., Hypertransport, Intel QuickPath

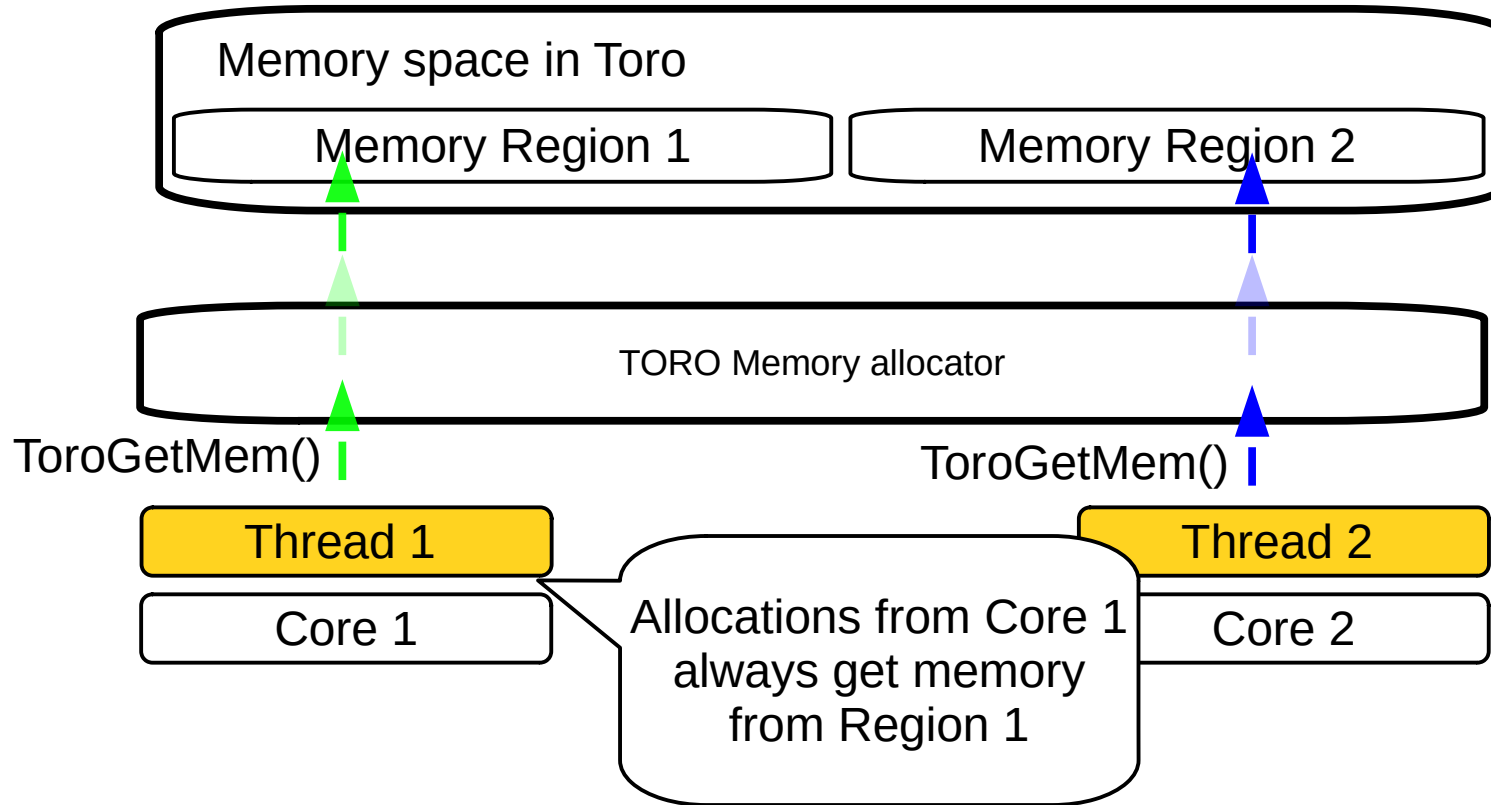
Dedicated Resources Memory



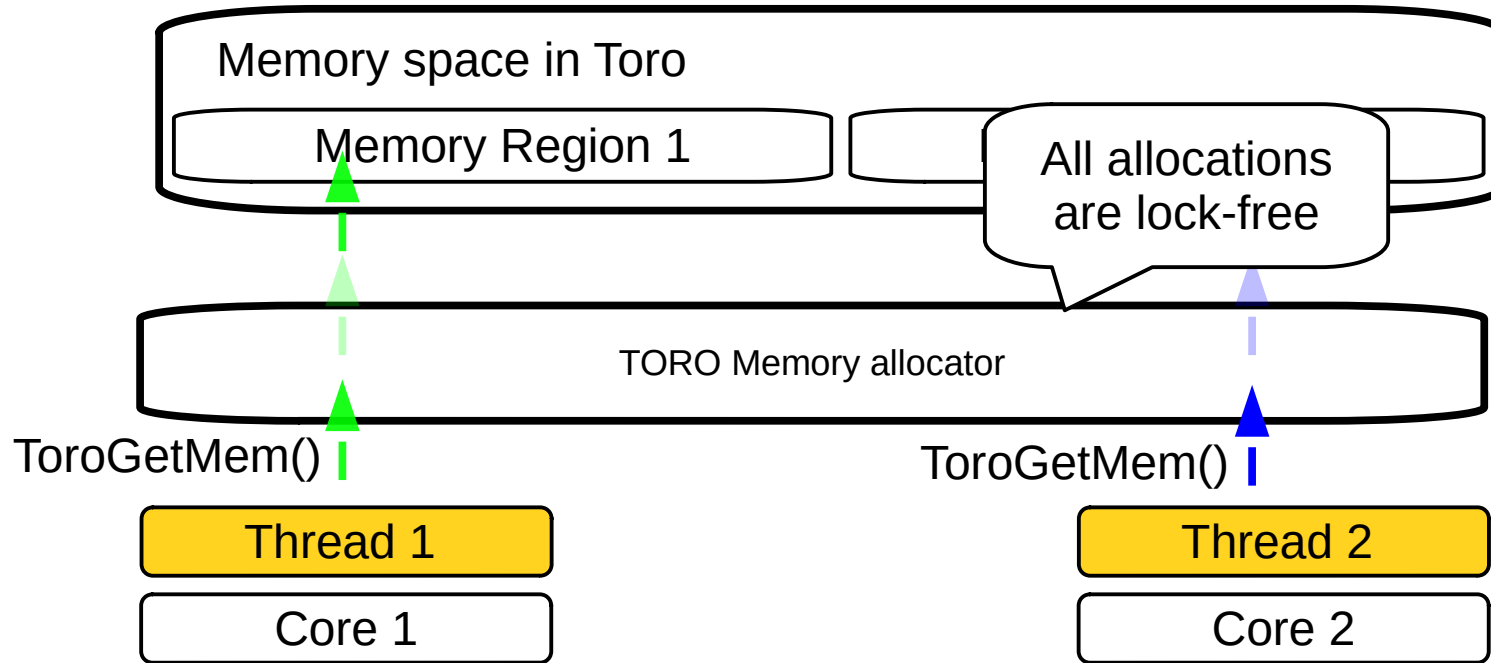
Dedicated Resources Memory



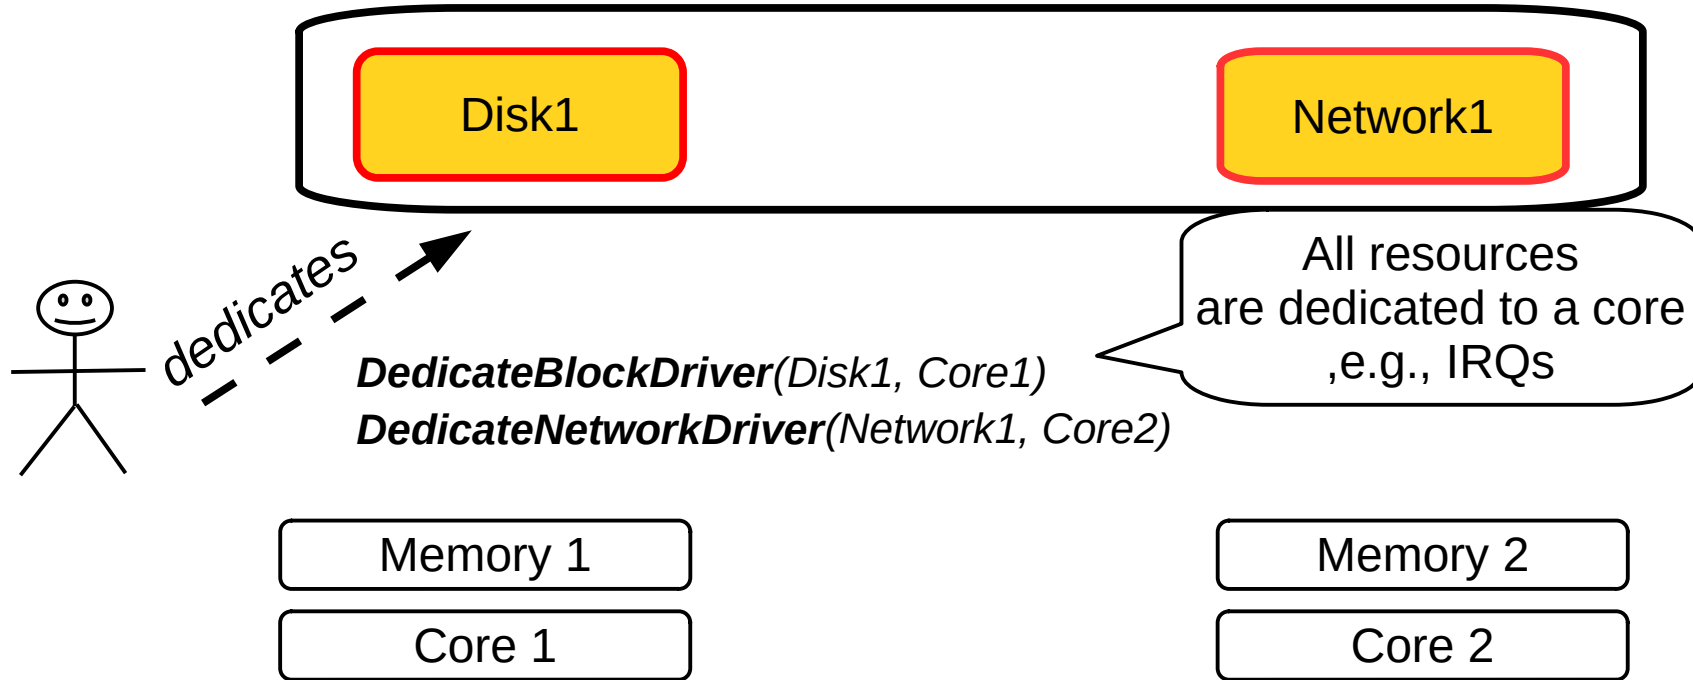
Dedicated Resources Memory



Dedicated Resources Memory

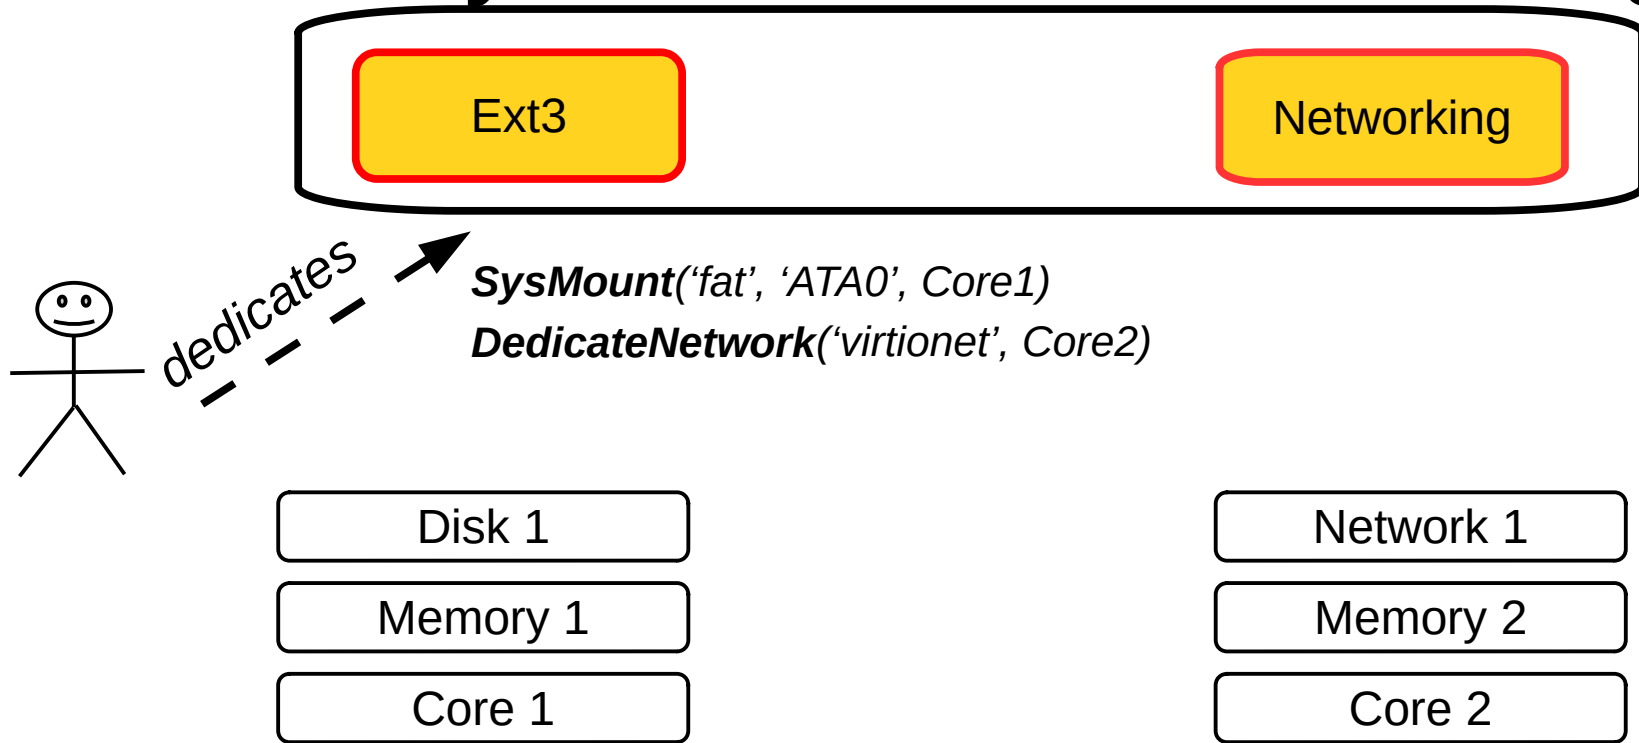


Dedicated Resources Devices



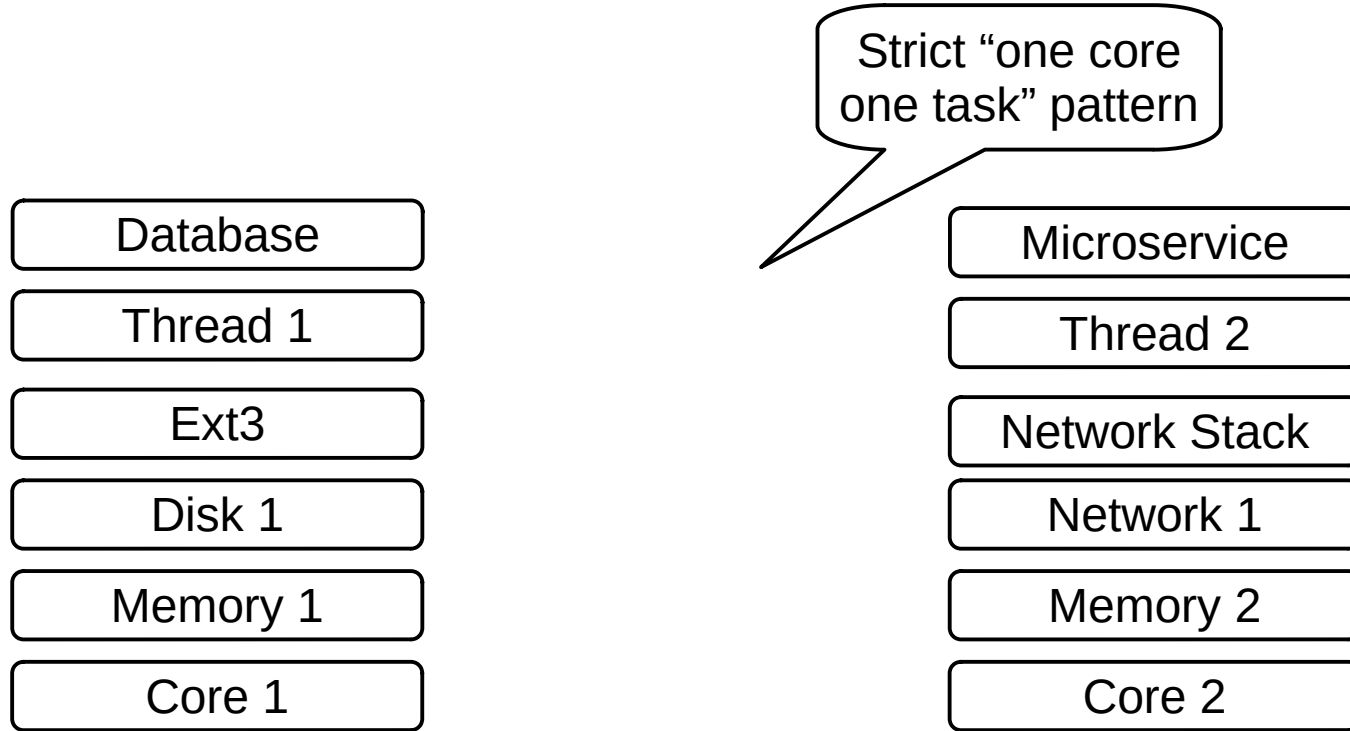
Dedicated Resources

Filesystems and Networking



Dedicated Resources

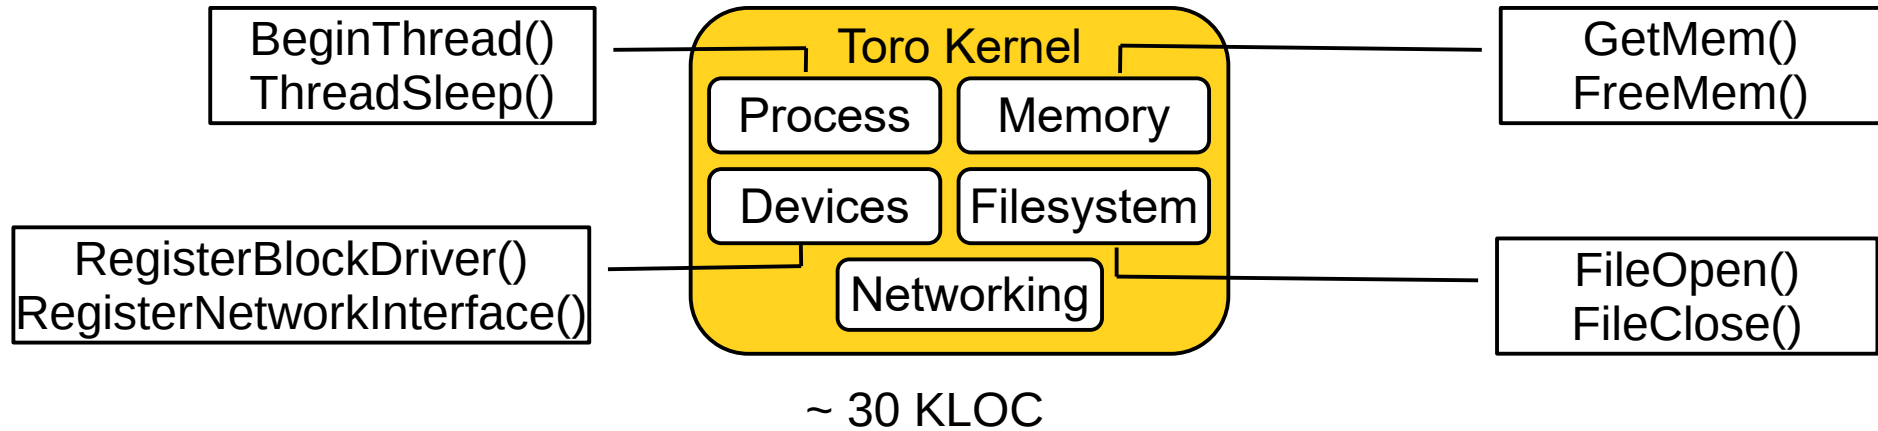
Filesystems and Networking



Ingredients for ToroKernel

- A kernel's architecture that improves the execution of a single user application which is based on:
 - Simple Scheduler
 - Dedicated Resources
- A kernel's architecture that improves the execution of a single user application as a guest by providing:
 - Fast instantiation of user application
 - Reduced guest's footprint
 - Improved networking and filesystem by relying on VirtIO devices

Application-oriented Kernel

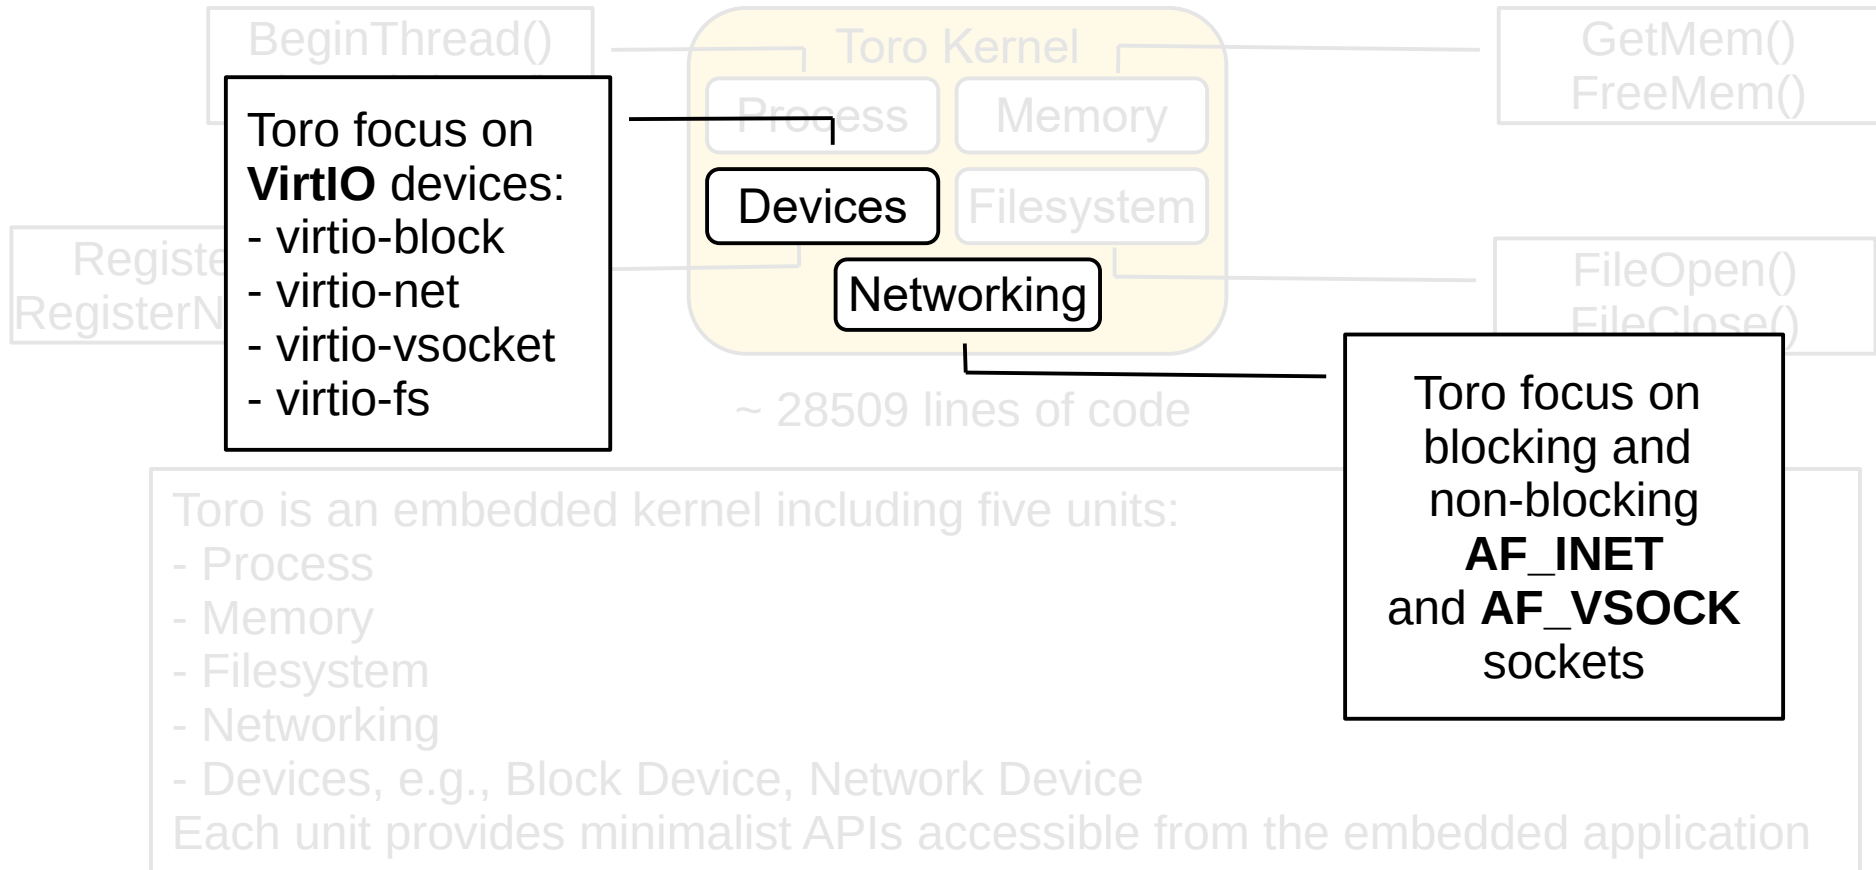


Toro is an embedded kernel including five units:


- Process
- Memory
- Filesystem
- Networking
- Devices, e.g., Block Device, Network Device

Each unit provides minimalist APIs accessible from the embedded application

Application-oriented Kernel




Application-oriented Kernel



- User application and kernel units are compiled in a single binary
- The application includes only the component required

Application-oriented Kernel



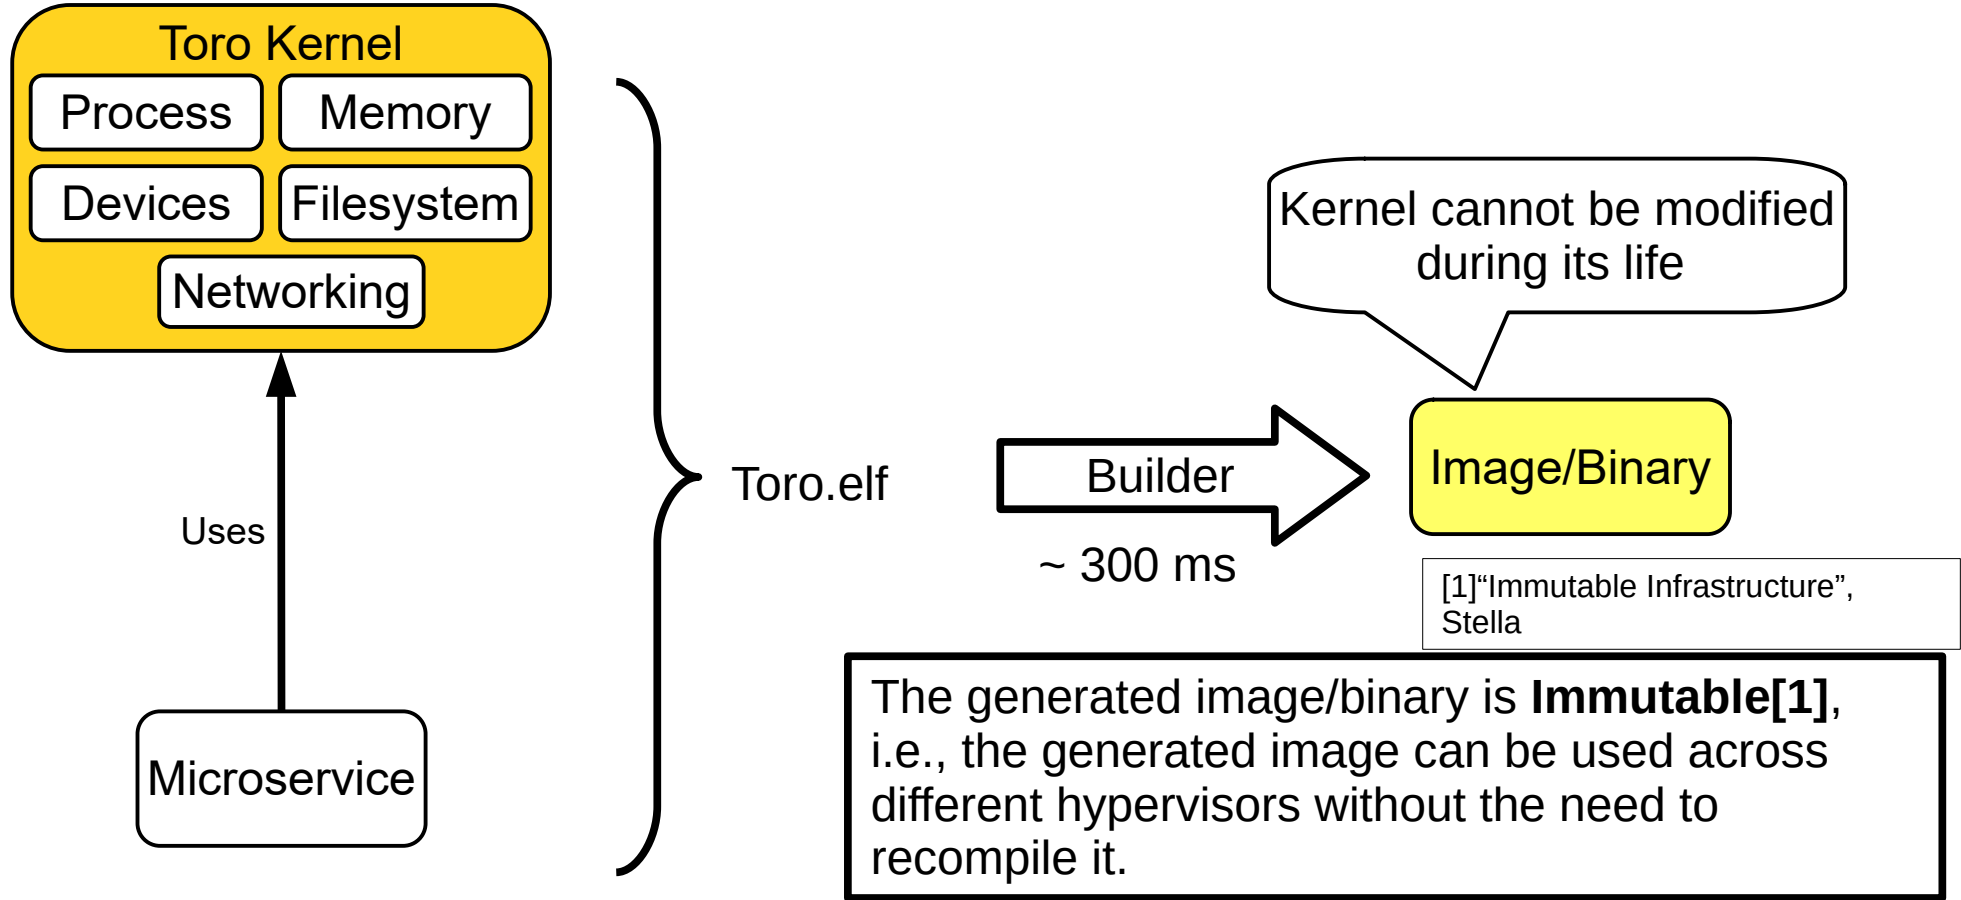
Uses

Microservice

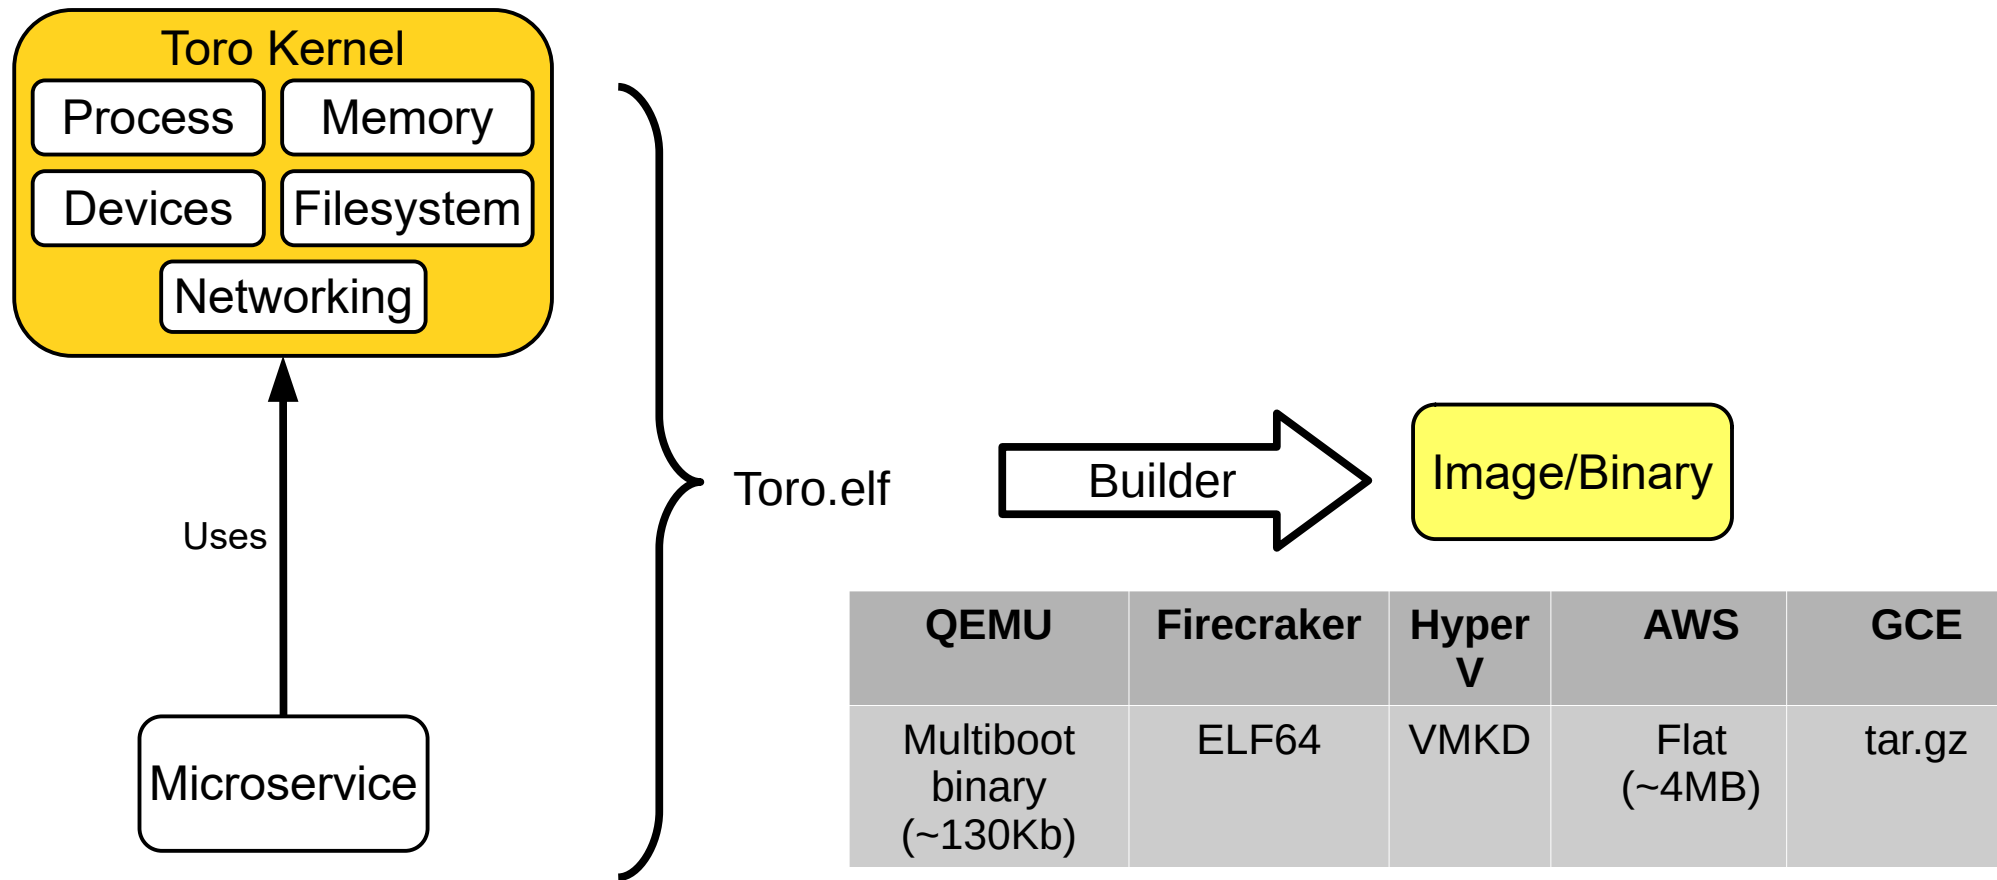
- User application and kernel units are compiled in a single binary
- The com

```
program MyMicroservice;  
uses  
    Memory,  
    Filesystem,  
    Process,  
    Networking,  
    Ext3,  
    E1000;  
Begin  
    //  
    // Your Code  
    //  
End.
```

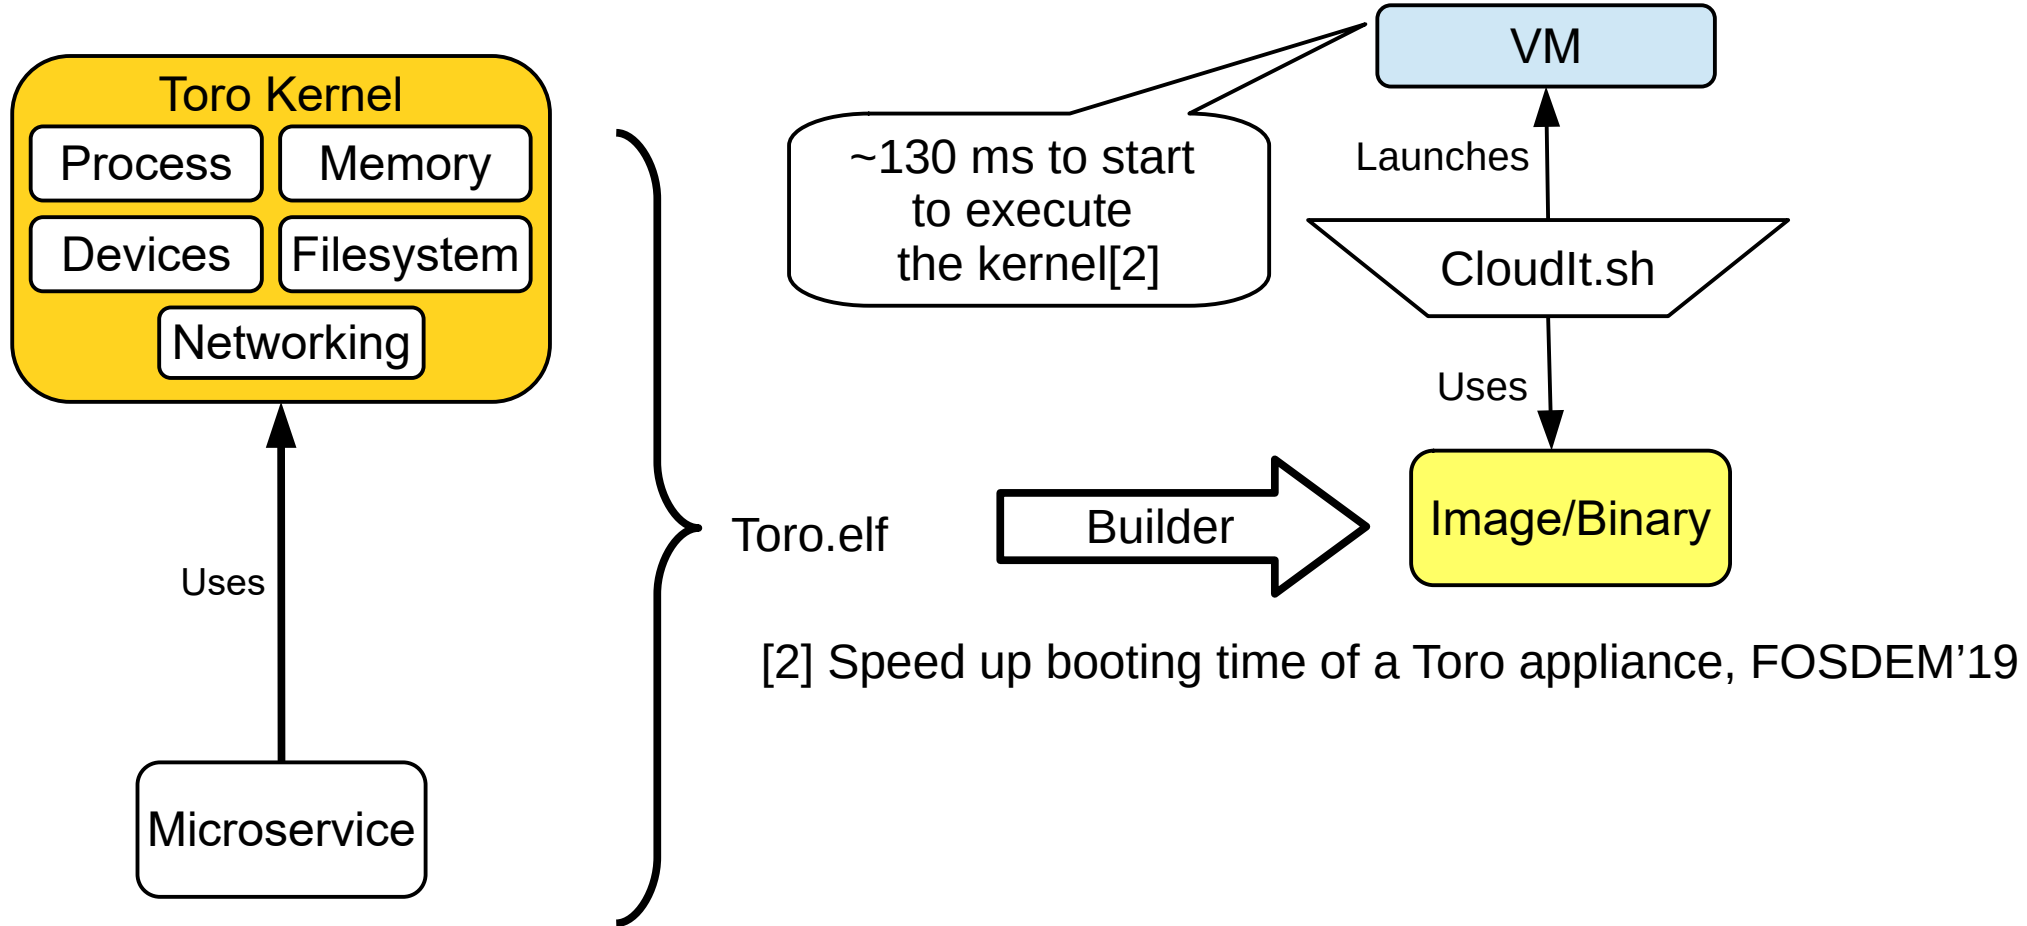

Application-oriented Kernel



Application-oriented Kernel



Application-oriented Kernel



Time until kernel starts to execute

4 cores Intel(R) Atom(TM) CPU C2550 @ 2.40GHz
8 GB of physical memory

Approach	Image (~ 4MB)	Binary (~ 130kB)	Binary with QBoot
QEMU/KVM (2.5.0)	1457 ms	452 ms	132 ms
NEMU (#39af42)		309 ms	95 ms
Firecracker (0.14.0)		17ms	

```
$ echo "Hello World!"  
avg: 2.629263ms
```

<https://blog.iron.io/the-overhead-of-docker-run/>

Benefits of Application-oriented Kernel

- Security is based on the hypervisor which ensures a reduced attack surface
- The kernel size is smaller since it only includes the units required for the application
- The smaller kernel size reduces the booting time and eases image manipulation, i.e., your program is your kernel!
- No communication overhead since application is using the kernel APIs

Benefits of Application-oriented Kernel

- Security is based on the hypervisor which ensures a reduced attack surface
 - The kernel is responsible for security
 - The application is responsible for security
- "It's all talk until the code runs." - Ward Cunningham
- No communication overhead since application is using the kernel APIs

Example: The Static WebServer

- Simple microservice that serves files by using the HTTP protocol
 - <https://github.com/torokernel/torokernel/tree/master/examples/StaticWebServer> (among other examples ;))
- For development, we base on Windows
- For deployment, we base on a baremetal Linux Host (KVM guest) on Scaleway

Development Workbench

Lazarus IDE v2.1.0 r61784 - StaticWebServer

File Edit Search View Source Project Run Package Tools Window Help

Standard Additional Common Controls Dialogs Data Controls Data Access System SQLldb Misc LazControls SynEdit Chart

Source Editor

StaticWebServer

```
//  
1 program StaticWebServer;  
.  
{$IFDEF FPC}  
{$mode delphi}  
25 {$ENDIF}  
.  
2 {%RunCommand qemu-system-x86_64 -m 256 -smp 1 -drive format=raw,file=StaticWebServer.img -net nic,model=virtio -net tap,ifname=TAP2 -drive file=fat  
3 {%RunFlags BUILD-}  
.  
3 uses  
Kernel in '..\..\rtl\Kernel.pas',  
Process in '..\..\rtl\Process.pas',  
Memory in '..\..\rtl\Memory.pas',  
Debug in '..\..\rtl\Debug.pas',  
35 Arch in '..\..\rtl\Arch.pas',  
Filesystem in '..\..\rtl\Filesystem.pas',  
Pci in '..\..\rtl\drivers\Pci.pas',  
// Ide in '..\..\rtl\drivers\IdeDisk.pas',  
{$IFDEF UseVirtIOFS}  
40 VirtIOFS in '..\..\rtl\drivers\VirtIOFS.pas',  
VirtIOVSocket in '..\..\rtl\drivers\VirtIOVSocket.pas',  
{$ELSE}  
VirtIOBlk in '..\..\rtl\drivers\VirtIOBlk.pas',  
// Ext2 in '..\..\rtl\drivers\Ext2.pas',  
45 Fat in '..\..\rtl\drivers\Fat.pas',  
VirtIONet in '..\..\rtl\drivers\VirtIONet.pas',
```

During development, we use as block driver **virtio-blk** and for networking **virtio-net**

270: 1 INS C:\Users\Matias\Desktop\torokernel\master-debug\examples\StaticWebServer\StaticWebServer.pas

Development Workbench

Lazarus IDE v2.1.0 r61784 - StaticWebServer

File Edit Search View Source Project Run Package Tools Window Help

Standard Additional Common Controls Dialogs Data Controls Data Access System SQLdb Misc LazControls SynEdit Chart

Source Editor

StaticWebServer

```
//
20 program StaticWebServer;
.
. {$IFDEF FPC}
. {$mode delphi}
25 {$ENDIF}
.
. [%RunCommand qemu-system-x86_64 -m 256 -smp 1 -drive format=raw -hda tap,ifname=TAP2 -drive file=fat
. [%RunFlags BUILD-}
```

1

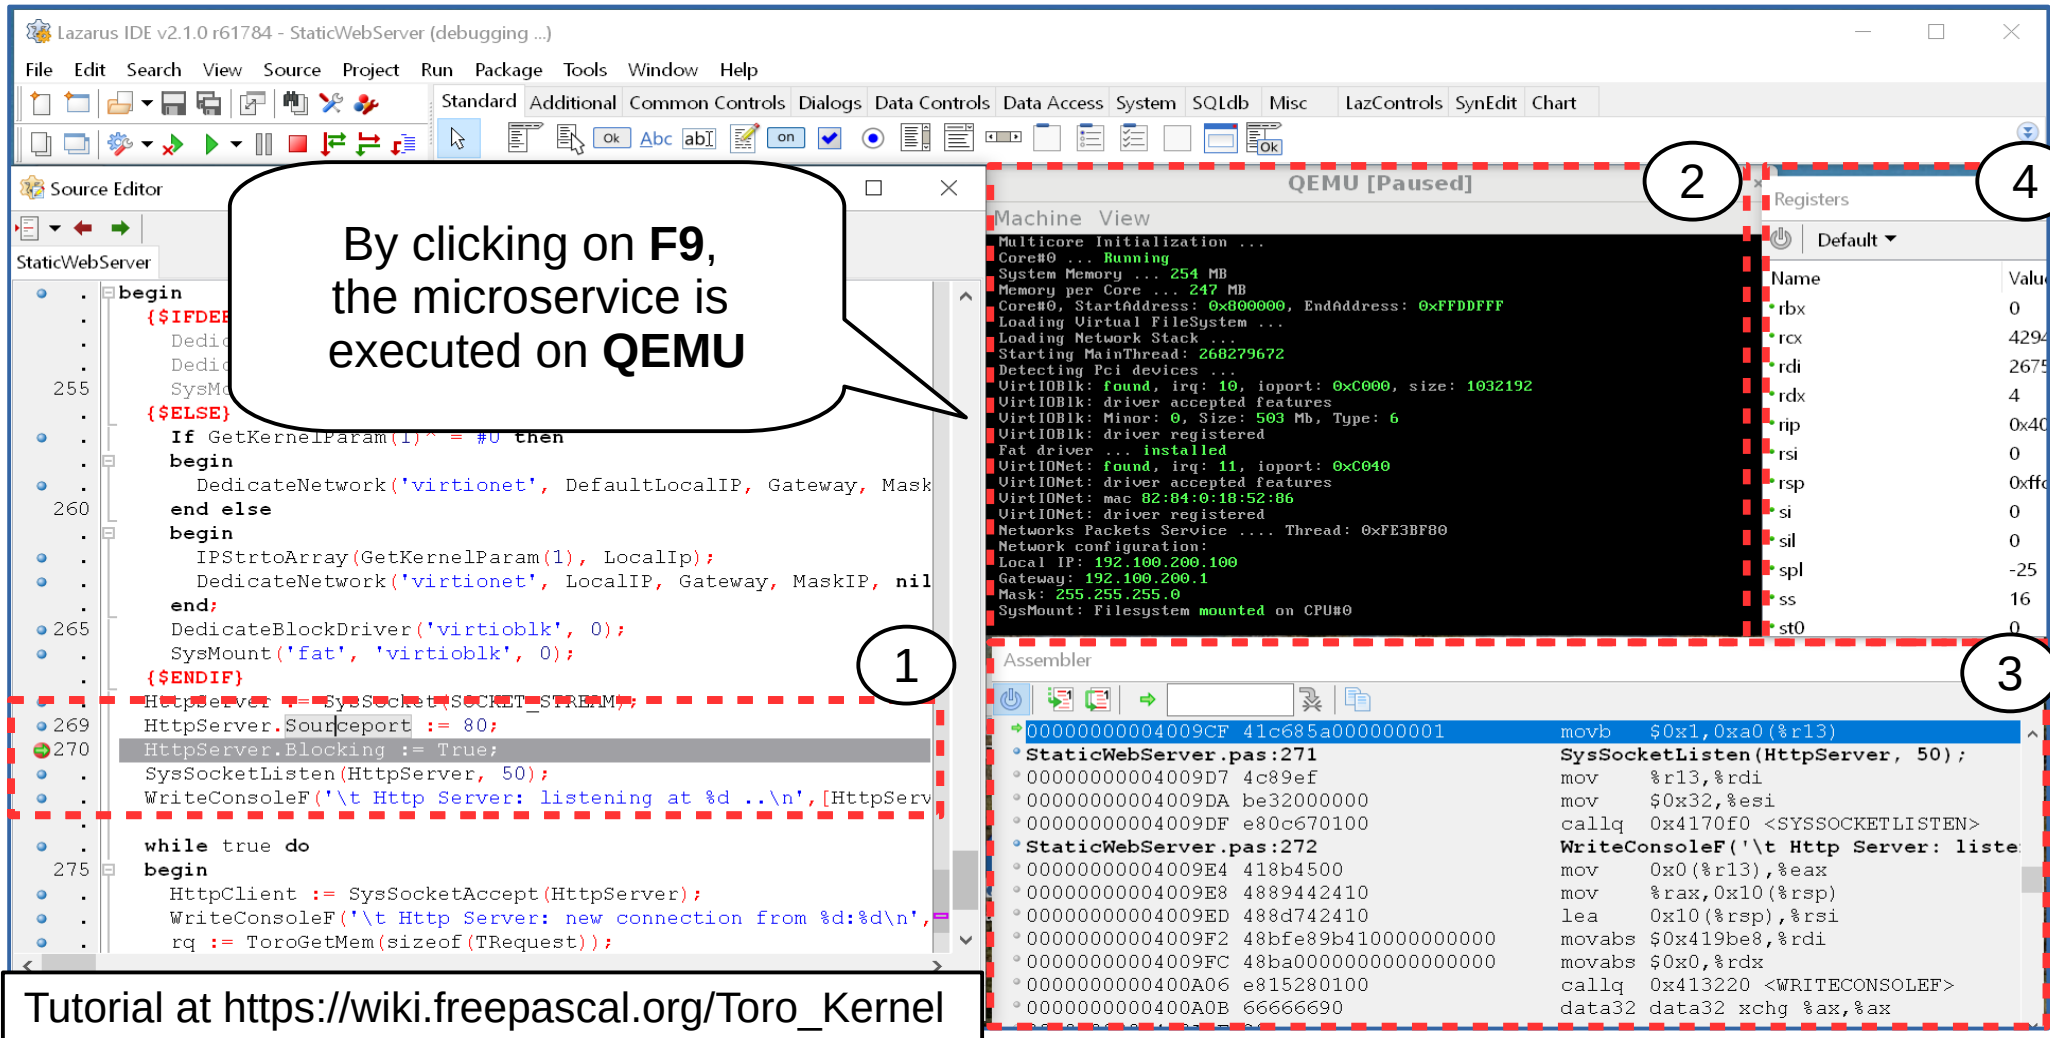
Messages

Project: Executing command before: Success
target OS: win64 for x64
Compiling ..\..\builder\build.pas
Linking ..\..\builder\build.exe
66 lines compiled, 2.1 sec, 80688 bytes code, 4900 bytes data
Project checks, Hints: 1
Note: passing compiler option -Cg twice with different values
Compile Project, OS: linux, Target: StaticWebServer: Success, Hints: 1
StaticWebServer.pas(165,3) Note: Local variable "indexSize" is assigned but never used
Project: Executing command after: Success
StaticWebServer: Detected ELF64 format
Writing .text section ...
Writing .data section ...
Building Image ...
Entry point : 4194544

35: 42 INS C:\Users\Matias\Desktop\torokernel\master-debug\examples\StaticWebServer\StaticWebServer.pas

By clicking on **Ctrl + F9**, the application and the kernel are compiled and an image is generated

Debugging the kernel from the IDE



Deployment Workbench

KVM Guest (**Scaleway** Baremetal Host)

\$ **CloudIt.sh StaticWebServer** "-append virtiovsocket,virtiofs,myfstoro"

```
Booting from ROM..Loading Toro ... HEAD:a990437
Multicore Initialization ...
Core#0 ... Running
System Memory ... 4094 MB
Memory per Core ... 3063 MB
Core#0, StartAddress: 0x800000, EndAddress: 0xBFFDFFFF
Loading Virtual FileSystem ...
Loading Network Stack ...
Starting MainThread: 3221077912
Detecting Pci devices ...
VirtIOFS: Detected device tagged: myfstoro, queues: 1
VirtIOFS: Device accepted features
VirtIOFS: Queue 0, size: 1024, initiated, irq: 11
Fat driver ... installed
VirtIOVSocket: found, irq: 10, ioport: 0xC040
VirtIOVSocket: Guest ID: 3
VirtIOVSocket: RX_QUEUE was initiated
VirtIOVSocket: EVENT_QUEUE was initiated
VirtIOVSocket: TX_QUEUE was initiated
VirtIOVSocket: driver registered
Networks Packets Service .... Thread: 0xBE83BFA8
DedicateNetworkSocket: success on core #0
SysMount: Filesystem mounted on CPU#0
26/10/2019-21:07:33 Http Server: listening at 80 ..
```

Check on <http://www.torokernel.io>

Summary

- Toro design is improved in five main points:
 - Booting time and building time
 - Kernel API with zero overhead function call
 - Access to shared memory
 - Networking
 - CPU usage

Future Work

- Improve tooling to build, test and debug microservices
- Investigate new use-cases
- Port networking applications to provide starter kits: SMTP relay, HTTP proxy, Web tracking
- Investigate the use of microVM technologies, e.g., Firecracker, NEMU.
- Investigate the use of OpenStack to ease the deployment of Toro appliances
- Compare with other approaches, e.g., unikernels, containers [3]

[3] Toro, a dedicated kernel for microservices, OSSummit'18

QA

